


Motivation: Root cause analysis



From: alice@xyz.com
To: Admin (bob@xyz.com)
Title: Help!

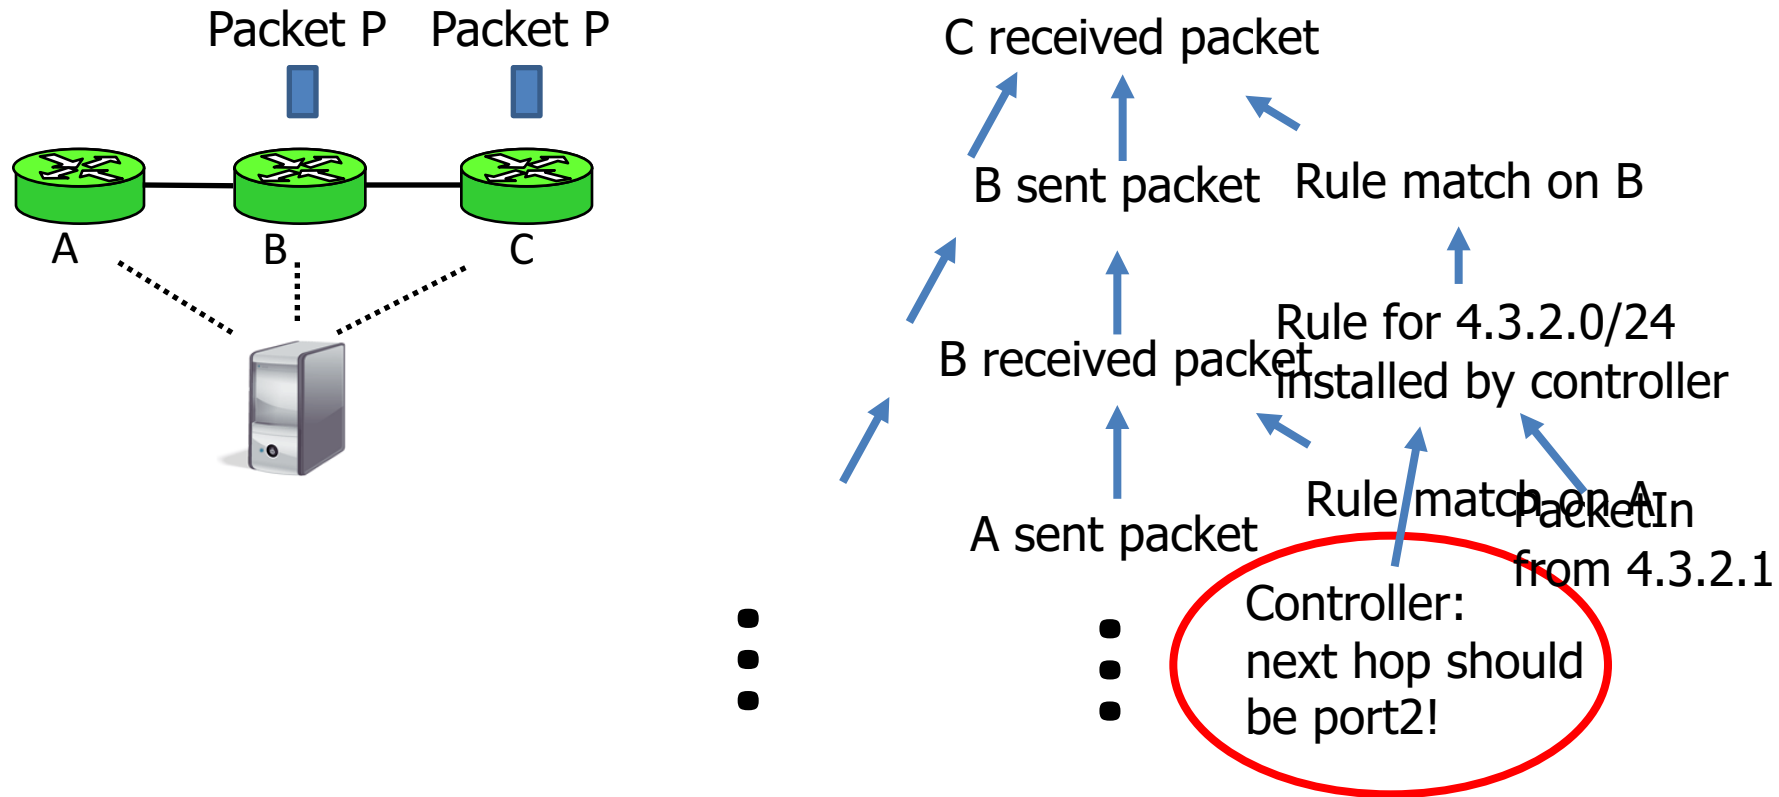
My server is receiving suspicious traffic from 4.3.2.0/24--it should have been sent to the low-security server. Packets from 4.3.3.0/24 are still being routed correctly. Can you help?

... is arriving
the wrong
server !?!




- Networks can (and frequently do!) have bugs
- We need a good debugger!

Debugging networks with provenance

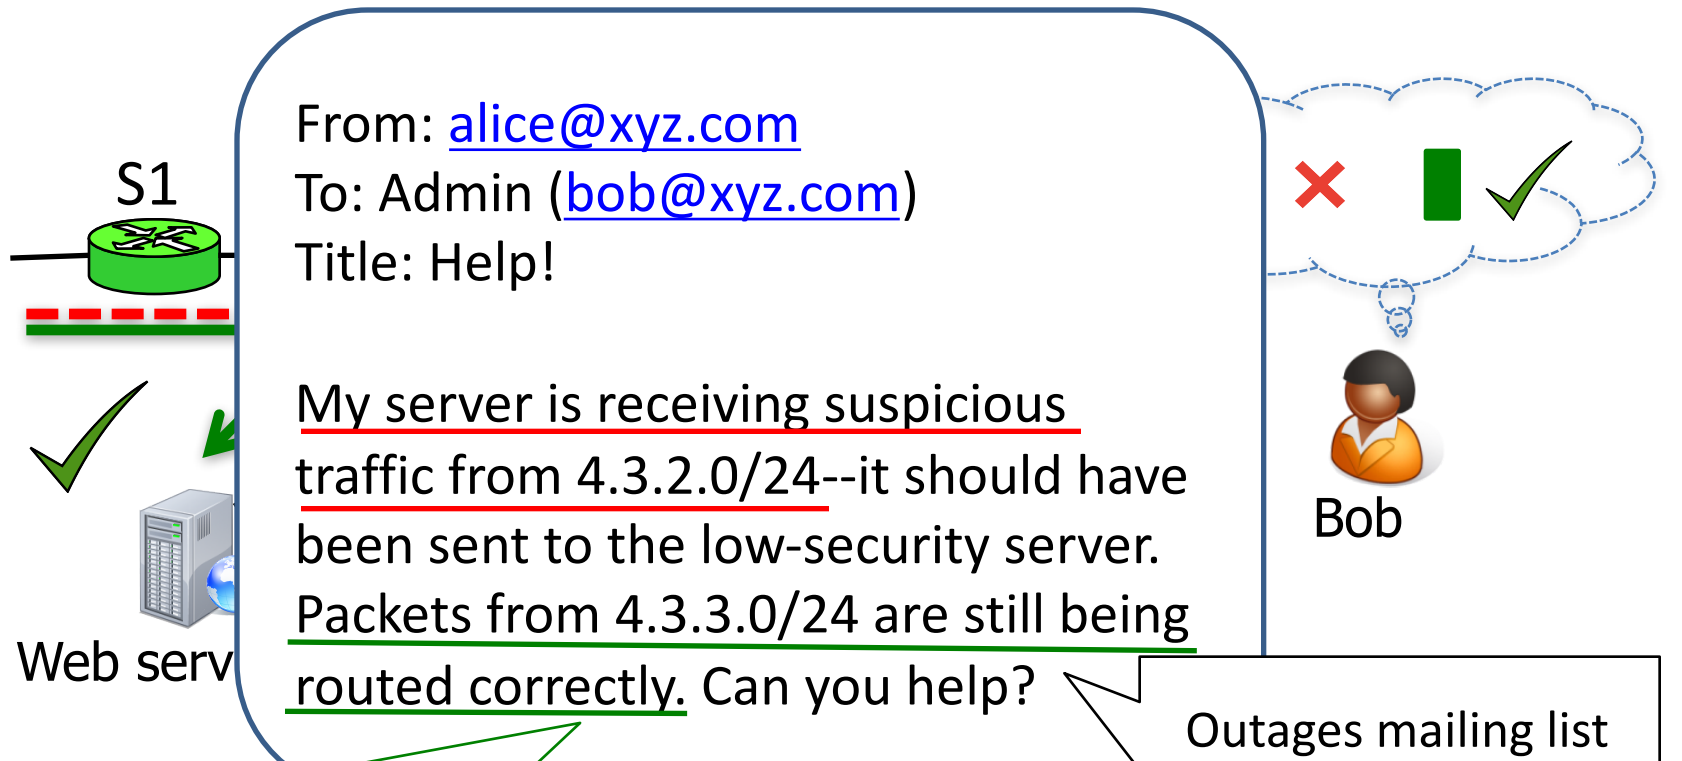


- Existing debuggers tell us what happened
 - Example: NetSight [NSDI'14]
- **Provenance** offers a richer explanation
 - Example: Y! [SIGCOMM'14]

Problem: The explanation can be too big!



What can we do?




Working reference!

Outages mailing list
Sept.—Dec. 2014:
66% have references!

... Reason about the differences bet
and the reference

Differential provenance

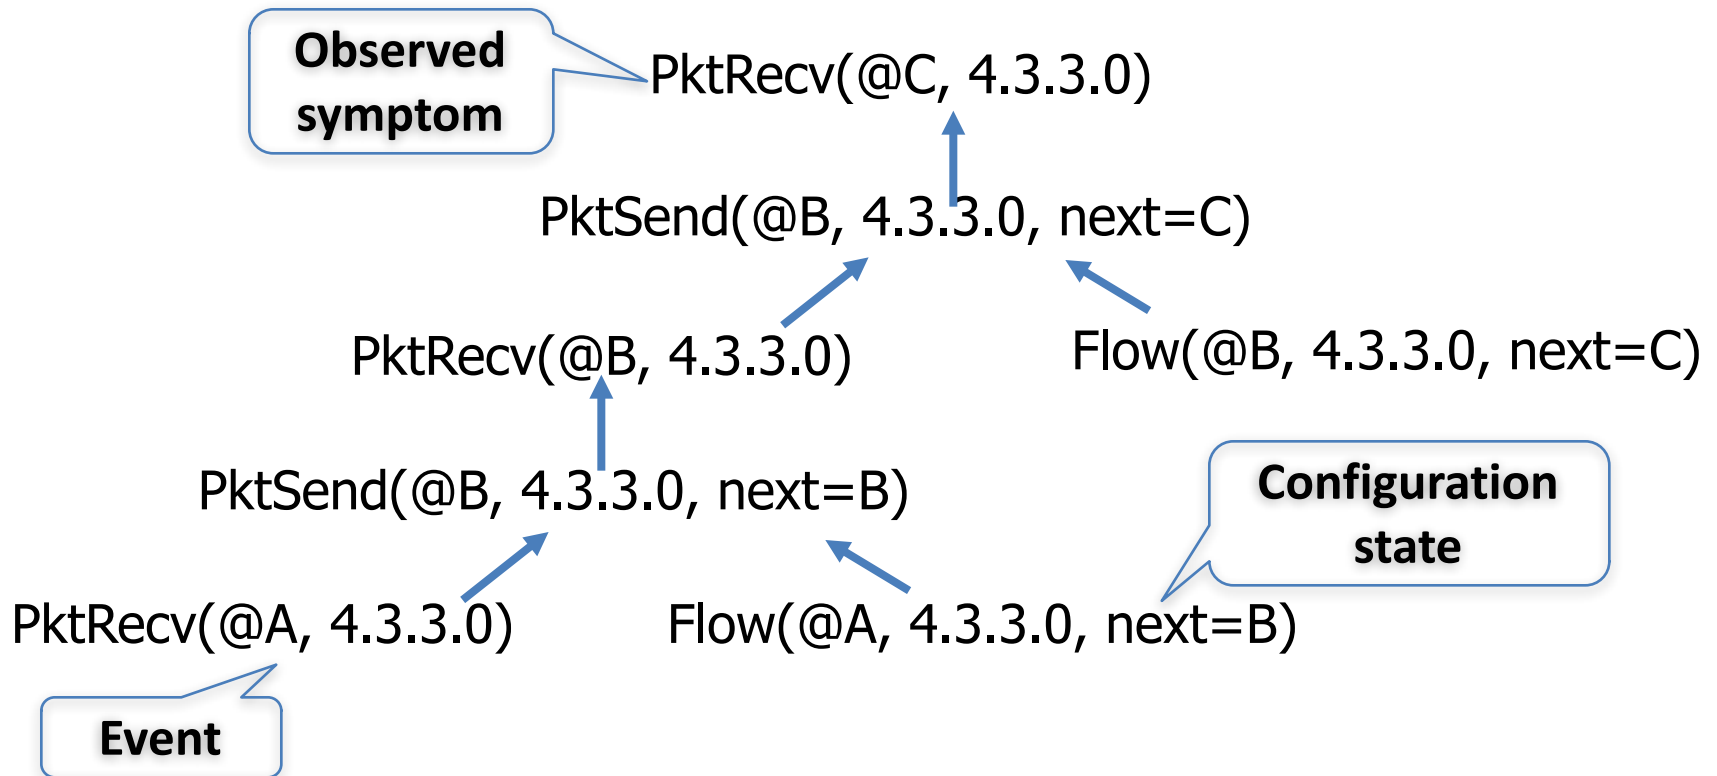


- Input: a bad symptom and a good reference
- Debugger reasons about the differences
- Output: root cause

Outline


- ✓ • Motivation: Root cause analysis
- ➡ • Differential provenance
 - Background: Provenance
 - Strawman solution
 - Algorithm
- Evaluation
 - Prototype implementation
 - Usability
 - Query processing speed
 - Complex network diagnostics
- Conclusion

Background: Provenance



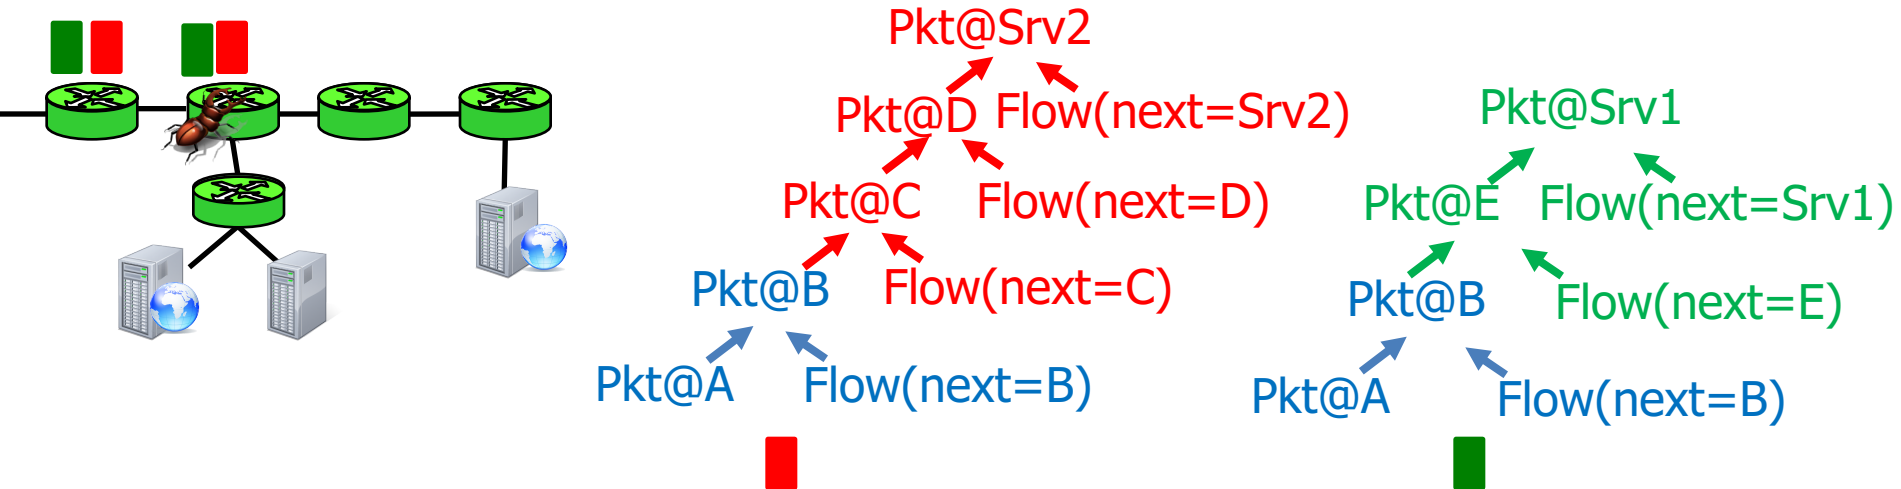
- Provenance tracks causal connections between network events and state [ExSPAN-SIGMOD'10]
 - Provenance graph: Vertexes → event/state. Edge → causality
 - Provenance tree: Recursive explanation of an event/state

Strawman solution



- Strawman solution: Find vertexes that are different in the two trees
- Problem: The **diff can be larger than the individual trees!**

Why does the strawman solution not work?

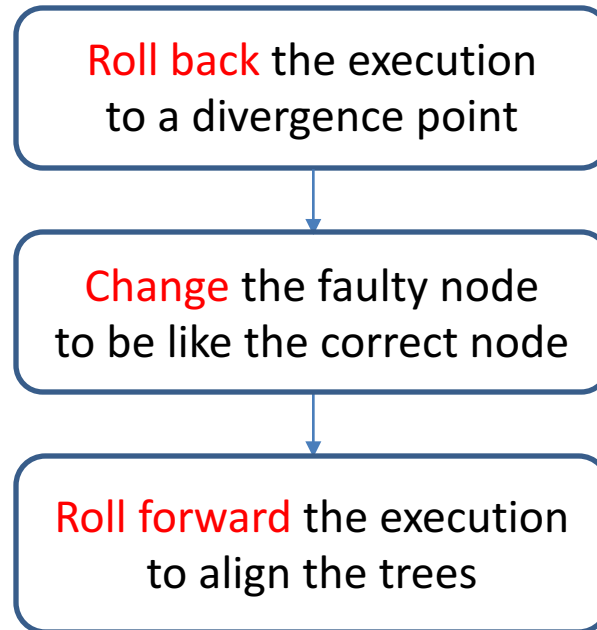


- Observation: The diff can be larger than the individual trees
- Reason: "Butterfly effect"
 - A small initial difference can lead to drastically different events later on

Outline


- ✓ • Motivation: Root cause analysis
- ✓ • Key insight
- ✓ • Differential provenance
 - ✓ • Background: Provenance
 - ✓ • Strawman solution
- ➡ • **Algorithm**
- Evaluation
 - Prototype implementation
 - Usability
 - Query processing speed
 - Complex network diagnostics
- Conclusion

Algorithm: Refinement #1



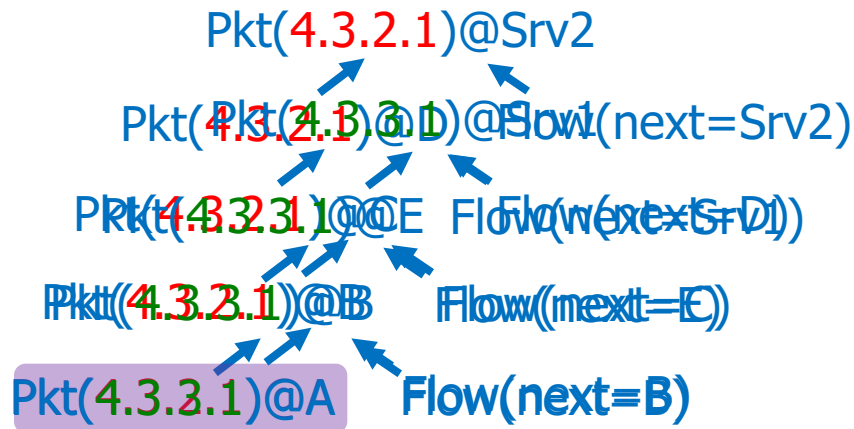
- This approach finds only the (small) initial differences
 - The (potentially large) consequences are ignored

Algorithm: Refinement #1 (Cont'd)

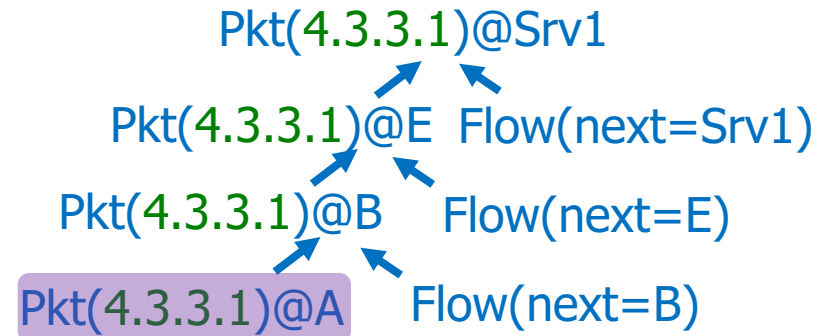


- Approach: Roll back the execution, change the first faulty node, and roll forward again to align the trees

How to preserve crucial differences?



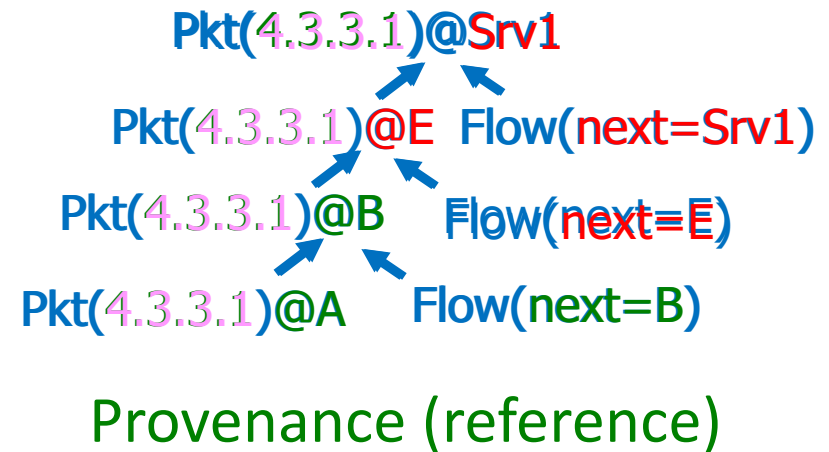
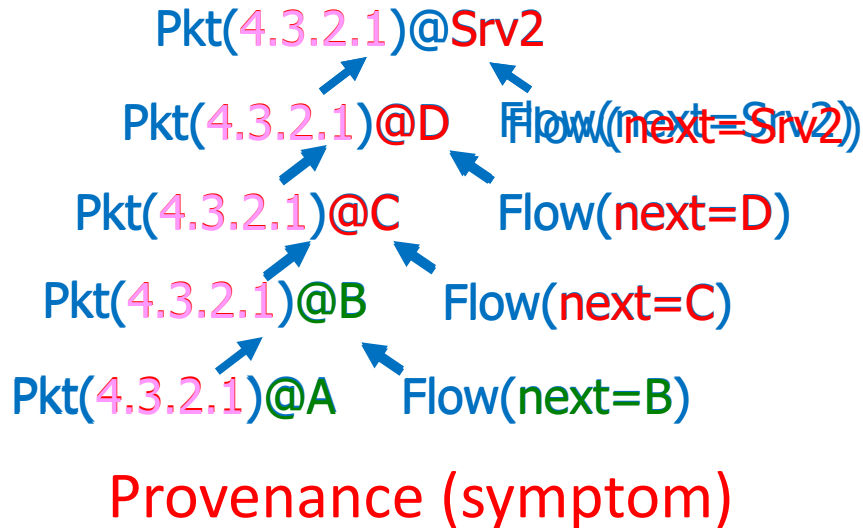
Provenance (symptom)



Provenance (reference)

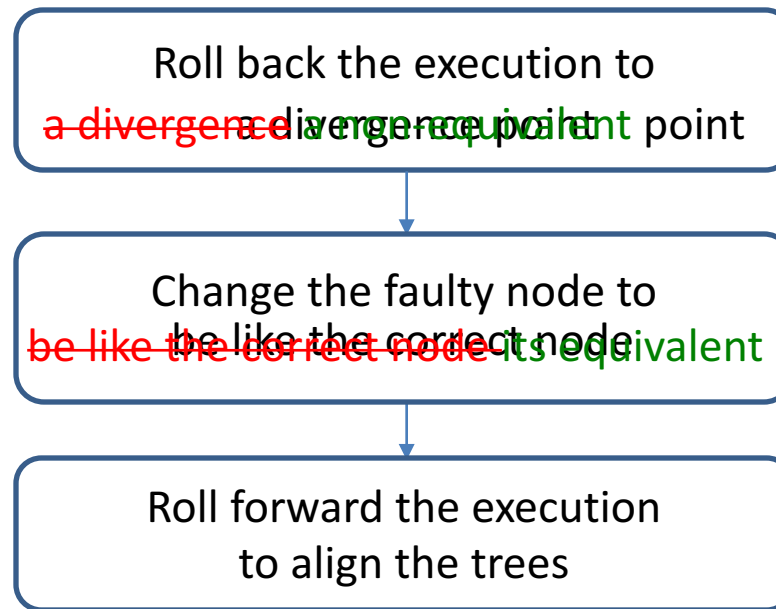
- Problem: There are differences that we need to preserve
 - Example: The packets whose provenance we are looking at

Solution: Establish equivalence



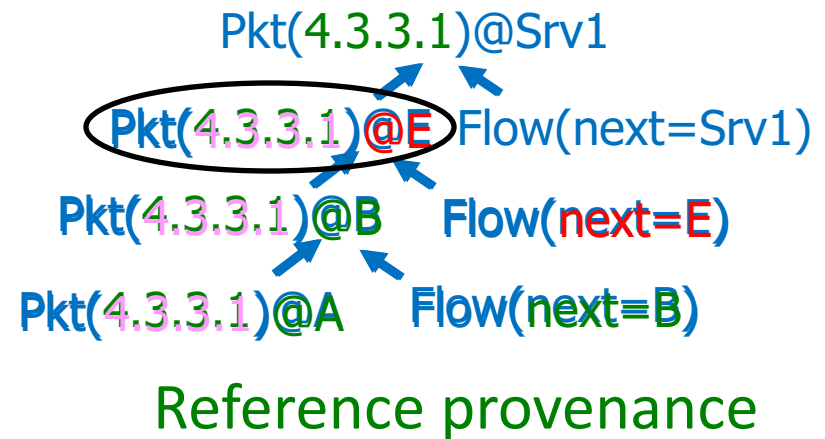
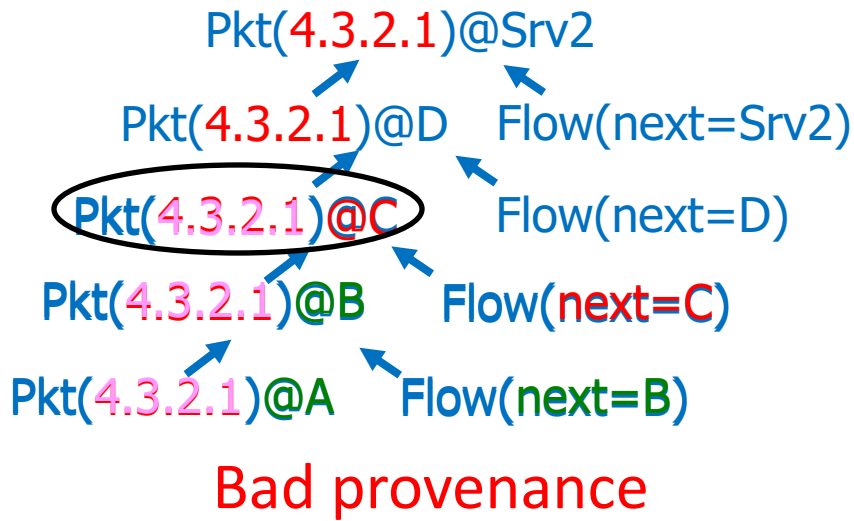
- Establish an equivalence relation between the trees
 - Example: IP addresses 4.3.2.1 and 4.3.3.1
 - Values on the trees can be identical, equivalent, or different
- Goal: Make the trees equivalent, not necessarily identical!

Algorithm: Refinement #2



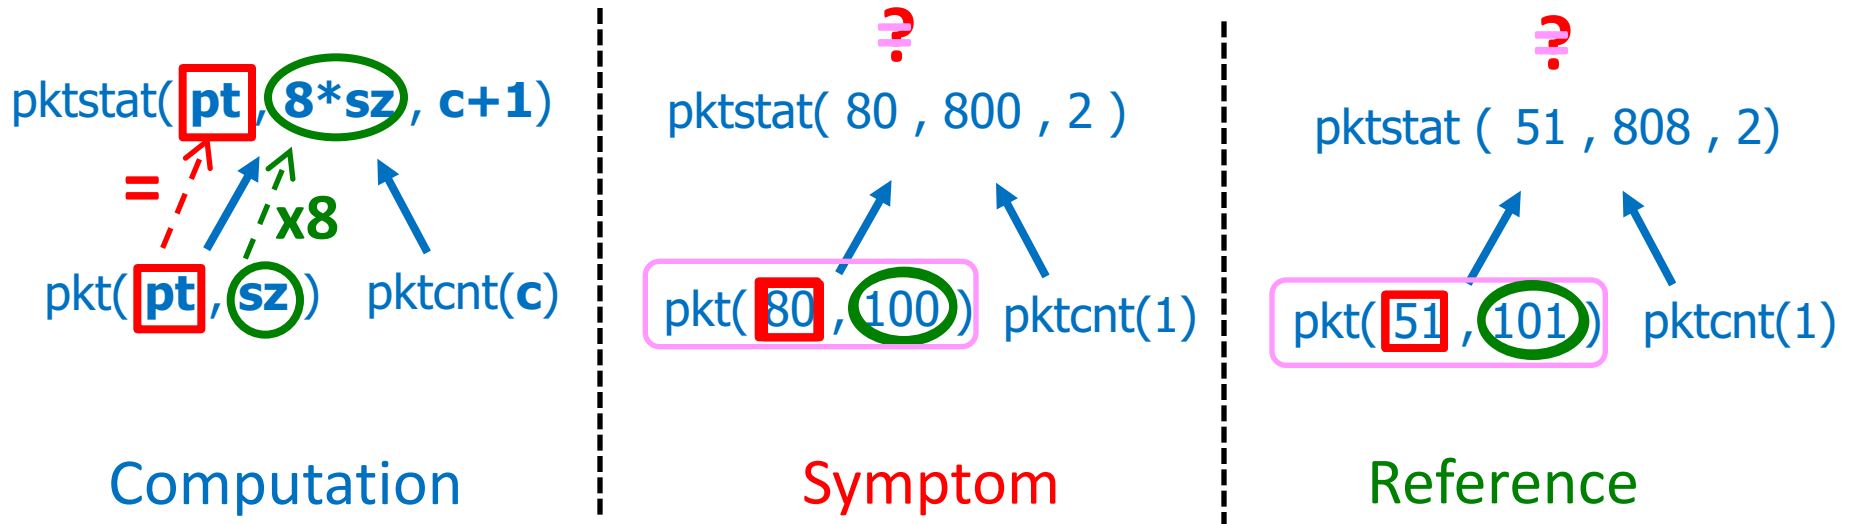
- Benefit: Preserves the crucial differences between the trees

Establishing and propagating equivalence



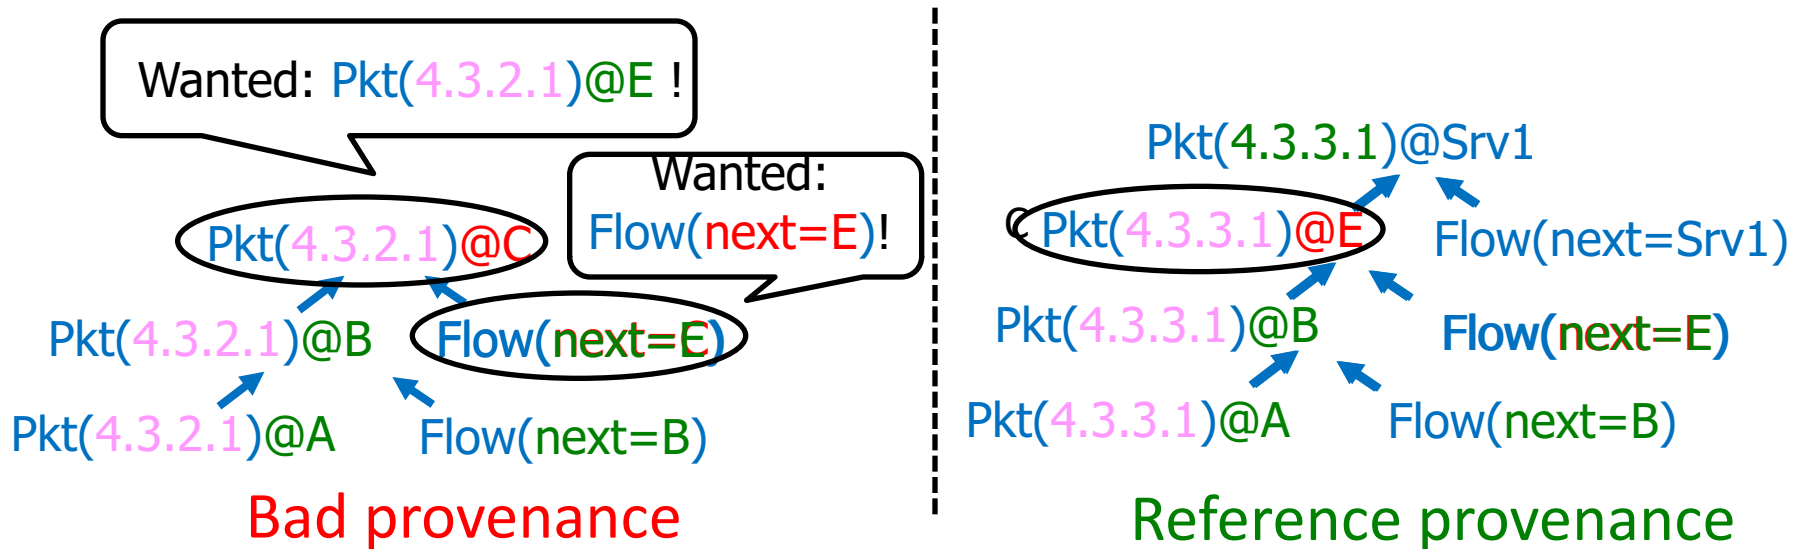
- Start with an initial equivalence relation between the packets
 - Establish a mapping between packet fields that are different
- Keep track of the mapping while going up the tree
 - Stop at the first non-equivalent(!) node
 - More general approach: taint analysis

Propagating equivalence with taints




- Approach:
 - Create **taints** for equivalent fields
 - Propagate taints up the tree
 - Repeat until we find a non-equivalent node

Changing the faulty node




- Change the faulty node to its equivalent: `Pkt(4.3.2.1)@C` → `Pkt(4.3.2.1)@E`
 - Have dependent nodes → Create their equivalents recursively
 - Example: `Flow(next=C)` → `Flow(next=E)`
 - No dependent nodes → Insert its equivalent
 - Example: Insert `Flow(next=E)`
 - See paper for how to propagate taints in the reverse direction

Problem: Multiple faults



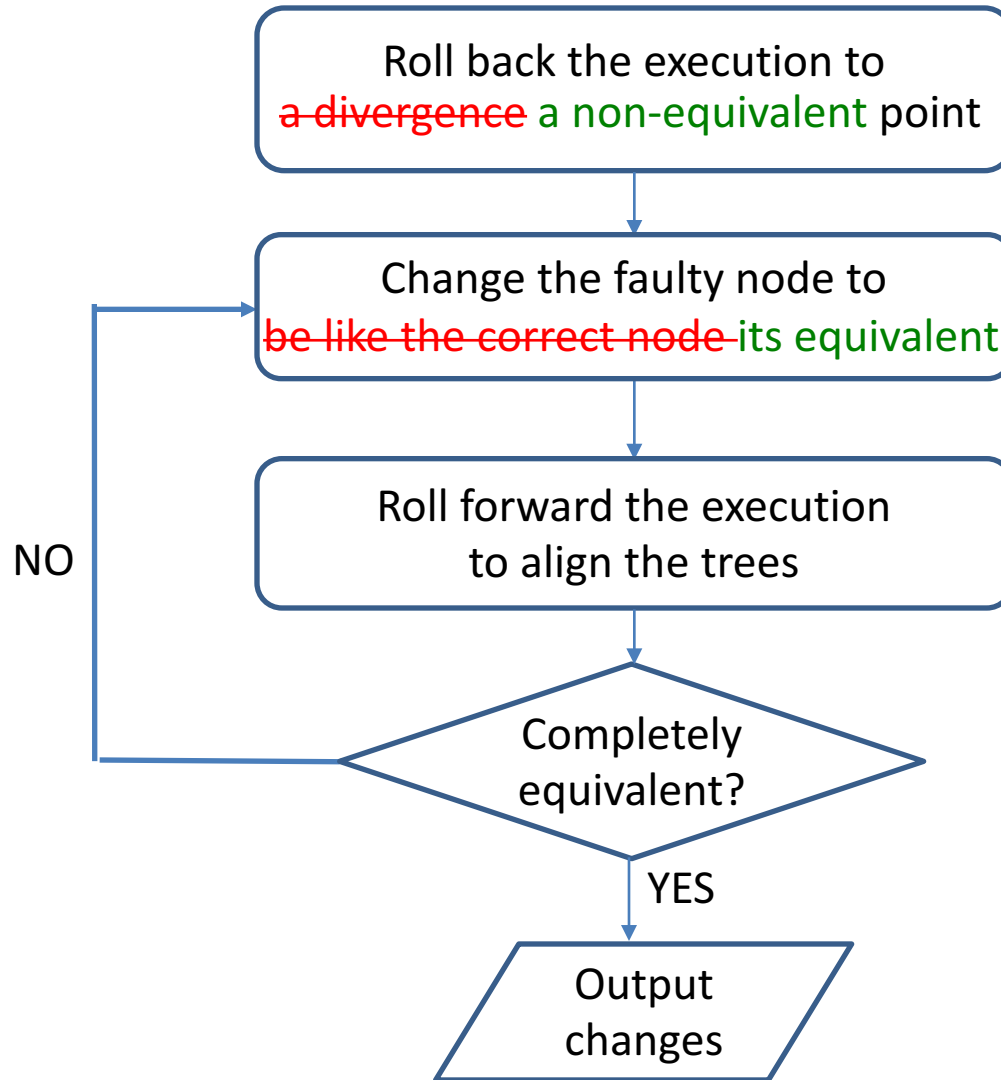
Bad provenance



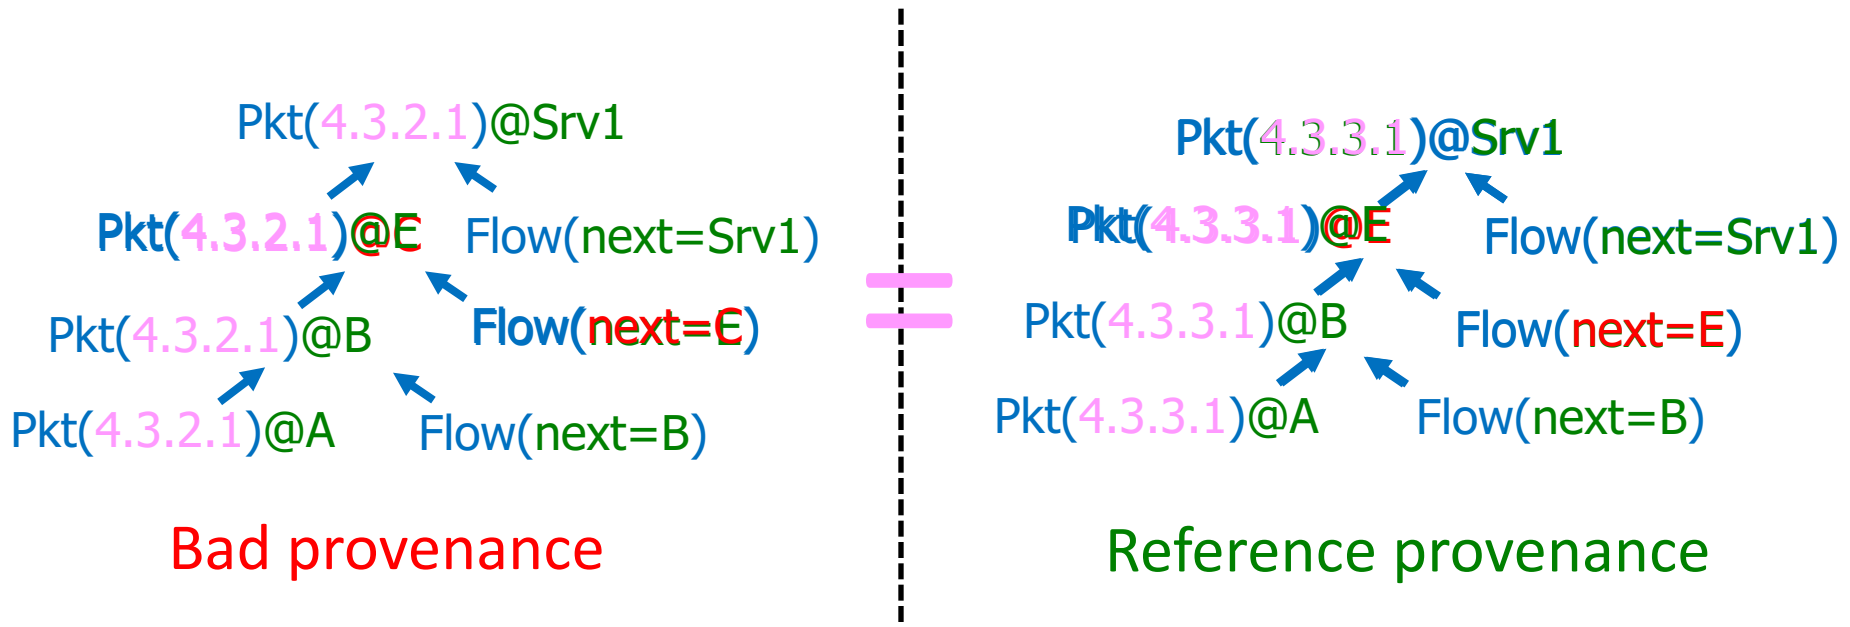
Reference provenance

- Problem: There could be more than one difference between the two trees
- Solution: Repeat until the trees are completely aligned

Refinement #3: Final algorithm



Rolling forward the execution



- Roll the execution forward to align the trees
 - Output the accumulated change(s): Flow(next=C) → Flow(next=E)!

Outline

- ✓ • Motivation: Root cause analysis
- ✓ • Differential provenance
 - ✓ • Background: Provenance
 - ✓ • Strawman approach
 - ✓ • Algorithm
- ➡ • Evaluation
 - Prototype implementation
 - Usability
 - Query processing speed
 - Complex network diagnostics
- Conclusion

Prototype implementation: DiffProv

- Mostly focuses on Network Datalog (NDlog) [CACM '2009] programs, where provenance is easy to see
- NetCore [NSDI '13] programs are also supported
- Applicable beyond SDN: Hadoop MapReduce
- Integrated with Mininet + the Beacon controller; based on Rapidnet

Evaluation: Overview

✓ Q1: How well does DiffProv find the root cause?

Q2: How much overhead does DiffProv incur at runtime?

✓ Q3: How quickly does DiffProv answer diagnostic queries?


Q4: How well does DiffProv recognize bad reference events?

✓ Q5: How well does DiffProv work for complex networks?

Experimental setup


- We adapted seven diagnostic scenarios:
 - SDN1: Broken flow entry [Empr.Soft.Eng.'09]
 - SDN2: Multi-controller inconsistency [CoNEXT'14]
 - SDN3: Unexpected rule expiration [P2P'13]
 - SDN4: Multiple faulty entries [Empr.Soft.Eng.'09]
 - Complex network diagnostics [CoNEXT'12]
 - MR1: Configuration changes [Industry collaborators]
 - MR2: Code changes [Industry collaborators]
- Baseline: Y!, a provenance debugger without reference support [SIGCOMM'14]

How well does DiffProv find the root cause?




Provenance
(symptom)

201 nodes



Provenance
(reference)

156 nodes



Naïve diff

278 nodes



Next hop of
rule 7 is wrong



DiffProv

1 node!




How well does DiffProv find the root cause? (Cont'd)

Query	SDN1	SDN2	SDN3	SDN4
Num. of faults	1	1	1	2
DiffProv	1	1	1	2
Reference	156	156	156	201/201
Symptom	201	156	201	156/145
Plain tree diff	278	238	74	278/218

Query	MR1-D	MR2-D	MR1-I	MR2-I
Num. of faults	1	1	1	1
DiffProv	1	1	1	1
Reference	1051	1001	588	588
Symptom	1051	848	588	438
Plain tree diff	164	306	240	216

- DiffProv finds **one or two nodes (the faulty rules or MapReduce configuration entries)**, which are the actual root cause

How long does DiffProv take to find the root causes?



- DiffProv answered most of our queries **within one minute!**

How well does DiffProv work in complex networks?

- Setup: larger topology, complex config, background traffic
 - 'Forwarding error' scenario [ATPG-CoNEXT'12]
 - Stanford network: 757,000 forwarding entries and 1,500 ACL rules
 - Multiple faults: Injected 20 additional faulty entries
 - Background traffic: 12GB traffic, 69 protocol types
- Results:
 - DiffProv: the faulty entry for misconfigured subnet – one node
 - Identified the root cause despite heavy interference
- Why is DiffProv not confused by the interference?
 - Provenance captures causality, not merely correlations!

Summary

- Debugging networks is hard
 - Need good debuggers to **find root causes!**
- Key insight: **We can use reference events**
 - We often have more information than we are using
 - Idea: Reason about the differences between bad events and reference events
- **Approach: Differential provenance**
 - We have built a prototype debugger for SDNs
 - Applicable to other distributed systems beyond SDNs
- Result: **Very precise diagnostics**
 - Differential provenance can often identify a single root cause

Thank you!