



Globally Synchronized time via Datacenter Networks


Ki Suh Lee

Cornell University

Joint work with Han Wang, Vishal Shrivastav and Hakim Weatherspoon


Synchronized Clocks

- Fundamental for network and distributed systems
 - OWD, Monitoring, Coordination, Snapshots, Updates, ...
- Goal: *Minimized and bounded precision with scalability*
 - Minimized and bounded precision: *hundreds of nanoseconds*
 - Scalability: *Entire datacenter*



Clock Synchronization Protocol

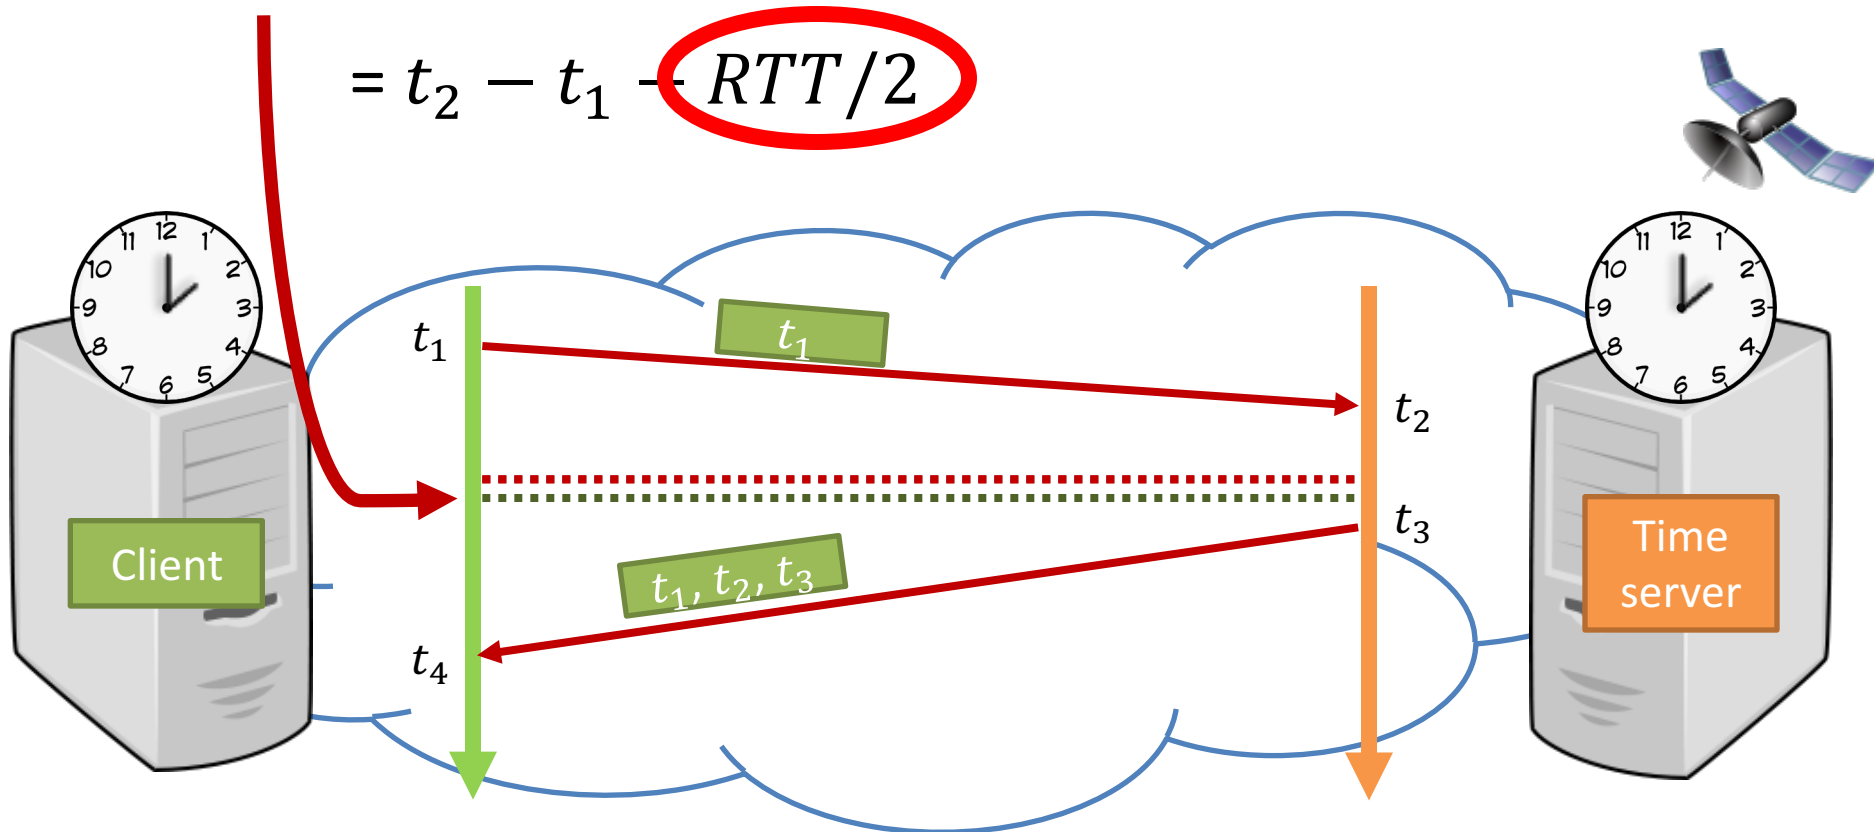
- Offset: Time difference between two clocks
- Precision: The worst case of offset



Clock Synchronization Protocol

- $RTT = (t_4 - t_1) - (t_3 - t_2)$
- Offset =

$$= t_2 - t_1 - RTT/2$$






Current time protocols
do NOT provide *bounded precision*,
due to *uncertainty* in measured RTT!


Challenge: RTT is not accurate

- Errors from
 - Oscillator skew
 - Inaccurate Timestamping
 - Network Stack
 - Network Jitter




Challenge: RTT is not accurate

- Errors from
 - Oscillator skew
 - Inaccurate Timestamping
 - Network Stack
 - Network Jitter
- PTP
 - Hardware timestamping
 - PTP-enabled switches
 - Filtering / Smoothing




Challenge: Scalability

- Re-synchronization period vs. Network overhead
- Limited number of clients



Synchronization Protocols

	Precision	Scalability	Overhead	Extra Hardware
NTP	us	Good	Moderate	None
PTP	sub-us	Good	Moderate	PTP-enabled devices
GPS	ns	Bad	None	Timing signal receivers, cables



Solution: Use *the PHY* to synchronize clocks

Application



Transport

Network

Data Link

Physical

- Protocol in the PHY
 - Each physically link is already synchronized!
 - No protocol stack overhead
 - No network overhead
 - Scalable: peer-to-peer and decentralized



DTP: Datacenter Time Protocol

Application



Transport

Network

Data Link

Physical

- ***Highly Scalable with bounded precision!***
 - ~25ns (4 clock ticks) between peers
 - ~150ns for a datacenter with six hops
 - No Network Traffic
 - *Internal* Clock Synchronization
- End-to-End: ~200ns precision!





Outline

- Introduction
- **Design**
- Evaluation
- Discussion
- Conclusion

DTP: Datacenter Time Protocol

Application


Transport

Network

Data Link

Physical

- 10G Background
 - Continuous /I/s when there is no packet
 - At least 12 /I/s between two Ethernet frames



DTP: Datacenter Time Protocol

Application



Transport

Network

Data Link

Physical

- 10G Background
 - Continuous /I/s when there is no packet
 - At least 12 /I/s between two Ethernet frames
 - 1 Control block (/E/, 66bit) = 8 /I/s
 - At least 1 /E/ between any two frames
 - The PHY is run by 156.25MHz
 - Period is 6.4ns



DTP: Datacenter Time Protocol

Application

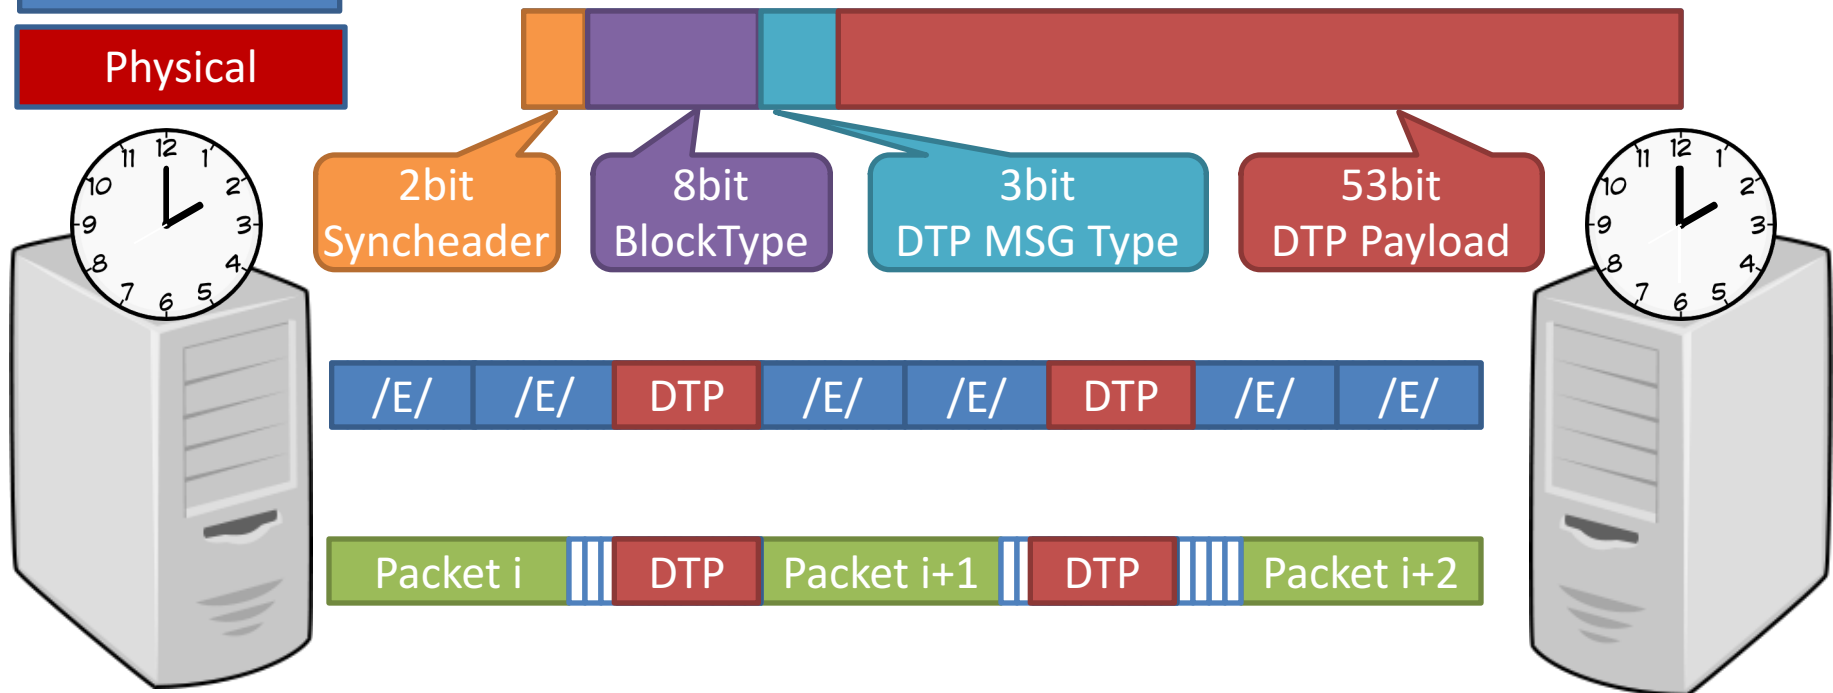
Transport

Network

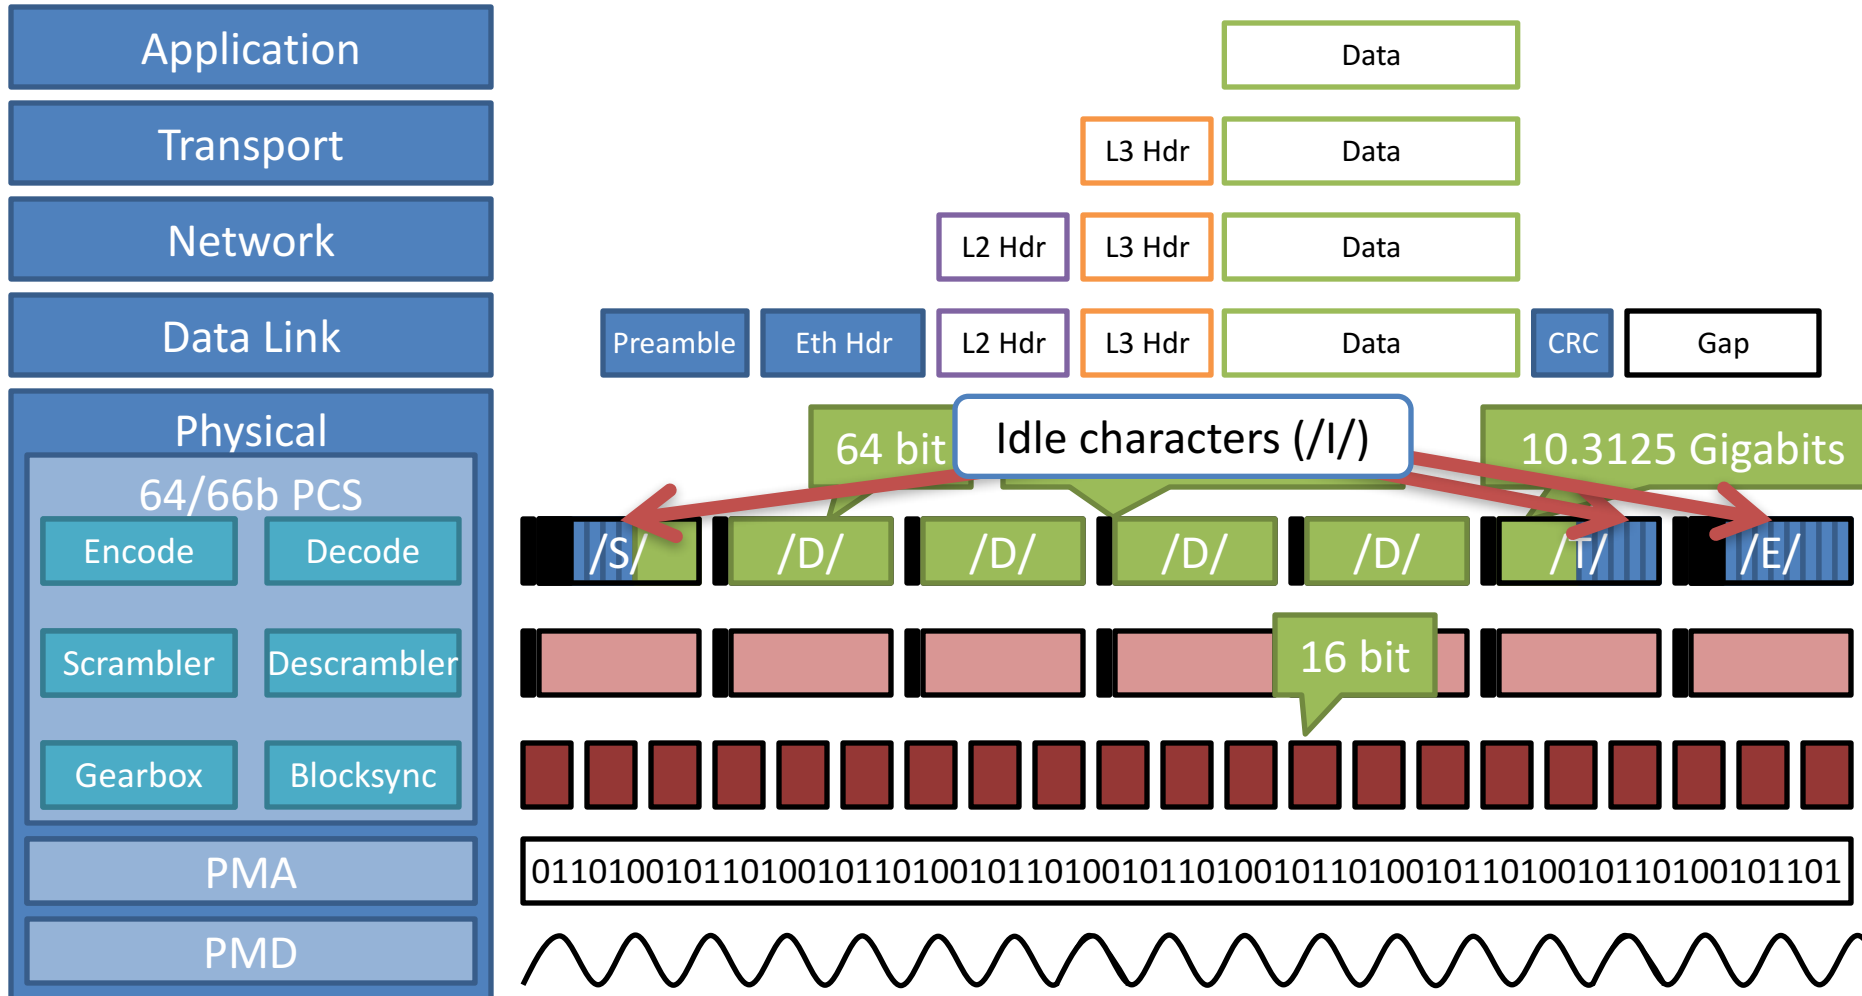
Data Link

Physical

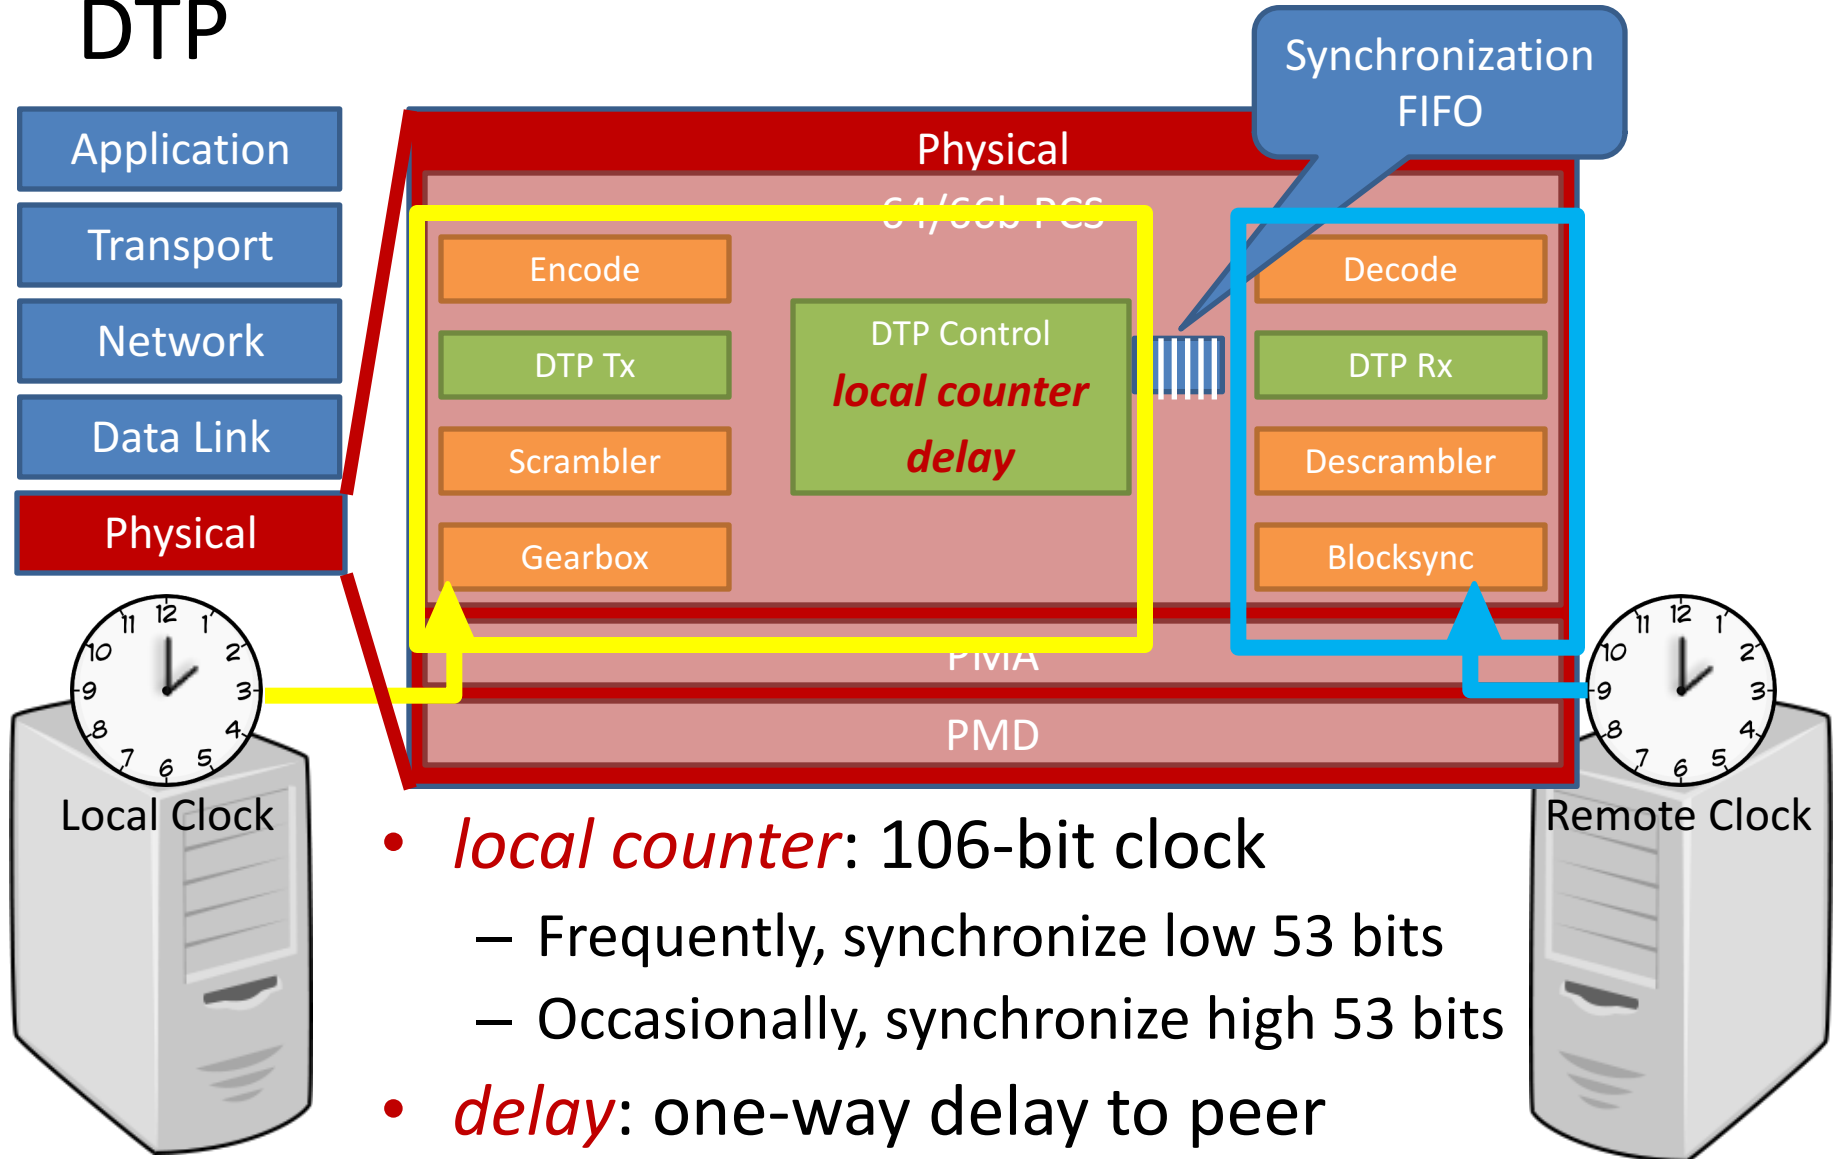
- DTP overwrites /E/ to send protocol messages
 - Frequent messaging
 - No overhead to Ethernet (L2)



10GbE Network Stack



DTP



DTP

Application

Transport

Network

Data Link


Physical

- Runs in two phases between two peers
 - Init Phase: Measuring OWD
 - Beacon Phase: Re-Synchronization



Physical

*local
delay*



Physical

*local
delay*

DTP: Init Phase

Application




Transport

Network

Data Link

Physical

- $delay = (t_4 - t_1 - \alpha)/2$
 - $\alpha=3$: Ensure *delay* is always less than actual delay
- Introduce 2 clock tick errors
 - Due to oscillator skew, timing and Sync FIFO



DTP: Beacon Phase

Application

Transport

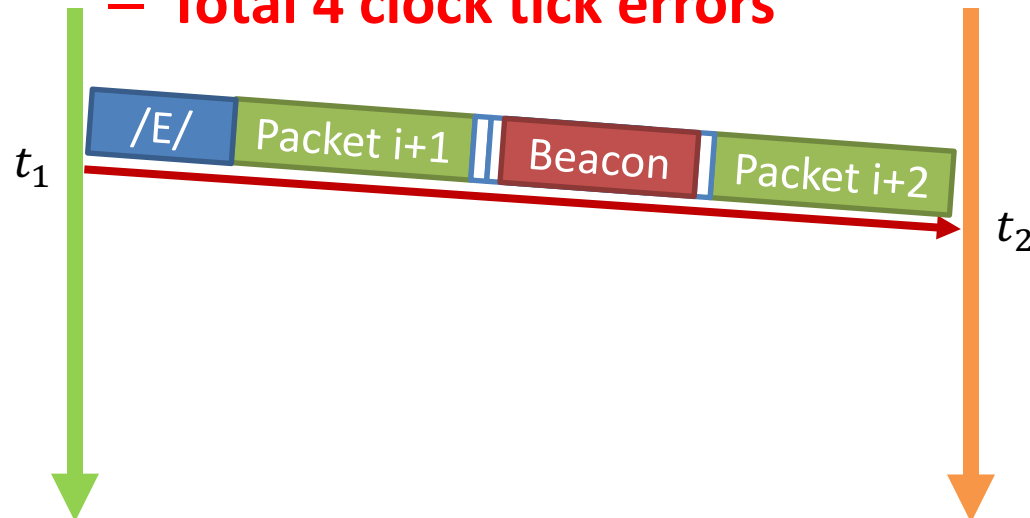
Network

Data Link

Physical

- $local = \max(local, remote + delay)$
- Frequent messages
 - Every 1.2 us (200 clock ticks) with MTU packets
 - Every 7.2 us (1200 clock ticks) with Jumbo packets
- Introduces 2 clock tick errors

– **Total 4 clock tick errors**



Physical

local
delay

Physical

local
delay

DTP Switch

Application

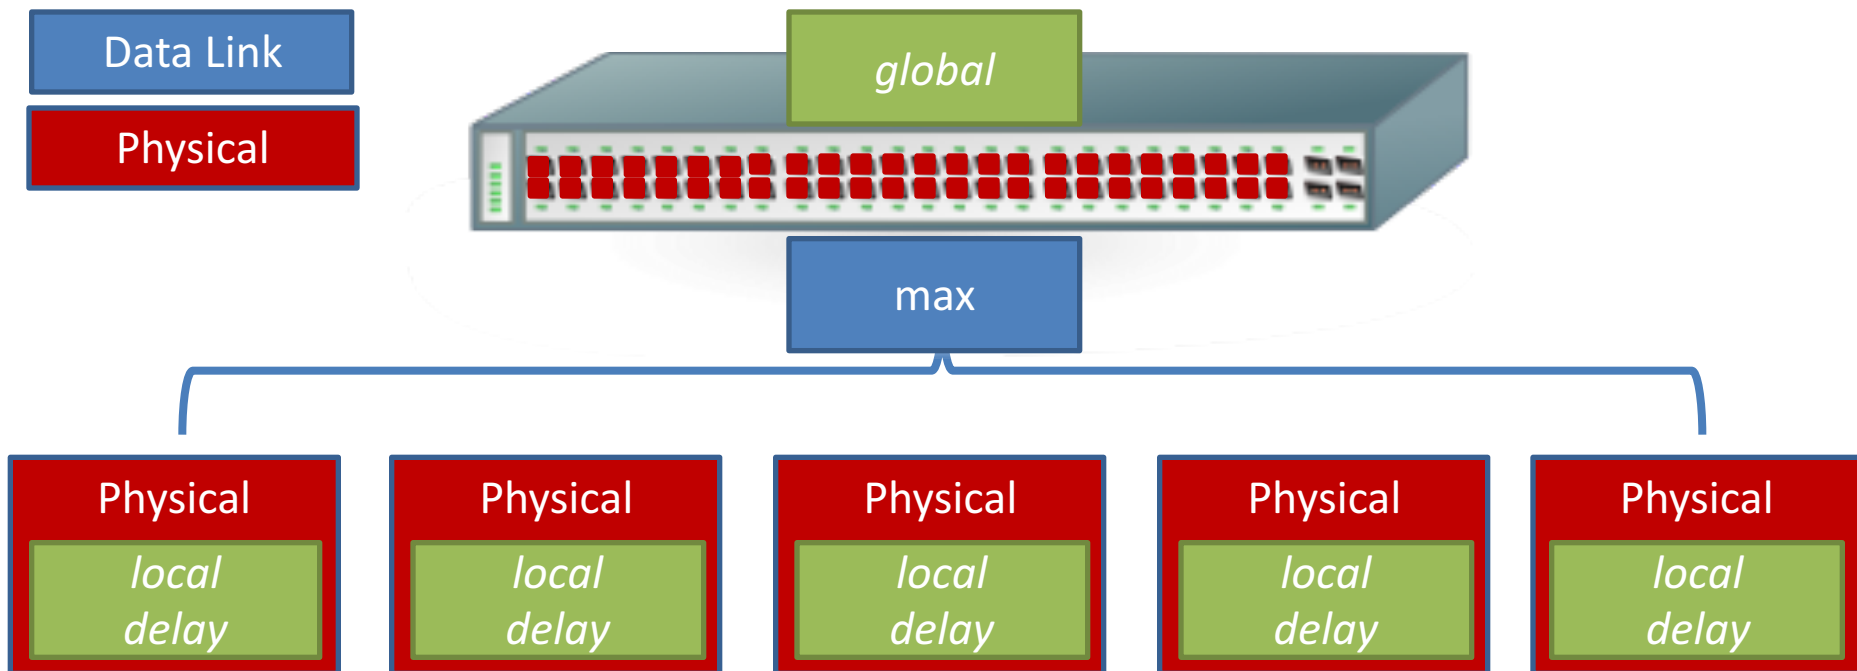
Transport

Network

Data Link

Physical

- $global = \max(local \text{ counters})$
- Propagates *global* via Beacon messages





DTP Daemon

- End-to-End precision
- Access the DTP counter via PCIe
- Estimate DTP time using invariant TSC counter



DTP Property

- Bounded Precision in hardware
 - Bounded by $4T$ ($=25.6\text{ns}$, $T=\text{oscillator tick is } 6.4\text{ns}$)
 - Network precision bounded by $4TD$
 - D is network diameter in hops
- Requires NIC and switch modifications
 - PTP also requires PTP-enabled devices



DTP vs PTP

	PTP	DTP
Oscillator Skew		
Timestamping	HW - timestamping	PHY timestamping
Network Stack	Not involved	Not involved
Network Jitter	Transparent Clock Boundary Clock	No jitter
Precision	Unbounded Tens to Hundreds ns (When Idle)	Bounded



DTP: Topics discussed in paper

- Handling failure
- Different standards: 1GbE, 25GbE, 40GbE, 100GbE, etc
- External synchronization (i.e. synchronizing to true time)
- Incremental deployment



Handling failure

- Bit Errors
 - Ignores Bit errors in MSBs
 - Appends checksum for low LSBs
- Faulty Devices
 - When too many jumps outside the bound



Different Standards

Data Rate	Encoding	Data Width	Frequency	Period	Δ
1 GbE	8b/10b	8bit	125MHz	8ns	25
10 GbE	64b/66b	32bit	156.25MHz	6.4ns	20
40 GbE	64b/66b	64bit	625MHz	1.6ns	5
100 GbE	64b/66b	64bit	1562.5MHz	0.64ns	2



External Synchronization

- *A master server*
 - Connected to a reference time
 - Broadcasts the mapping between *DTP* and wall time
- Client servers
 - Interpolates time using *DTP* counters



Incremental Deployment

- Updates per rack
 - DTP-enabled switch
 - DTP-enabled NICs
 - One server acting as a *master* for wall time
- Synchronizing Racks
 - DTP-enabled switch
 - DTP *beacon-join* message for synchronizing DTP counters
 - Select a new *master*




Outline


- Introduction
- Design
- **Evaluation**
- Discussion
- Conclusion

Evaluation

- DTP Prototype
 - Terasic DE5 board with Altera Stratix V
 - Using Bluespec and Connectal framework



Evaluation: DTP Topology




Measured offsets between peers

Evaluation: Logger

- Offset between peers: $t_3 - t_2 - \text{OWD}$
- Offset between SW and HW: $t_2 - t_1$

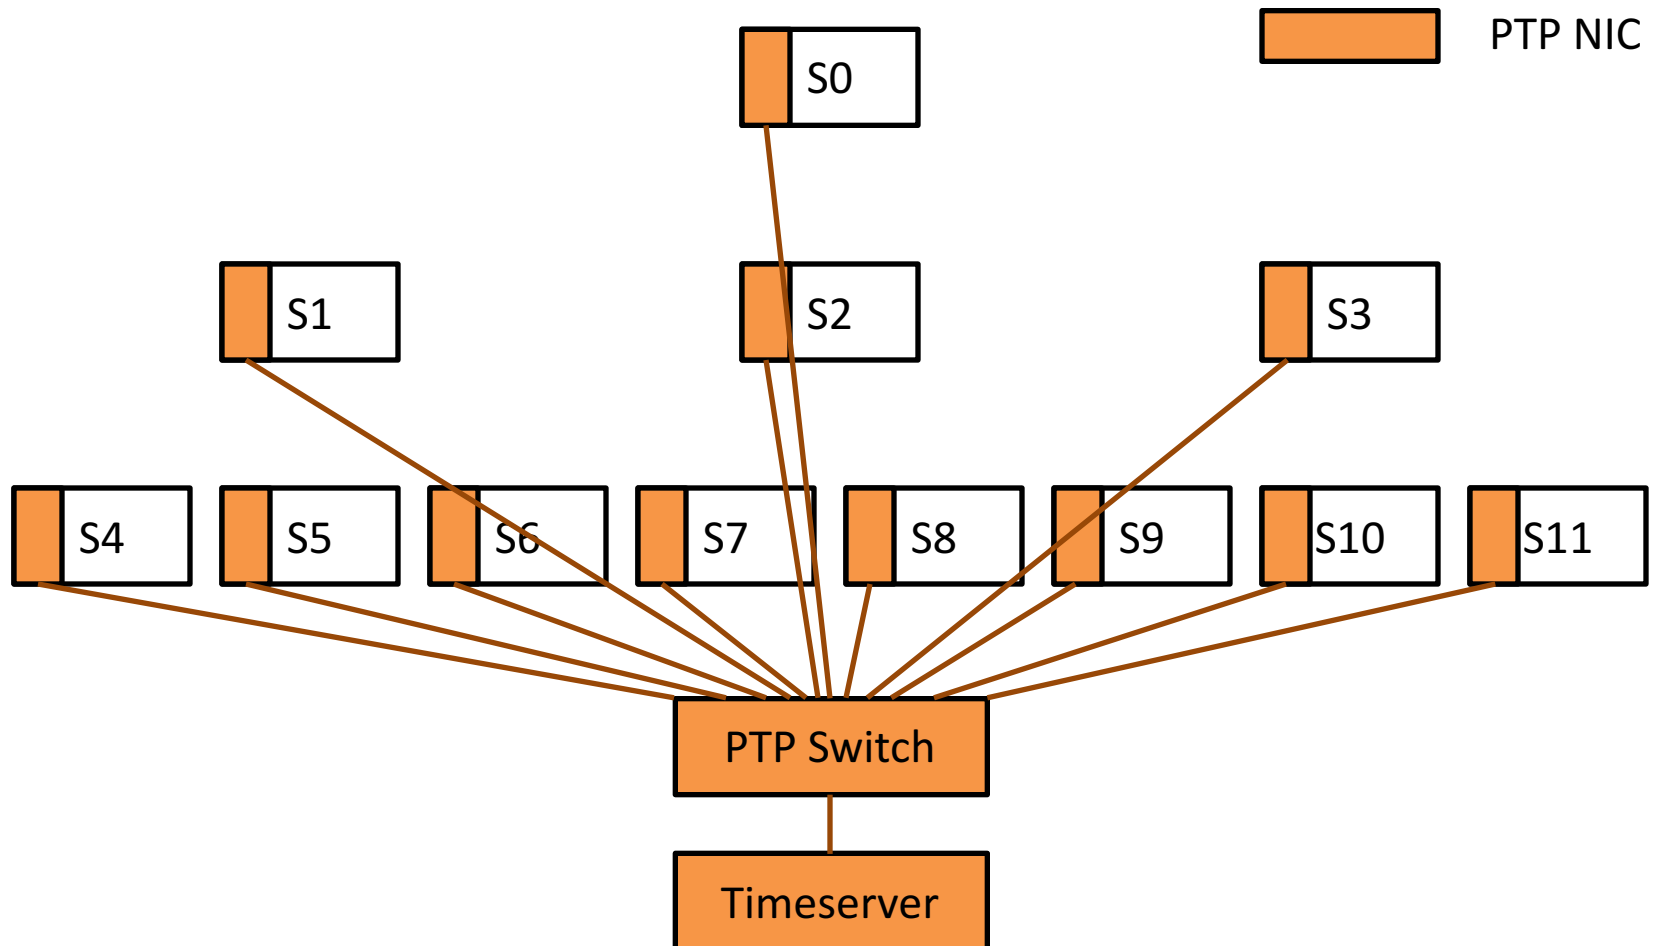


Evaluation: DTP Topology

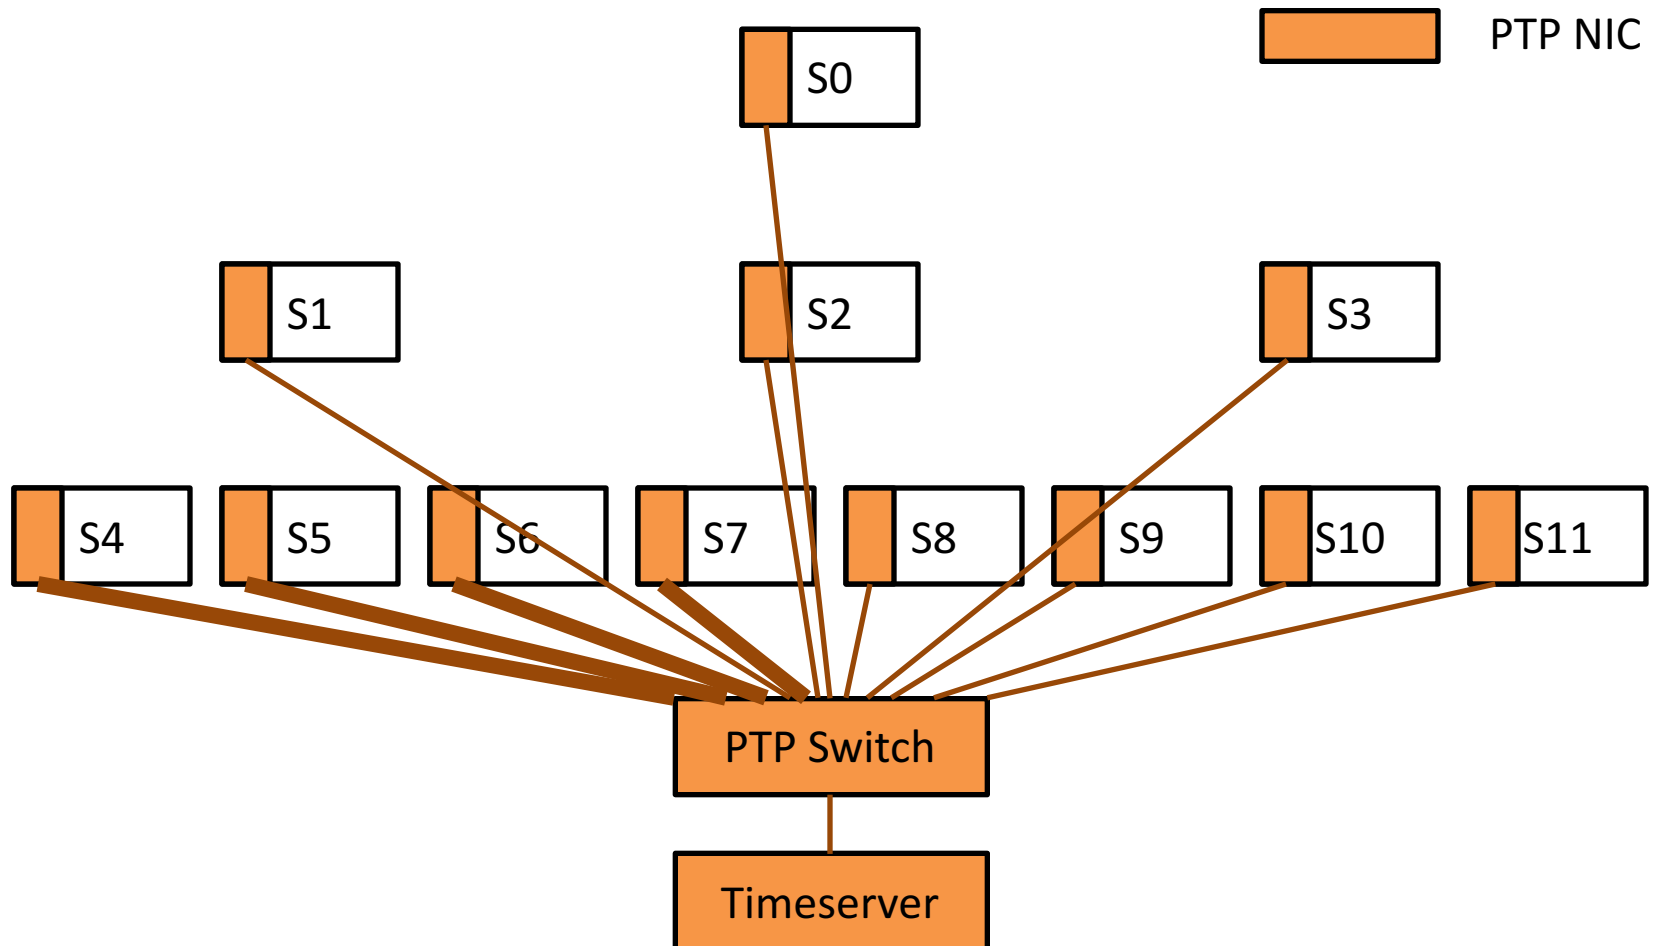


$$Offset = dtp_{SW} - dtp_{HW}$$

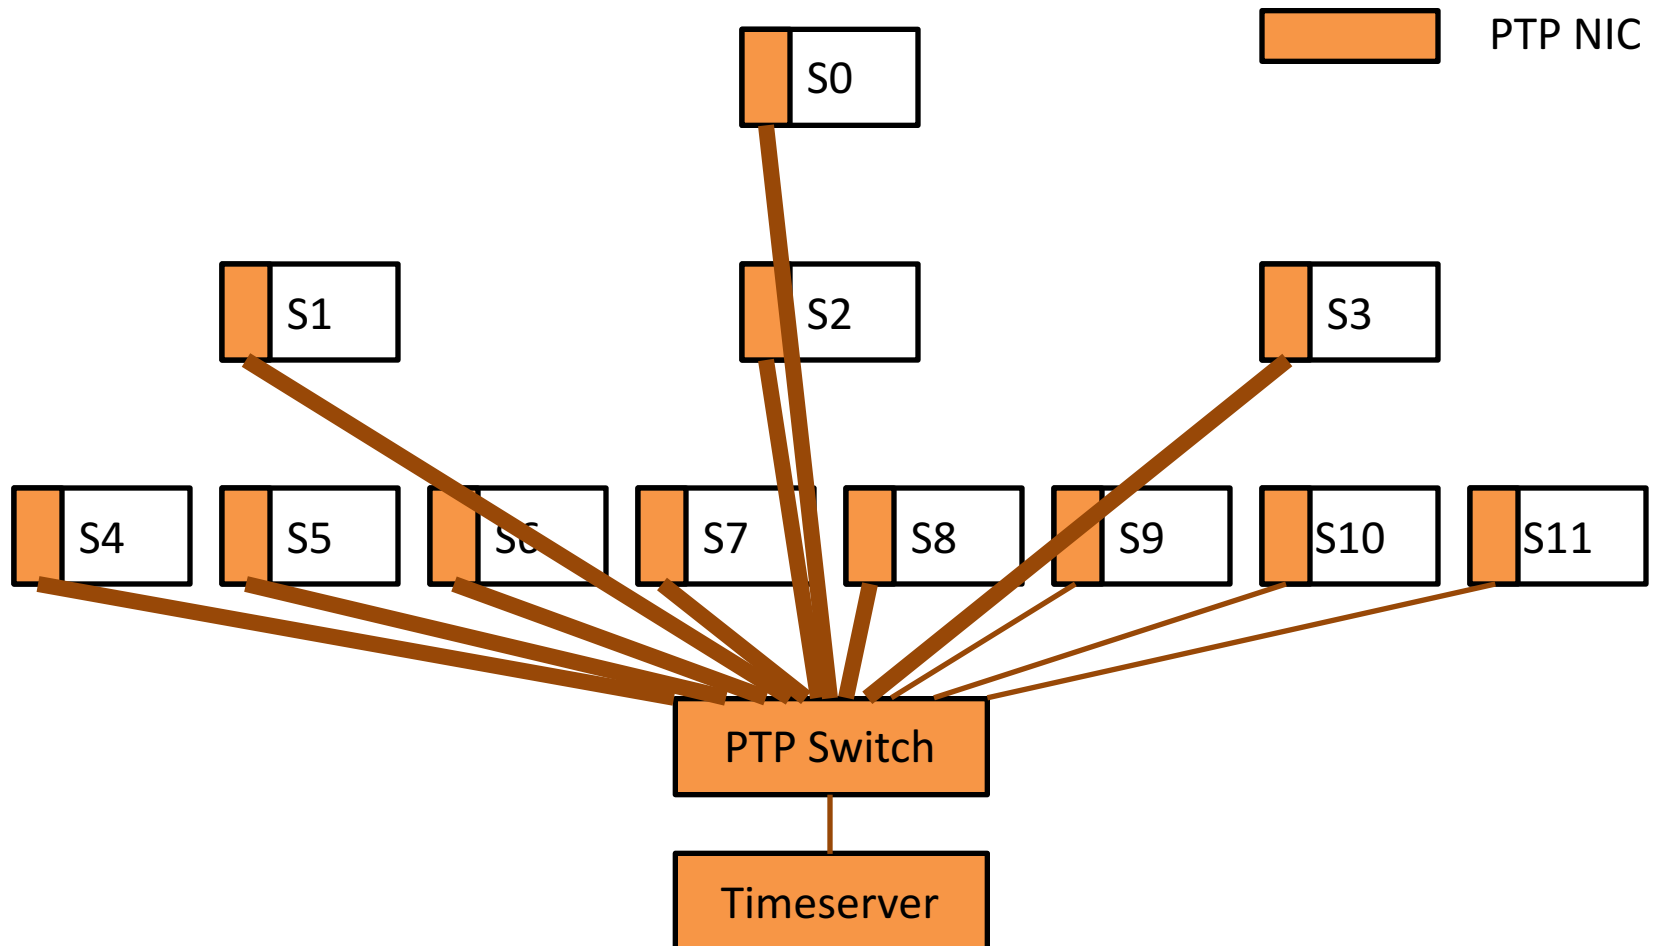
Evaluation: PTP Topology



Evaluation: PTP Topology




Evaluation: PTP Topology




PTP: Idle Network (No traffic)

- Tens to hundreds of nanosecond precision*




PTP: Medium Loaded (4 Gbps)

- *Tens of microseconds precision*




PTP: Heavily Loaded (9 Gbps)

- Tens to hundreds of microsecond precision*



DTP: Heavily Loaded

- Always within 25.6ns (4 clock ticks) between peers*






DTP Daemon

DTP Daemon (after smoothing)

- Usually can access the counter with 25.6 ns precision*






Outline

- Introduction
- Design
- Evaluation
- Discussion
- Conclusion

Next Steps

- Integration with OSNT (Open Source Network Tester)
 - NetFPGA SUME Board with Xilinx Virtex-7





Some Related Work

- Synchronous Ethernet (SyncE)
 - Synchronize the *frequency* of clocks
 - DTP, PTP synchronizes the *time* of clocks
- WhiteRabbit: PTP + SyncE
 - Sub-nanosecond precision
 - 1GbE only yet
- Commercial PTP + SyncE
 - Tens to hundreds of nanoseconds

Conclusion

- DTP provides *bounded precision* and *scalability*
 - Bounded precision: 4 clock ticks (25.6ns) between peers
 - Scalability: 153.6ns for a datacenter with six hops
 - Free: No Network Traffic
 - Applications: Usually within 25.6ns (without bounds)
 - End-to-End: $153.6 + 25.6 * 2 = 200\text{ns!}$



Questions?

- <http://github.com/hanw/sonic-lite>
- <http://sonic.cs.cornell.edu>
- Email: kslee@cs.cornell.edu
- Come to Poster session tomorrow!