



trilogy 2




Symnet: scalable symbolic execution for modern networks


University Politehnica of Bucharest

*Radu Stoenescu, Matei Popovici, Lorina Negreanu and
Costin Raiciu*

Networks are increasingly complex



Understand the network




Reachability

Packet modifications

Security policy violations

Static verification to the rescue



Data plane snapshot

Symbolic Execution Friendly Language (SEFL) - Network model

← Symnet Verification engine →



Choosing a modeling language

C code

- Expressive, well understood
- Symbolic execution captures many properties
- Very expensive to verify

Middle ground

Header Space Analysis

- Cheap, scalable
- No arbitrary protocol layering
- Only captures reachability

Symbol execution of firewall - C code

```
1: packet* filter(packet* p){  
2:     if (p->dst_port==80)  
4:         return p;  
5:     else {  
6:         free p;  
7:         return NULL;  
8:     }  
9: }
```

Symbol execution of firewall - C code

Path 1

*p=**

```
1: packet* filter(packet* p){  
2:     if (p->dst_port==80)  
4:         return p;  
5:     else {  
6:         free p;  
7:         return NULL;  
8:     }  
9: }
```

Symbol execution of firewall - C code

Path 1

$p = *$

```
1: packet* filter(packet* p){  
2:     if (p->dst_port==80)  
4:         return p;  
5:     else {  
6:         free p;  
7:         return NULL;  
8:     }  
9: }
```


Symbol execution of firewall - C code



```
1: packet* filter(packet* p){  
2:     if (p->dst_port==80)  
4:         return p;  
5:     else {  
6:         free p;  
7:         return NULL;  
8:     }  
9: }
```

p->dst_port=80 p->dst_port!=80

Symbol execution of firewall - C code

Path 1

Path 2

```
1: packet* filter(packet* p){  
2:     if (p->dst_port==80)  
4:         return p;  
5:     else {  
6:         free p;  
7:         return NULL;  
8:     }  
9: }
```

p->dst_port=80 p->dst_port!=80

Symbol execution of firewall - C code

Path 1

Path 2

```
1: packet* filter(packet* p){  
2:     if (p->dst_port==80)  
4:         return p;  
5:     else {  
6:         free p;  
7:         return NULL;  
8:     }  
9: }
```

p->dst_port!=80
p->dst_port=80
filter = p

Symbol execution of firewall - C code

Path 1

Path 2

```
1: packet* filter(packet* p){  
2:     if (p->dst_port==80)  
4:         return p;  
5:     else {  
6:         free p;  
7:         return NULL;  
8:     }  
9: }
```

p->dst_port!=80
p->dst_port=80
filter = p

Symbol execution of firewall - C code

Path 1

Path 2

```
1: packet* filter(packet* p){  
2:     if (p->dst_port==80)  
4:         return p;  
5:     else {  
6:         free p;  
7:         return NULL;  
8:     }  
9: }
```

p->dst_port!=80
p->dst_port=80
filter = p

Symbol execution of firewall - C code

Path 1

Path 2

```
1: packet* filter(packet* p){  
2:     if (p->dst_port==80)  
4:         return p;  
5:     else {  
6:         free p;  
7:         return NULL;  
8:     }  
9: }
```

*p->dst_port=80
filter = p*

p=NULL

Symbol execution of firewall - C code

Path 1

Path 2

```
1: packet* filter(packet* p){  
2:     if (p->dst_port==80)  
4:         return p;  
5:     else {  
6:         free p;  
7:         return NULL;  
8:     }  
9: }
```

p->dst_port=80
filter = p

p=NULL

Symbol execution of firewall - C code

Path 1

Path 2


```
1: packet* filter(packet* p){  
2:     if (p->dst_port==80)  
4:         return p;  
5:     else {  
6:         free p;  
7:         return NULL;  
8:     }  
9: }
```

p->dst_port=80
filter = p

p=NULL
filter=NULL


Symbol execution of firewall - C code

```
1: packet* filter(packet* p){  
2:     if (p->dst_port==80)  
4:         return p;  
5:     else {  
6:         free p;  
7:         return NULL;  
8:     }  
9: }
```



Two symbolic paths vs. one viable in the network
Non-packet processing being executed

Symbol execution of firewalls - C code



Symbolic execution of network data plane implementations does not scale

- A core IP router results in hundreds of thousands of paths
- For a TCP options-parsing middlebox, runtime depends on option length (<40):
 - 6B \sim 1 hour, 7B \sim 3 hours

Principles for scalable data plane symbolic execution

Fundamental tradeoff between fast symbolic execution and runtime efficiency [Wagner'13]

=> Use models of networks instead of real code

Only analyze relevant code

=> 1 execution path == 1 network packet

Complex data structures kill symbolic execution

=> Use symbolic-execution friendly data structures

Loops + conditionals are dangerous

=> Careful looping semantics with low branching factor

Our solution

SEFL symbolic execution friendly language

Symnet symbolic execution tool

Memory safety by design

- The memory space is the packet
- No pointers
- Memory access via concrete offsets; validated

Symbolic execution constructs part of the language

- Explicit forking of new execution paths
- Explicit stating of path constraints


No arbitrary data structures

- Only a map data structure

SEFL symbolic execution friendly language

- Variables are packet headers or metadata
 - **Packet headers** allocated at specific addresses in the packet header
 - **Metadata** are key/value pairs in a map data structure

The packet header in SEFL



```
CreateTag("L3",0)
```

```
Allocate(IpSrc,32)+96,IpDst) = Tag("L3")+96
```

```
Assign(IpSrc,"192.168.1.1")
```

```
Allocate(IpDst,32)
```

```
Assign(IpDst,Symbolic)
```

```
CreateTag("L2",Tag("L3")-112)
```

```
Assign(DstMac,Symbolic)
```

ERROR

Firewall

C

```
1: packet* filter(packet* p){  
2:   if (p->dst_port==80)  
4:     return p;  
5:   else {  
6:     free p;  
7:     return NULL;  
8:   }  
9: }
```

SEFL

```
1: filter(){  
2:   constrain(IpDst==80);  
3: }
```

Only relevant paths explored
Concise

Symnet symbolic execution tool

- 10K LOC of Scala; Z3 for constraint solving

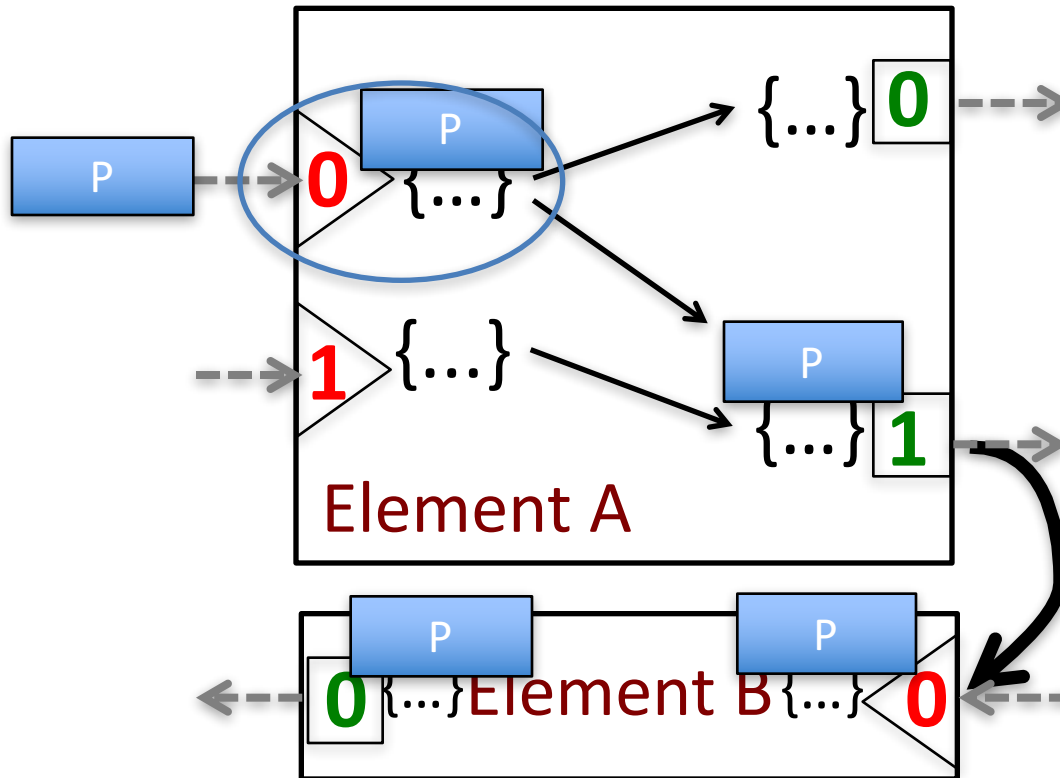
Input: SEFL network model

- SEFL models of individual network elements
- Connections between elements

Output: all feasible symbolic paths

- Values of header and metadata fields
- Path constraints

SEFL Network Models



Symbolic execution of filter + DNAT

Element A model

```
InputPort(0):  
  Constrain(IPDst==1.1.1.1),  
  If (Constrain(TcpDst==20),  
    InstructionBlock(  
      Assign(IPDst,192.168.0.1),  
      Assign(TcpDst,30),  
      Forward(OutputPort(0))  
    ),  
    Forward(OutputPort(1)),
```

Packet 1

```
IpDst=*      TcpDst=*  
IpDst=1,1... TcpDst=*  
IpDst=1.1... TcpDst=20  
IpDst=192... TcpDst=20  
IpDst=192... TcpDst=30  
CrtPort = 0
```

Packet 2

```
IpDst=1.1,  
  TcpDst != 20  
  
CrtPort = 1
```

- Reachability
- Loop detection
- Invariant header fields
- Header memory safety

Ready-made network models

Modeling network boxes is fairly difficult

We have developed parsers that output SEFL code from:

- Router/switch forwarding table snapshots
- CISCO ASA firewall configuration
- Click modular router configurations
- Openstack Neutron network configurations

Evaluation

Model correctness

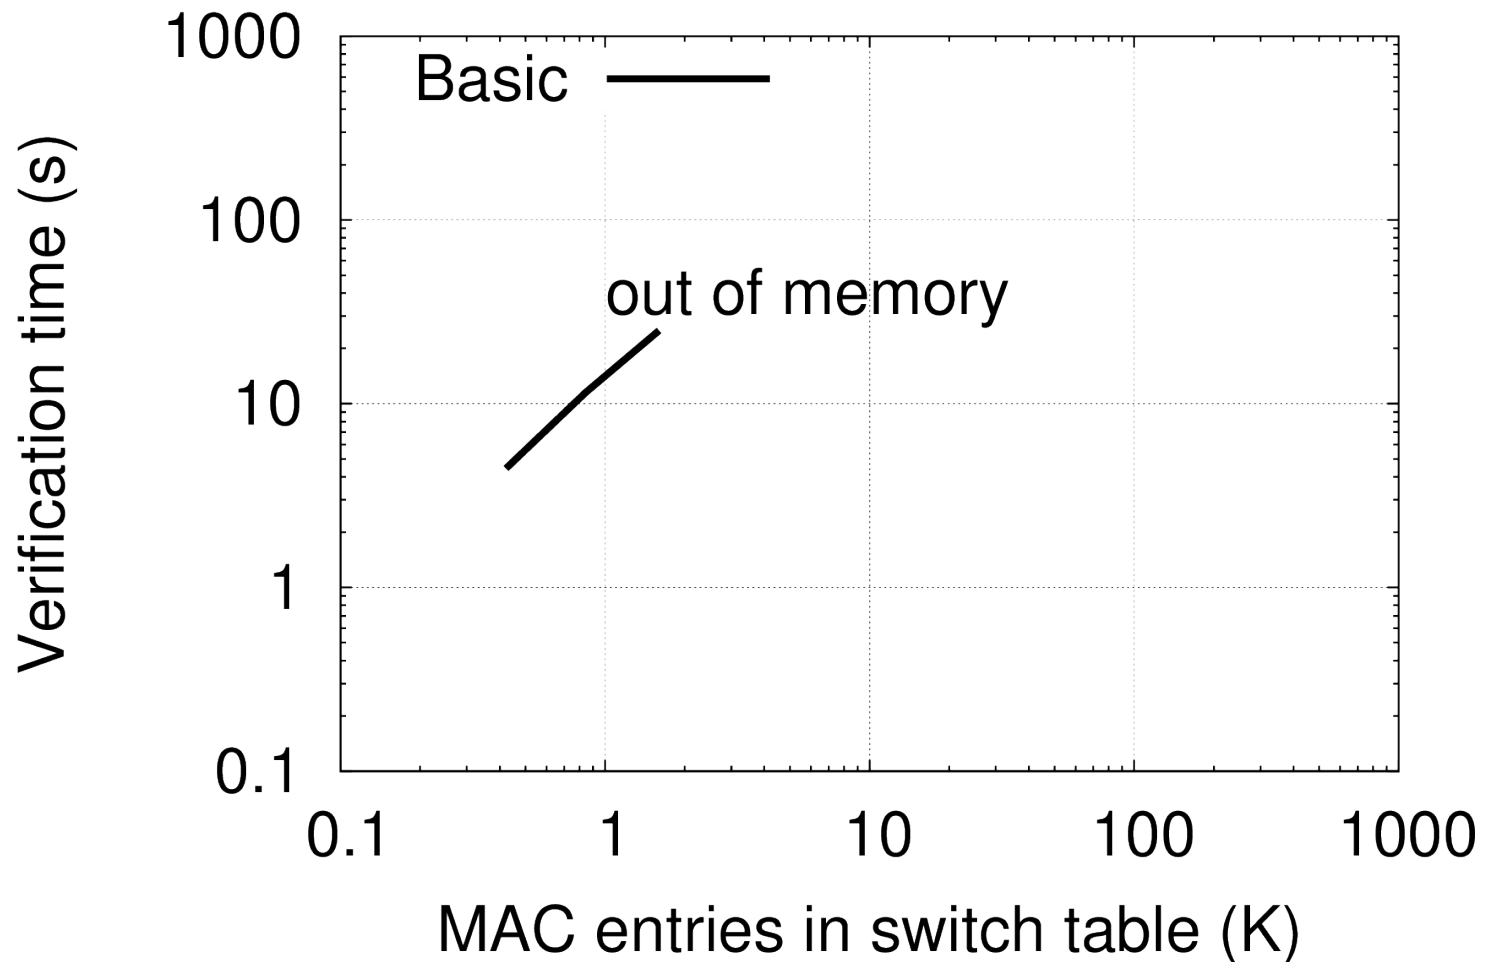
Functionality

Scalability


Verifiable properties

Property	HSA	NoD	SymNet
Reachability	✓	✓	✓
Loop Detection	✓	✗	✓
Header Field Invariance	✗	✗	✓
Arbitrary Packet Layout	✗	✓	✓
Tunneling	✗	✗	✓
Stateful Data Plane Processing	✗	✓	✓
Payload-sensitive Processing	✗	✗	✗
Properties Across Multiple Flows	✗	✗	✗


Does Symnet scale?



Does Symnet scale?



Does Symnet scale?



Analyzing bigger networks

- Stanford university backbone network
- Switches, routers and VLANs
 - Two-layer topology
 - Core routers have 180.000 entries in their FIBs

	HSA	Symnet
Model Generation Time	3.2 min	8.1 min
Runtime	24s	37s

Conclusions

SEFL + Symnet offers a deeper understanding of modern data planes at a low price.

Symnet is open-source

Check demo session tomorrow

Backup slides

TCP options parsing

Symbolic variable Path 1

```
int crt = 0;
while (crt >= 0 && crt < length &&
      options[crt]) {
    switch(options[crt]) {
        case 1:
            crt++; break;
        case 2: //MSS
        case 3: //WINDOW SCALE
        case 4: //SACK PERMITTED
        case 8: //TIMESTAMP
            crt += options[crt+1]; break;
        default:
            //unknown options, scrub
            int len = options[crt+1];
            for (i=crt; i<crt+len; i++)
                options[i] = 1;
            crt += len; break;
    }
}
```

TCP options parsing

```
int crt = 0;
while (crt >= 0 && crt < length &&
      options[crt]){
    switch(options[crt]){
        case 1:
            crt++; break;
        case 2://MSS
        case 3://WINDOW SCALE
        case 4://SACK PERMITTED
        case 8://TIMESTAMP
            crt += options[crt+1]; break;
        default:
            //unknown options, scrub
            int len = options[crt+1];
            for (i=crt;i<crt+len;i++)
                options[i] = 1;
            crt += len; break;
    }
}
```

Path 1

Path 2

Path 3

options[0]==1

options[0] in
{2,3,4,8}

options[0]
not in
{1,2,3,4,8}

TCP Options parsing

Leave the TCP options header outside of symbolic execution

Model TCP options as metadata instead

“OPT-x” models the presence of option x

“SZ-x” size of the option in bytes

“DATA-x” value of the option

```
Assign("OPT30", ConstantValue(0)),  
If (Constrain(TcpDst==80), Assign("OPT4", 0), NoOp),  
Assign("OPT2", 1),  
Assign("SIZE2", 4),  
If (Constrain("VAL2">1380), Assign("VAL2", 1380), NoOp)
```

Does Symnet scale?

Symbolic execution of a core router

Prefixes	Basic	Ingress	Egress
1600	25s	2.1s	0.4s
62500	DNF	23.1s	5.6s
188500	DNF	DNF	18s

Running Klee for options parsing

Length	1	2	3	4	5	6
Number of paths	4	67	140	464	1095	3081
Runtime (s)	0.3	8	20	420	1500	9120