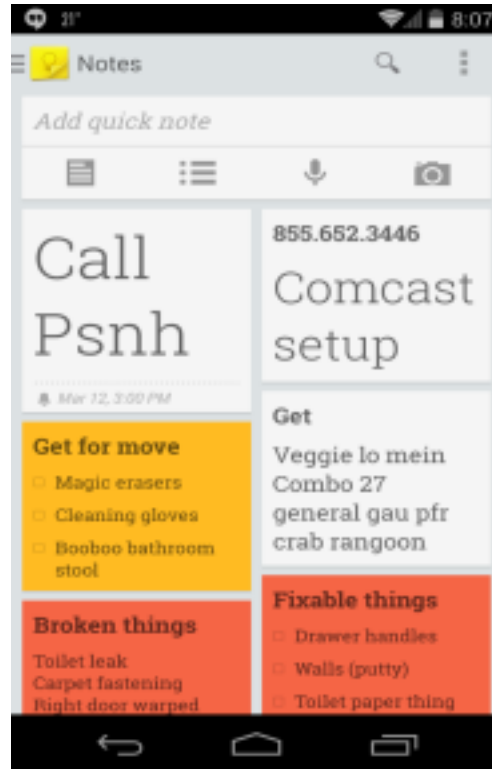


# **WebPerf: Evaluating “What- If” Scenarios for Cloud-hosted Web Applications**

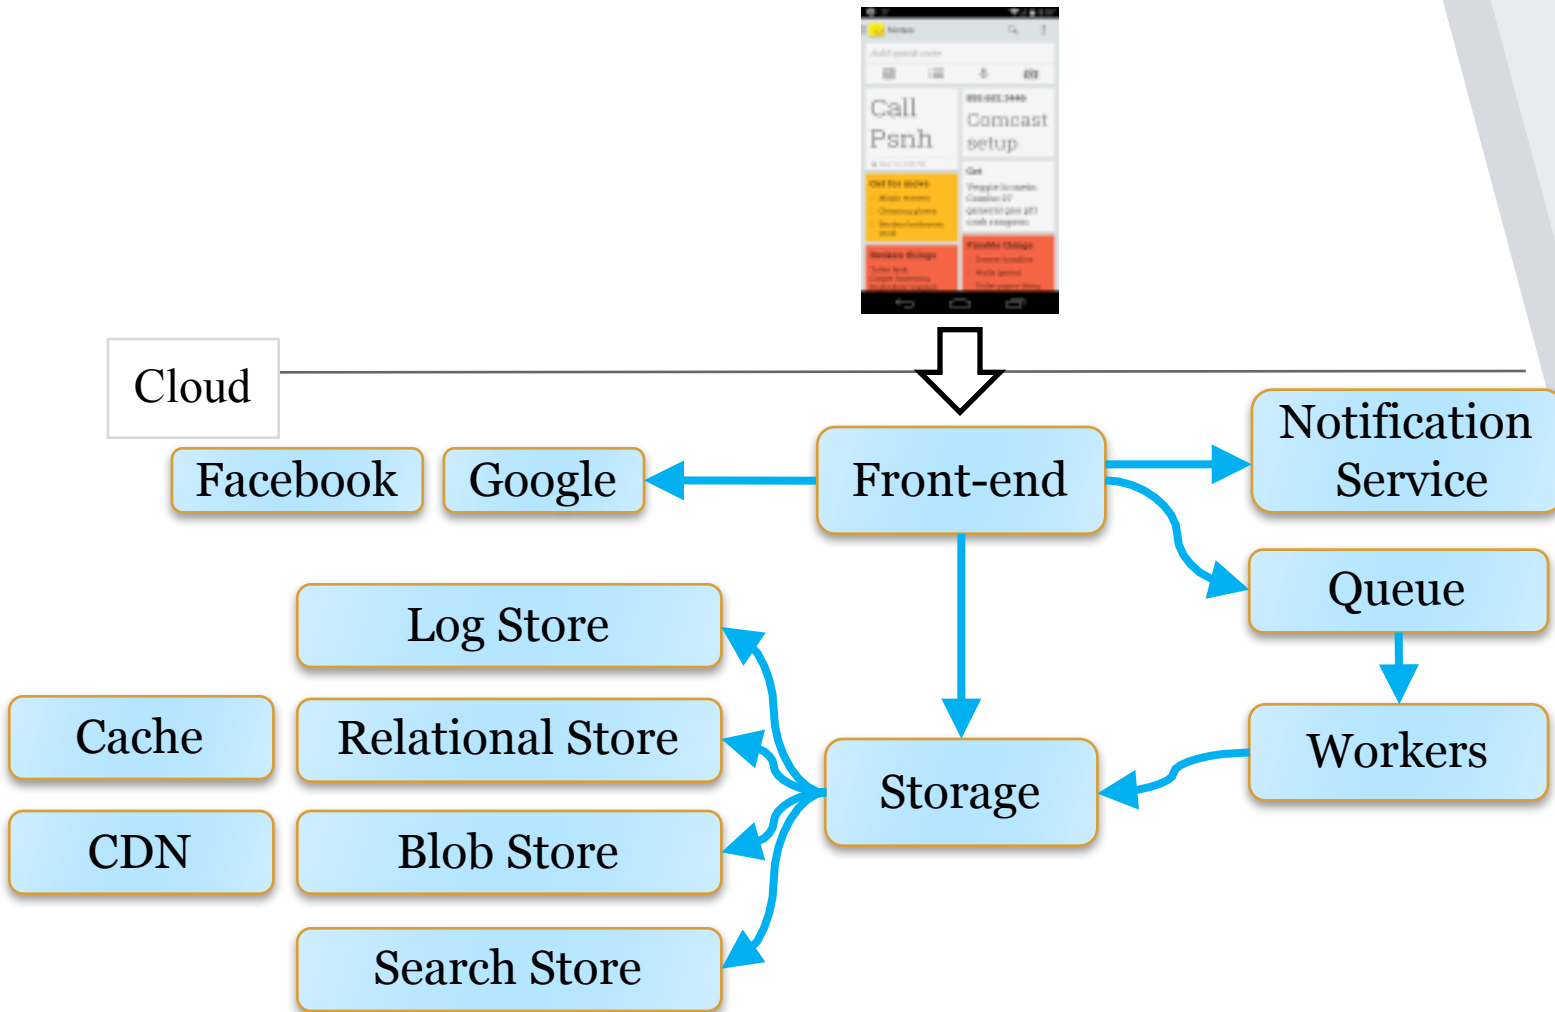
Yurong Jiang  
Lenin Ravindranath  
Suman Nath  
Ramesh Govindan



# A Cloud-Hosted Web Application



# Modern Cloud Applications are Complex



## Cloud-side Latency

Latency of these applications is critical for user experience

Developers find it hard to optimize cloud-side latency for cloud-hosted Web applications

## Front-end

| INSTANCE | CORES | RAM     | DISK SIZES | PRICE                    |
|----------|-------|---------|------------|--------------------------|
| A0       | 1     | 0.75 GB | 19 GB      | \$0.02/hr<br>(~\$15/mo)  |
| A1       | 1     | 1.75 GB | 224 GB     | \$0.08/hr<br>(~\$60/mo)  |
| A2       | 2     | 3.5 GB  | 489 GB     | \$0.16/hr<br>(~\$119/mo) |
| A3       | 4     | 7 GB    | 999 GB     | \$0.32/hr<br>(~\$238/mo) |
| A4       | 8     | 14 GB   | 2,039 GB   | \$0.64/hr<br>(~\$476/mo) |

*Each choice impacts latency*

Relational Store

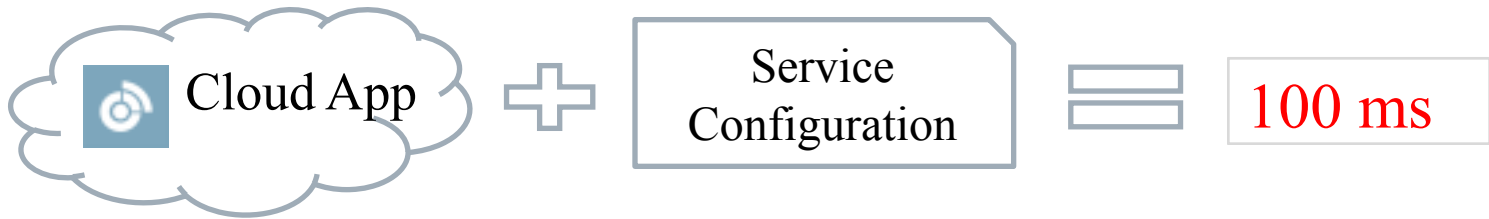
Azure SQL

| Premium |                   |                    |                    |
|---------|-------------------|--------------------|--------------------|
|         | DTUs <sup>2</sup> | MAX STORAGE PER DB | PRICE <sup>1</sup> |
| P1      | 125               | 500 GB             | ~\$465/mo          |
| P2      | 250               | 500 GB             | ~\$930/mo          |
| P4      | 500               | 500 GB             | ~\$1,860/mo        |
| P6      | 1000              | 500 GB             | ~\$3,720/mo        |
| P11     | 1750              | 1 TB               | ~\$7,001/mo        |

*Latency implications hard to understand*

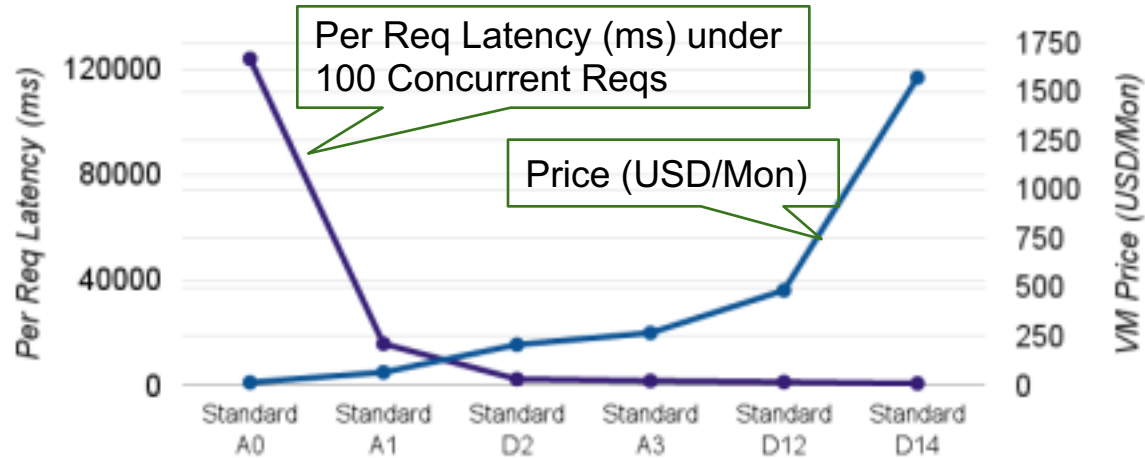
What if?

*What if I move the blob store  
from basic to standard tier?*



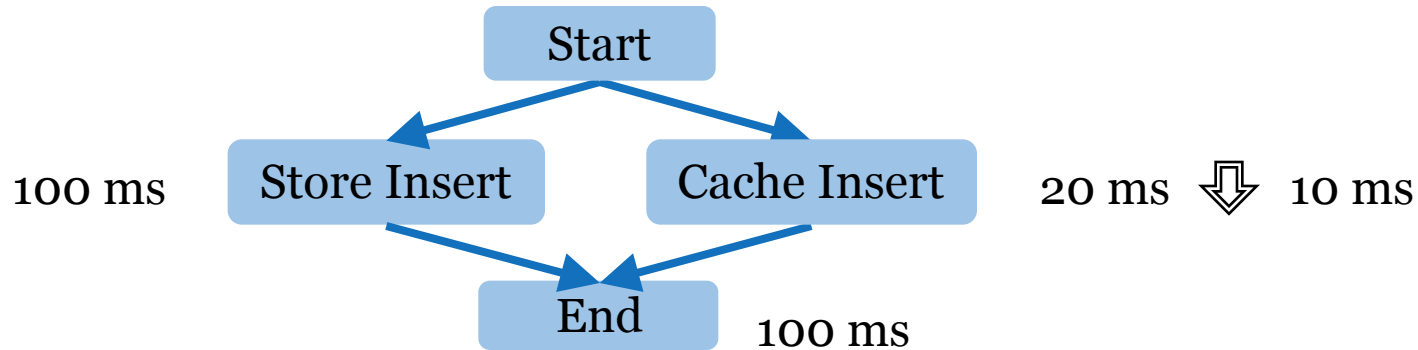
\$30  $\Rightarrow$  \$100

Answer to what-if question  
may depend on workload





Answer to what-if question may depend on causal dependencies



A what-if capability  
should be expressive

Relational Store



What if I re-locate this  
component?

Relational Store



What if I increase this  
component's load?

Table Store

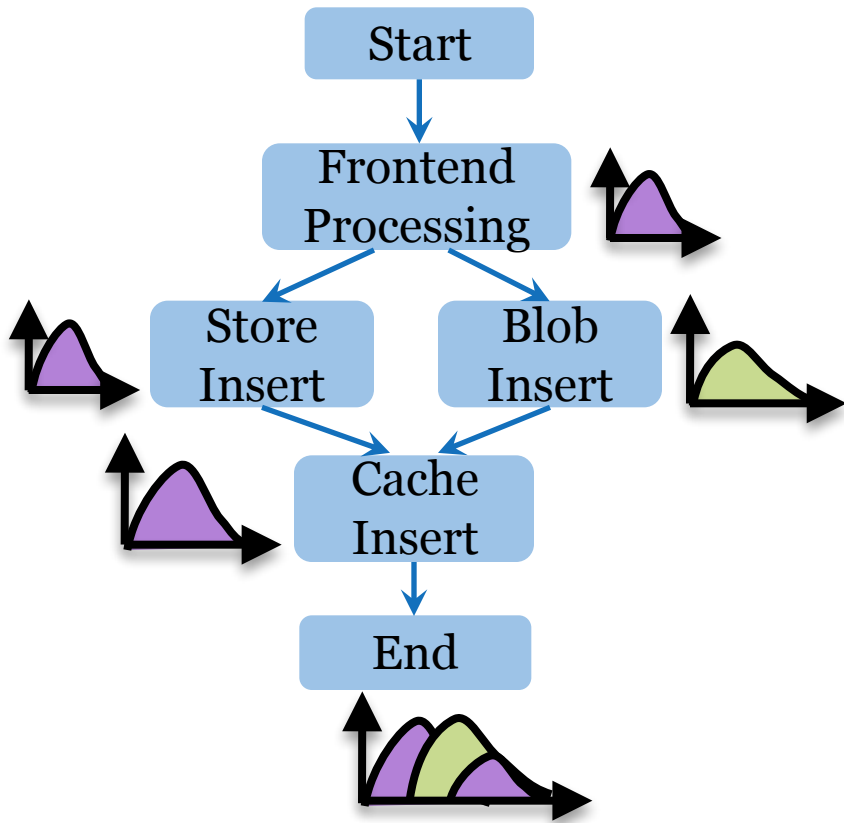


What if a replica fails?

## WebPerf is a what-if scenario evaluator

- ❖ Input: a what-if scenario
- ❖ Output: resulting **cloud-side latency distribution**

*What if I upgrade blob storage  
from basic to standard tier?*



Dependency graph  
extraction

Baseline latency  
estimation

Component Profiling

Cloud-side latency  
estimation

Developer  
supplies  
workload

Computed  
Offline

Why might WebPerf work?

## Cloud deployments well-engineered

- ❖ Components designed for predictable latency
- ❖ Often co-located in same datacenter

➤ Cheap      ➤ Fast      ➤ Low Effort

Automate  
most steps

Why?

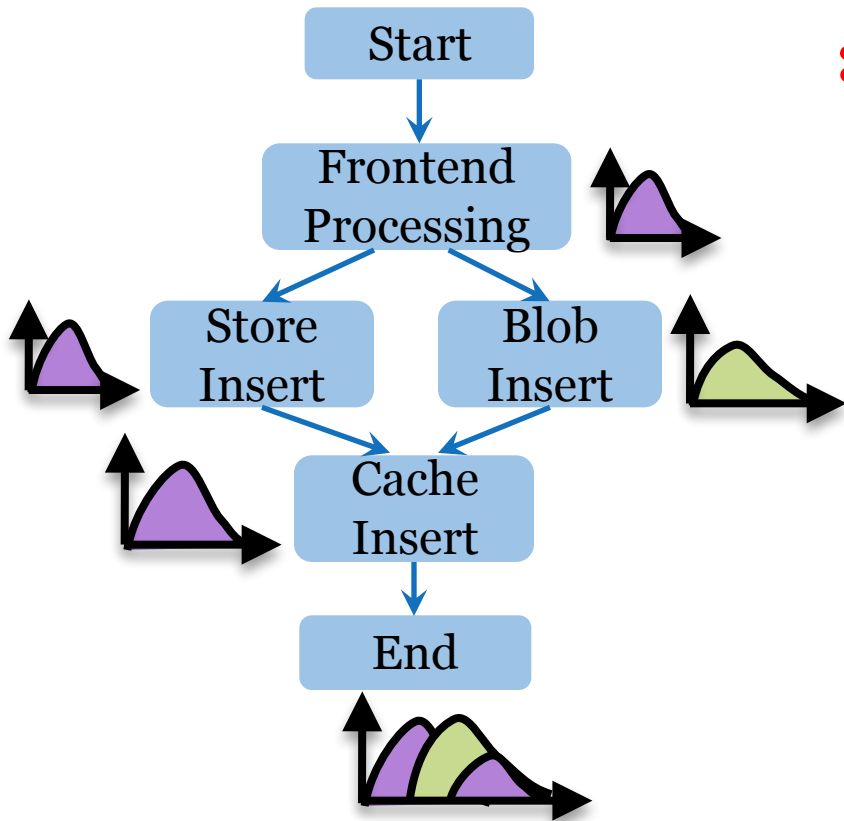
Many component profiles are application-independent

Compute  
offline,  
reuse

The dependency graph is usually independent of what-if scenario

Compute  
once

*What if I upgrade blob storage  
from basic to standard tier?*



Dependency graph  
extraction

Baseline latency  
estimation

Component Profiling

Cloud-side latency  
estimation

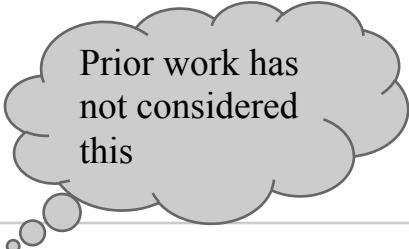
## Goal

*Fast, accurate* dependency extraction  
with *zero developer input*

## Approach

Track dependencies at run-time by  
*instrumenting binary*





Prior work has  
not considered  
this

## *Task Asynchronous Programming*

- ❖ Many cloud apps use this
- ❖ Only mechanism for asynchronous I/O in Azure
- ❖ AWS provides APIs for .NET

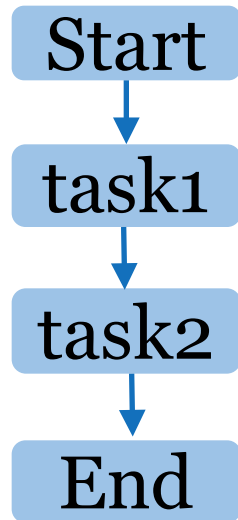
```
async processRequest (input)
{
    /* process input */
    task1 = store.get(key1);
    value1 = await task1;
    task2 = cache.get(key2);
    value2 = await task2;
    /* construct response */
    return response;
}
```

On front-  
end

Start task

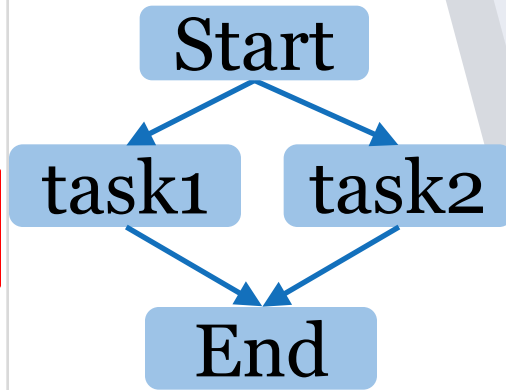
Continue

```
async processRequest (input)
{
    /* process input */
    task1 = store.get(key1);
    value1 = await task1;
    task2 = cache.get(key2);
    value2 = await task2;
    /* construct response */
    return response;
}
```



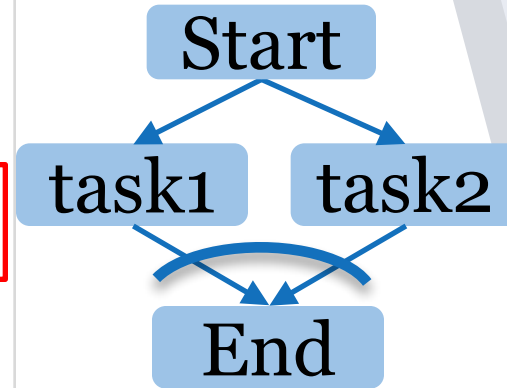
## WhenAll: Continue only when all tasks finish

```
async processRequest (input)
{
    /* process input */
    task1 = store.get(key1);
    task2 = cache.get(key2);
    value1, value2 = await
        Task.WhenAll(task1, task2);
    /* construct response */
    return response;
}
```



## WhenAny: Continue when *any one* task finishes

```
async processRequest (input)
{
    /* process input */
    task1 = store.get(key1);
    task2 = cache.get(key2);
    value1, value2 = await
        Task.WhenAny(task1, task2);
    /* construct response */
    return response;
}
```



```
async processRequest (input)
{
    /* process input */
    task1 = store.get(key1);
    value1 = await task1;
    task2 = cache.get(key2);
    value2 = await task2;
    /* construct response */
    return response;
}
```

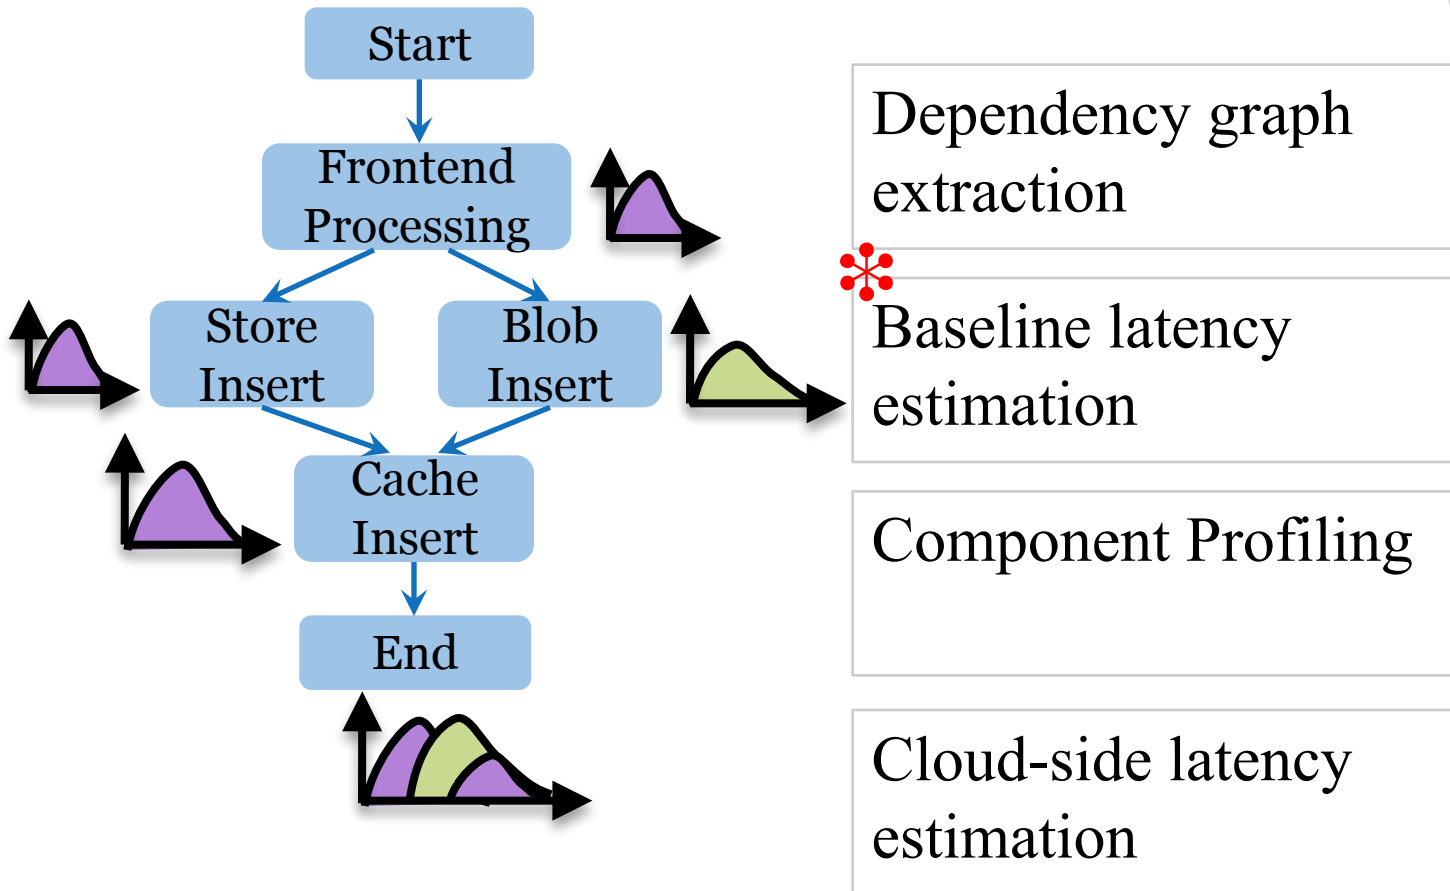
Init

Received  
value1Received  
value2.NET compiler  
generates this

Instrument state machine *binary* to  
*dynamically* track tasks and continuations

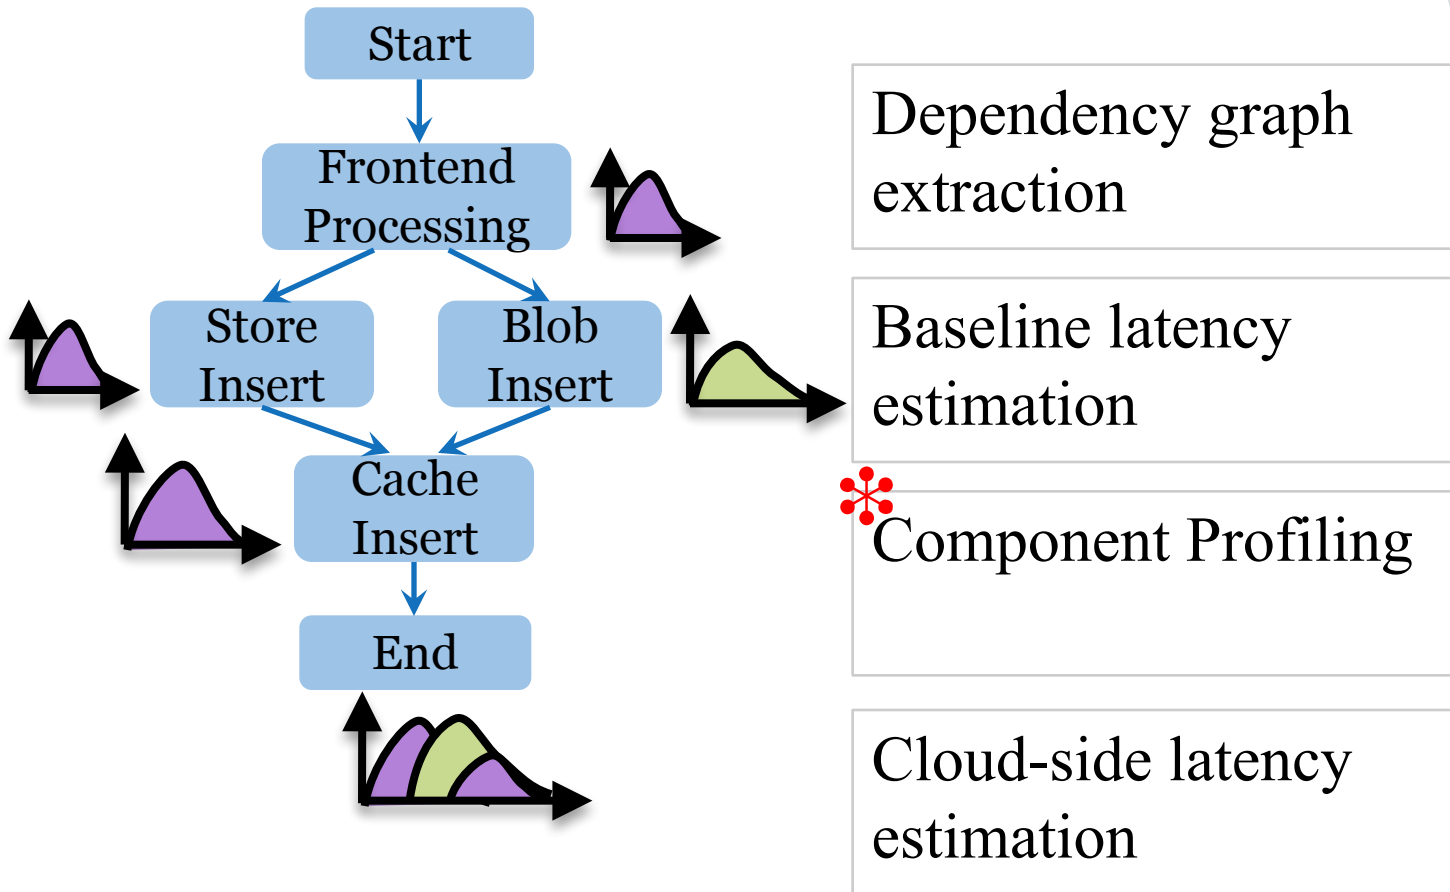
*What if I upgrade blob storage  
from basic to standard tier?*

## WebPerf Approach



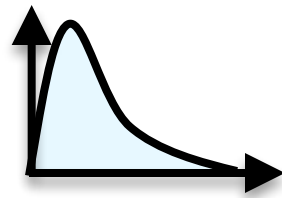
*What if I upgrade blob storage  
from basic to standard tier?*

## WebPerf Approach

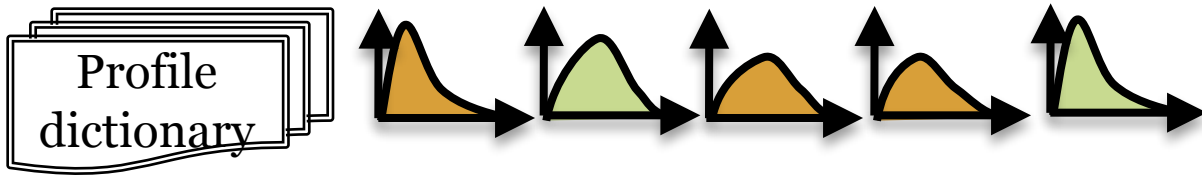




A component's *profile* contains latency distributions of API calls to component



WebPerf profiles commonly used components *offline*



Relational Store



Tiers

Relational Store



Location

Relational Store



Load

Table Store



Failure

# Not all profiles can be computed offline

Relational Store

Azure SQL

SQL join latency  
depends on size

Cache

Redis Cache

Cache latency  
depends on hit rate

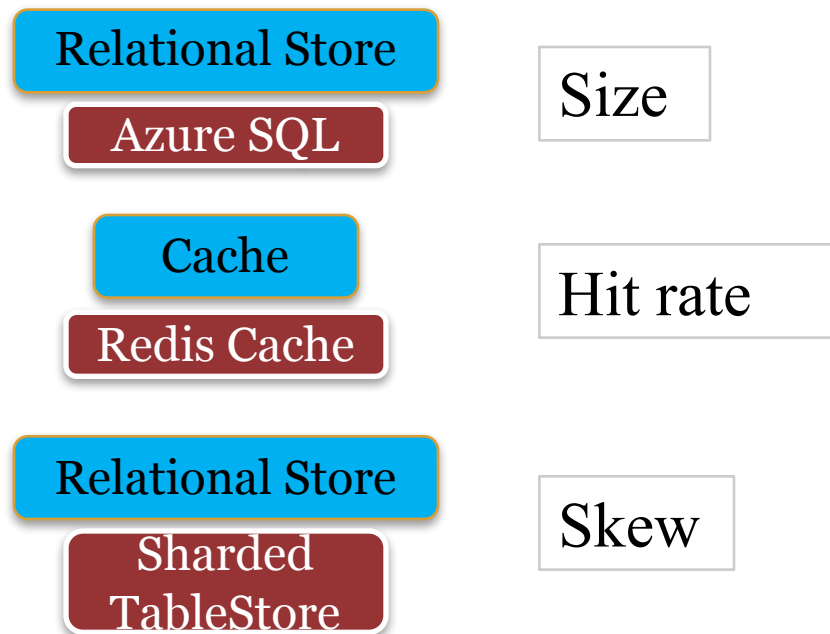
Relational Store

Sharded  
TableStore

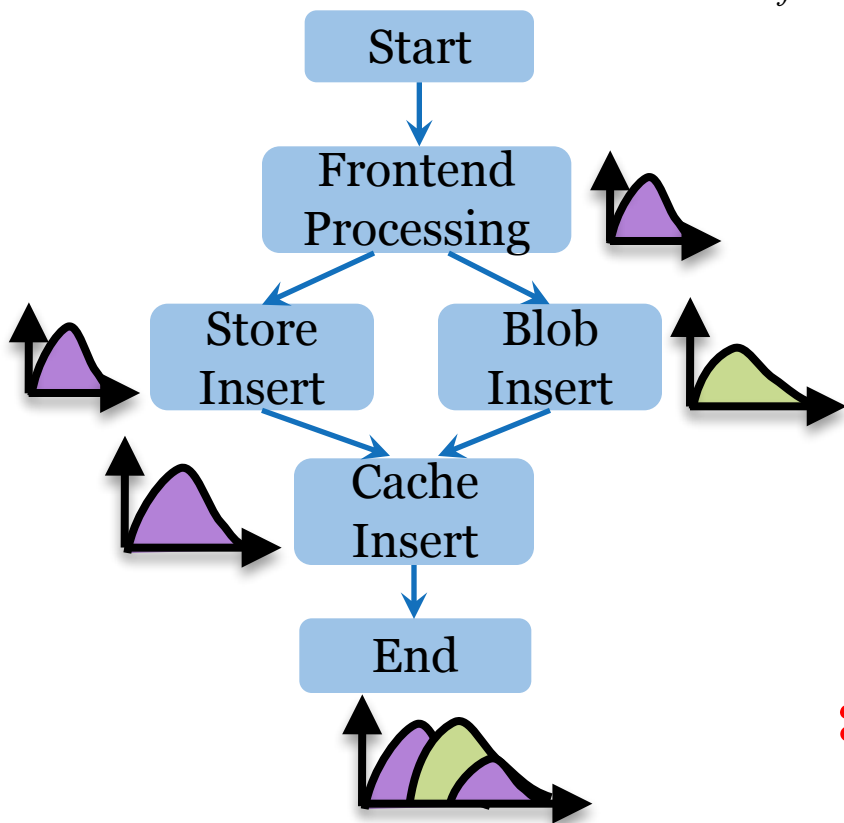
Access latency  
depends on skew

# WebPerf uses *parameterized profiles*

- ❖ User must specify *workload hint*



*What if I upgrade blob storage  
from basic to standard tier?*



Dependency graph  
extraction

Baseline latency  
estimation

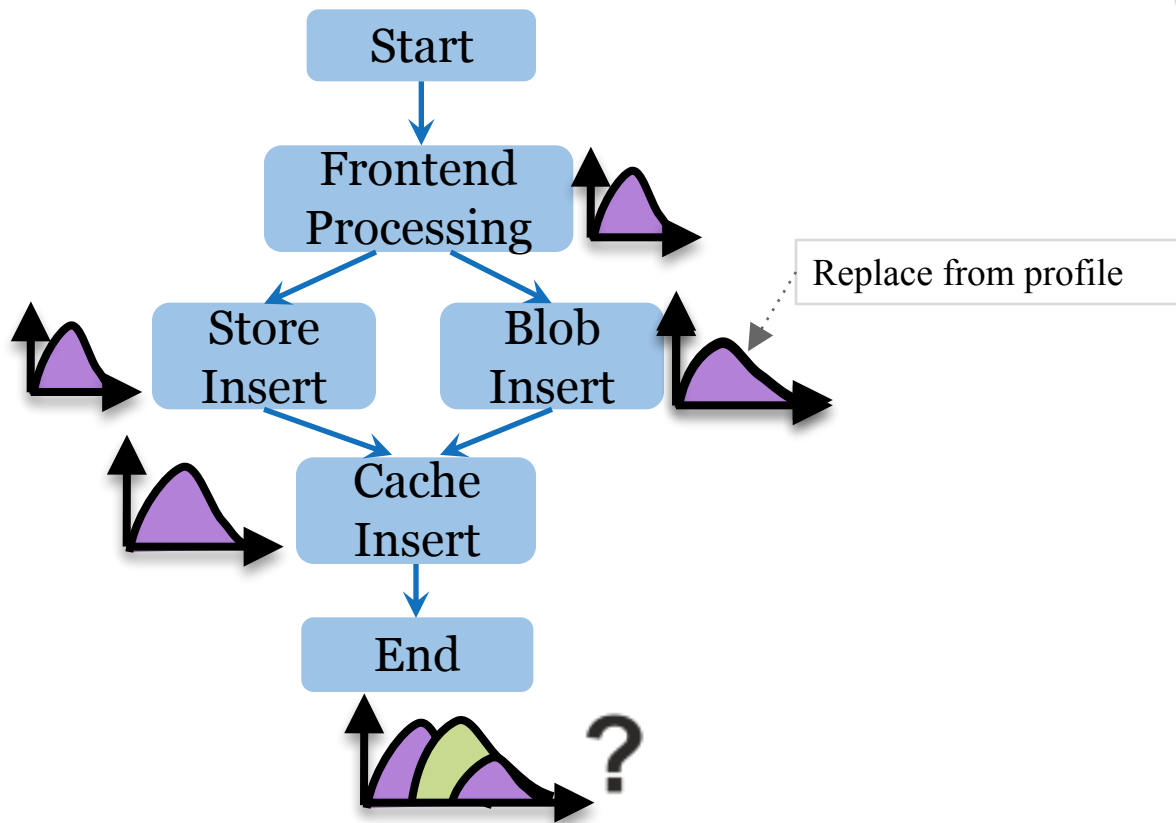
Component Profiling



Cloud-side latency  
estimation

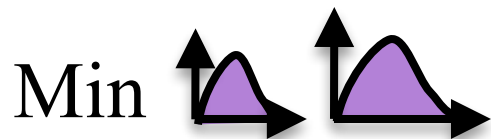
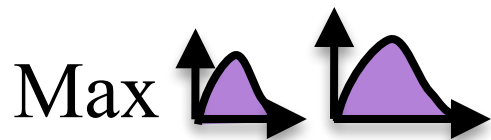
*What if I upgrade blob storage  
from basic to standard tier?*

## Cloud-Side Latency Estimation





Simple operations on  
distributions suffice



WebPerf is *accurate, fast, cheap, and requires low developer effort*

- ❖ How accurate is WebPerf?
- ❖ Are workload hints necessary?



Different  
functionality

Different  
components

Varying  
complexity

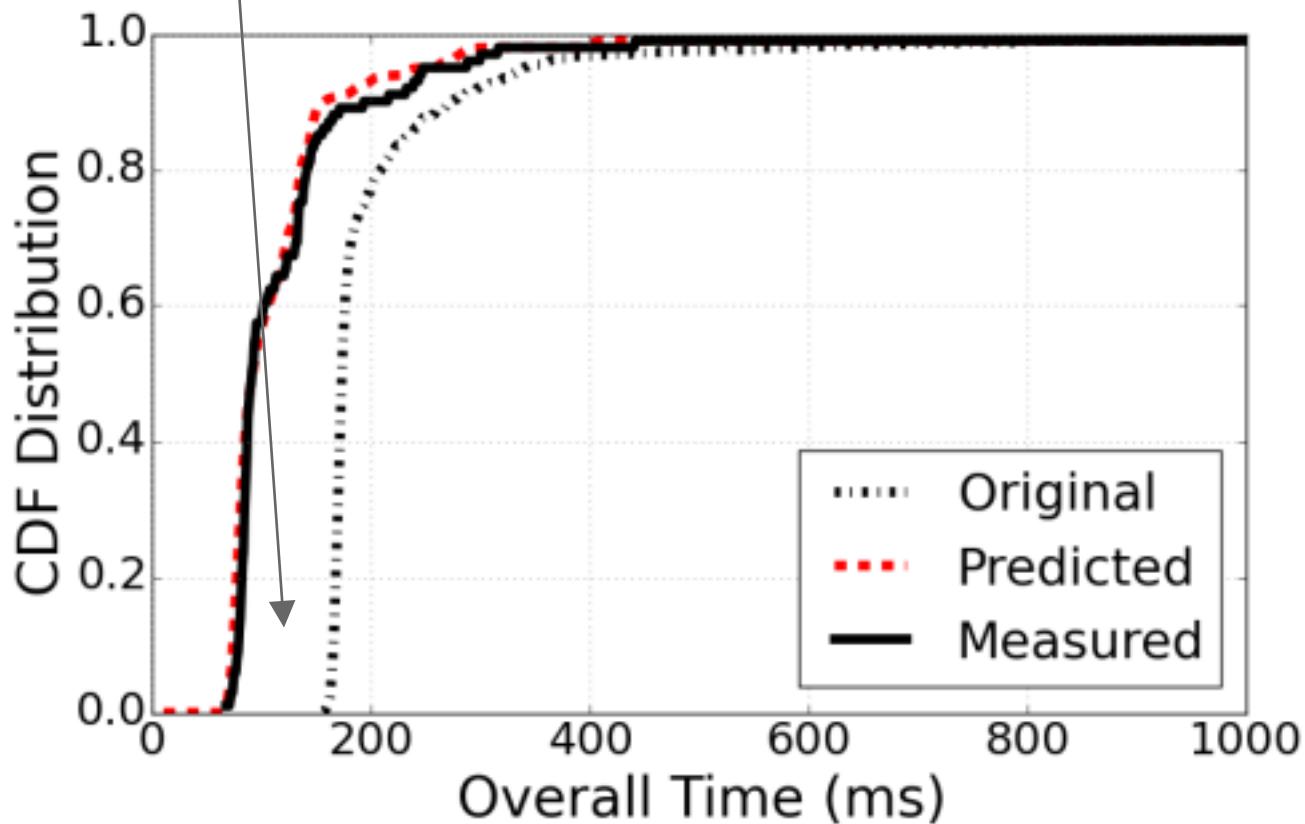
| Application     | Azure components used                                 | Average I/O Calls |
|-----------------|---|-------------------|
| SocialForum     | Blob storage, Redis cache, Service bus, Search, Table | 116               |
| SmartStore.Net  | SQL   | 41                |
| ContosoAds      | Blob storage, Queue, SQL, Search                      | 56                |
| EmailSubscriber | Blob storage, Queue, Table                            | 26                |
| ContactManager  | Blob storage, SQL                                     | 8                 |
| CourseManager   | Blob storage, SQL                                     | 44                |

| What-if scenario   | Example  |
|--|--|
| <b>Tier:</b> A component X is upgraded to tier Y                                     | X = A Redis cache, Y = a standard tier (from a basic tier) |
| <b>Load:</b> X concurrent requests to component Y                                    | X = 100 , Y = the application or a SQL database            |
| <b>Interference:</b> CPU and/or memory pressure, from collocated applications, of X% | X = 50% CPU, 80% memory                                    |
| <b>Location:</b> A component X is deployed at location Y                             | X = A Redis Cache or a front end, Y = Singapore            |
| <b>Failure:</b> An instance of a replicated component X fails                        | X = A replicated front-end or SQL database                 |

Configuration choices can significantly impact latency

*What if I move the Redis cache in SocialForum from basic to standard tier?*

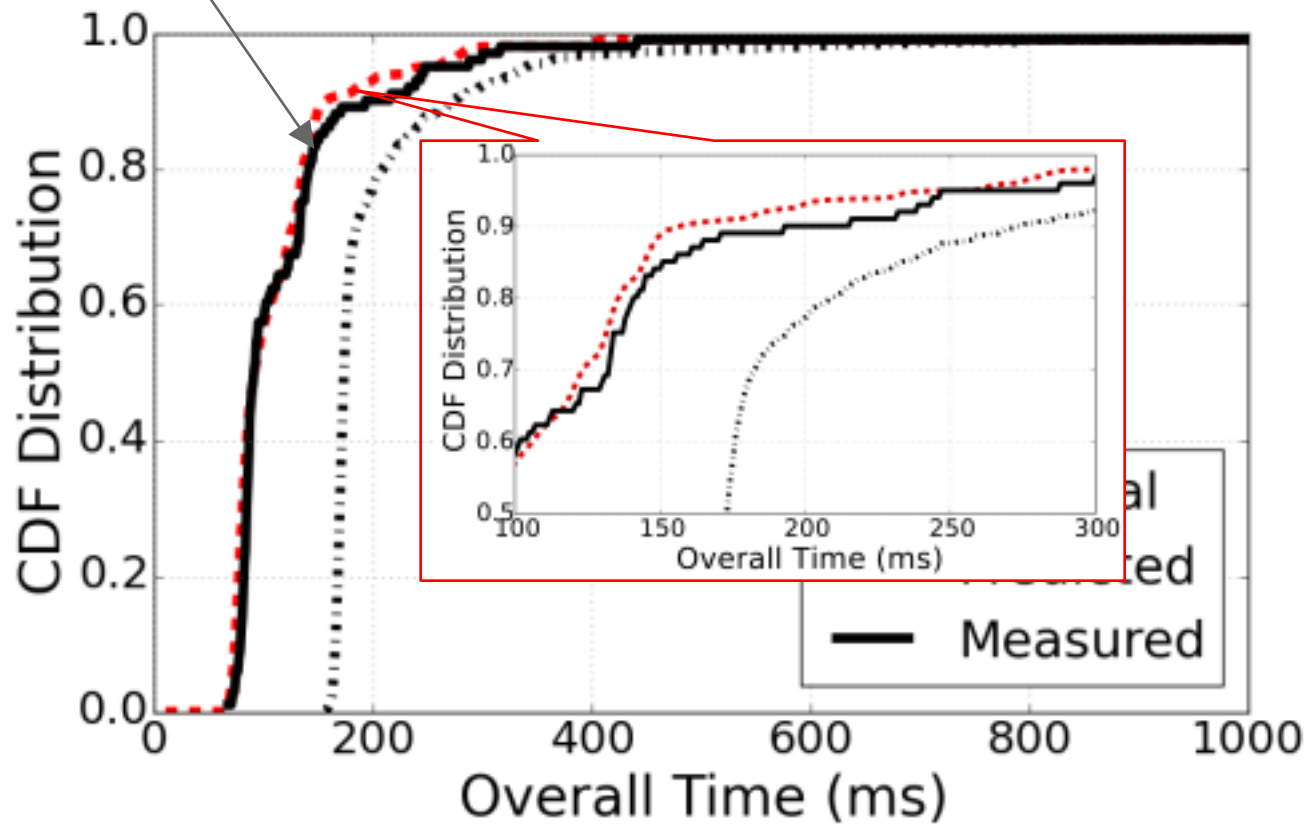
Accuracy



Prediction closely matches ground truth

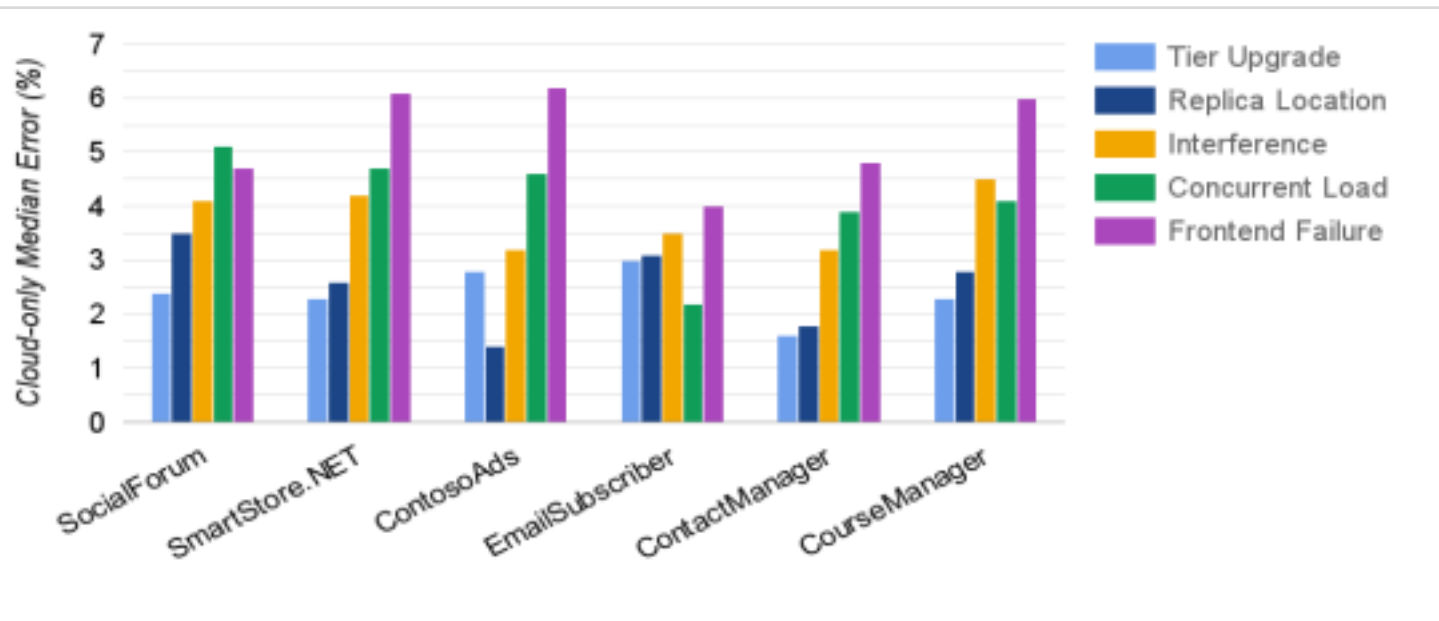
*What if I move the Redis cache from basic to standard tier?*

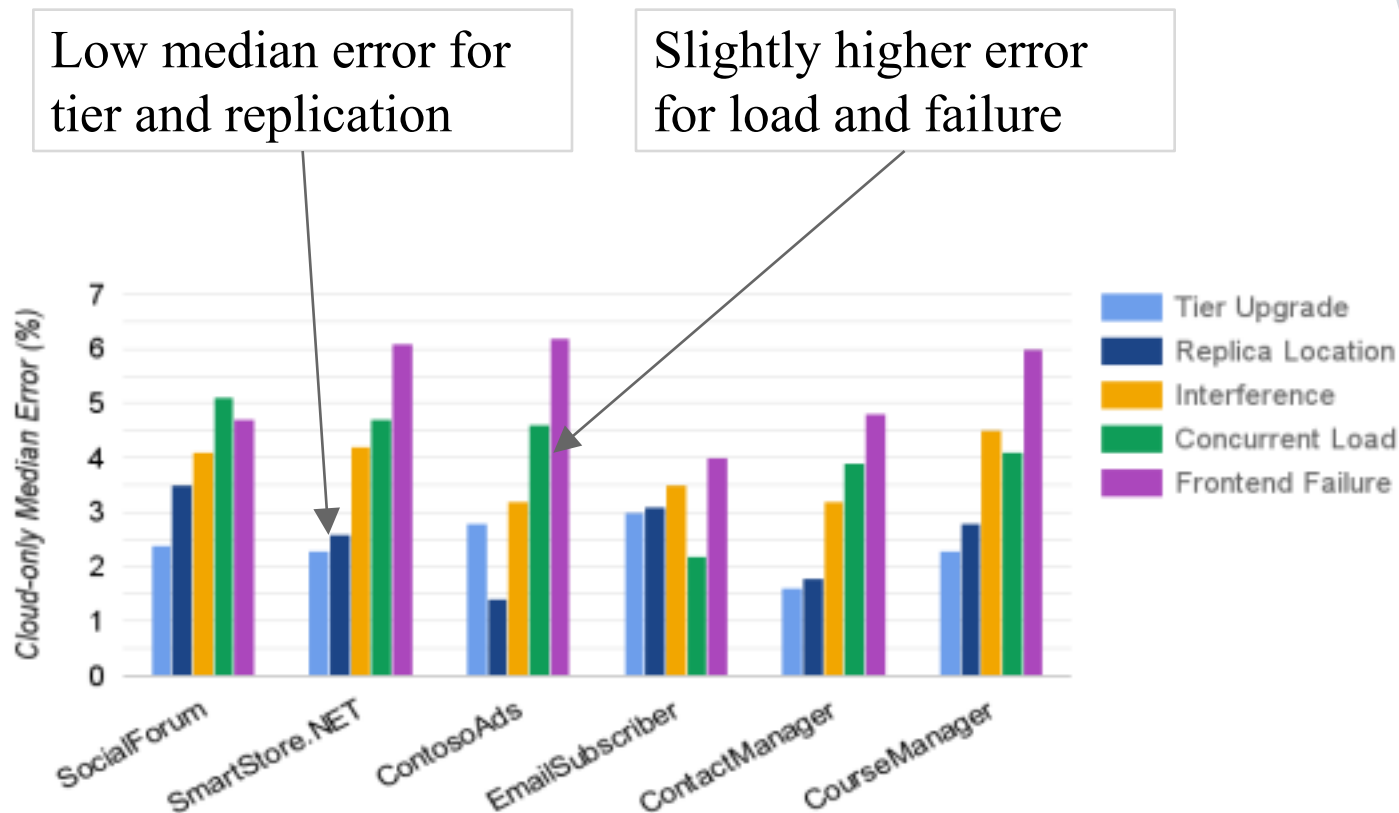
Accuracy

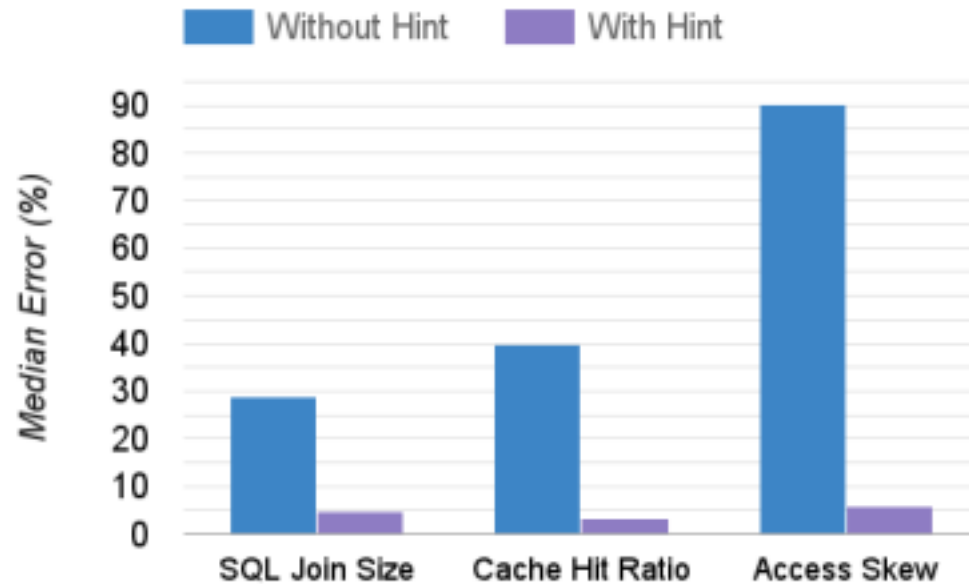


Median prediction error  
under 7%

Difference between predicted  
distribution and ground truth





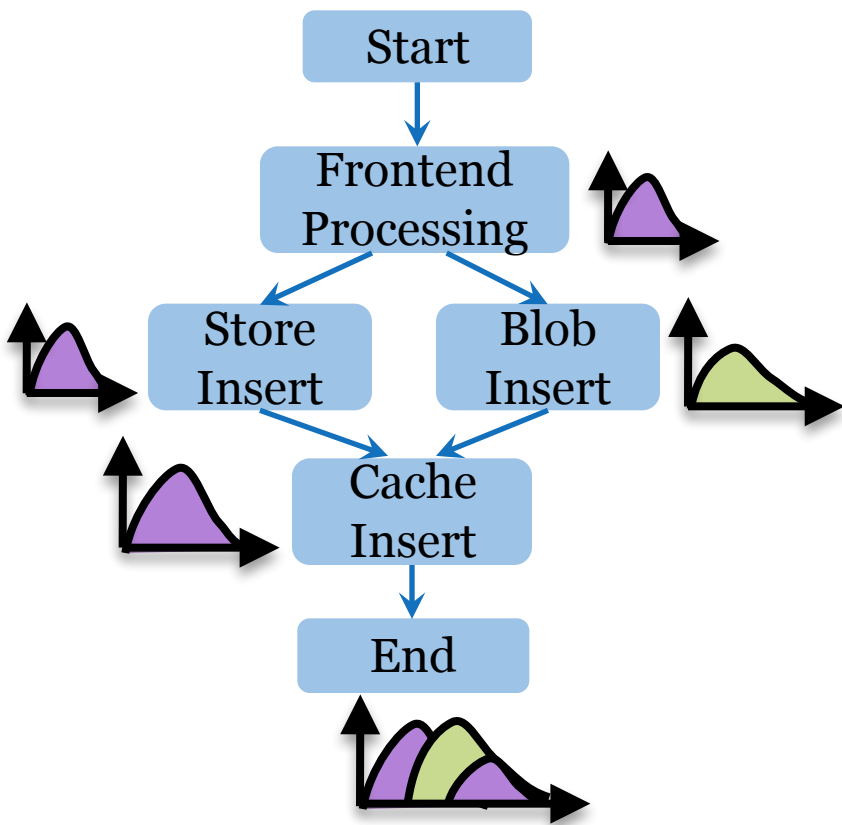


Workload hints can significantly improve accuracy

WebPerf predicts cloud-side latency distributions for different what-if scenarios

It accurately tracks dependencies and profiles components offline

Across six different applications and scenarios, its error is less than 7%





# WebPerf Contributions and Summary

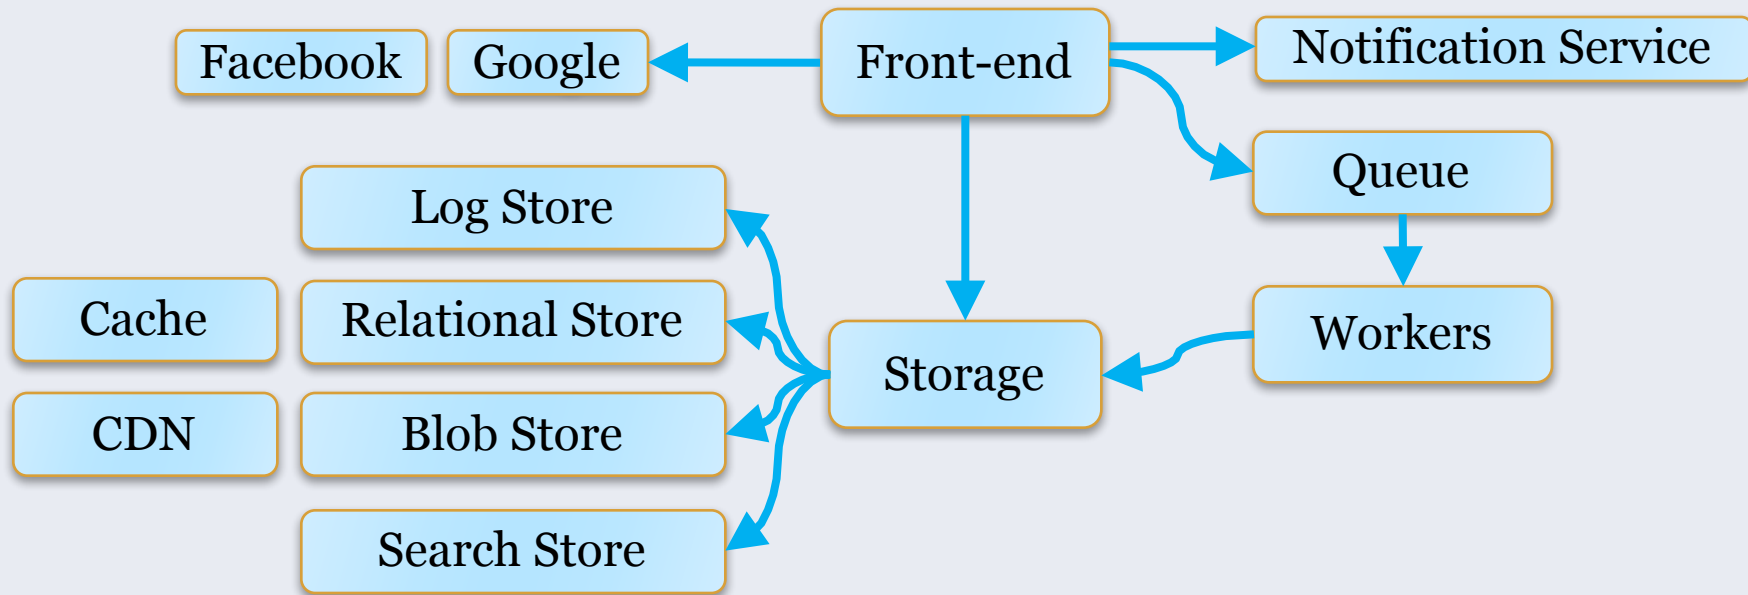
An automated tool to instrument web apps and capture both browser objects and front end cloud processing dependency

Predicting web app cloud latency and end-to-end latency in probabilistic setting under six different scenarios

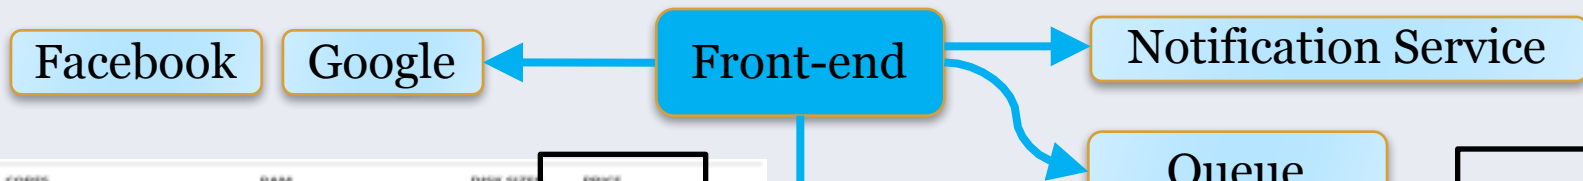
Evaluations with six real websites show WebPerf achieves < 7% median prediction error

# Thank you

# Large number of configuration choices



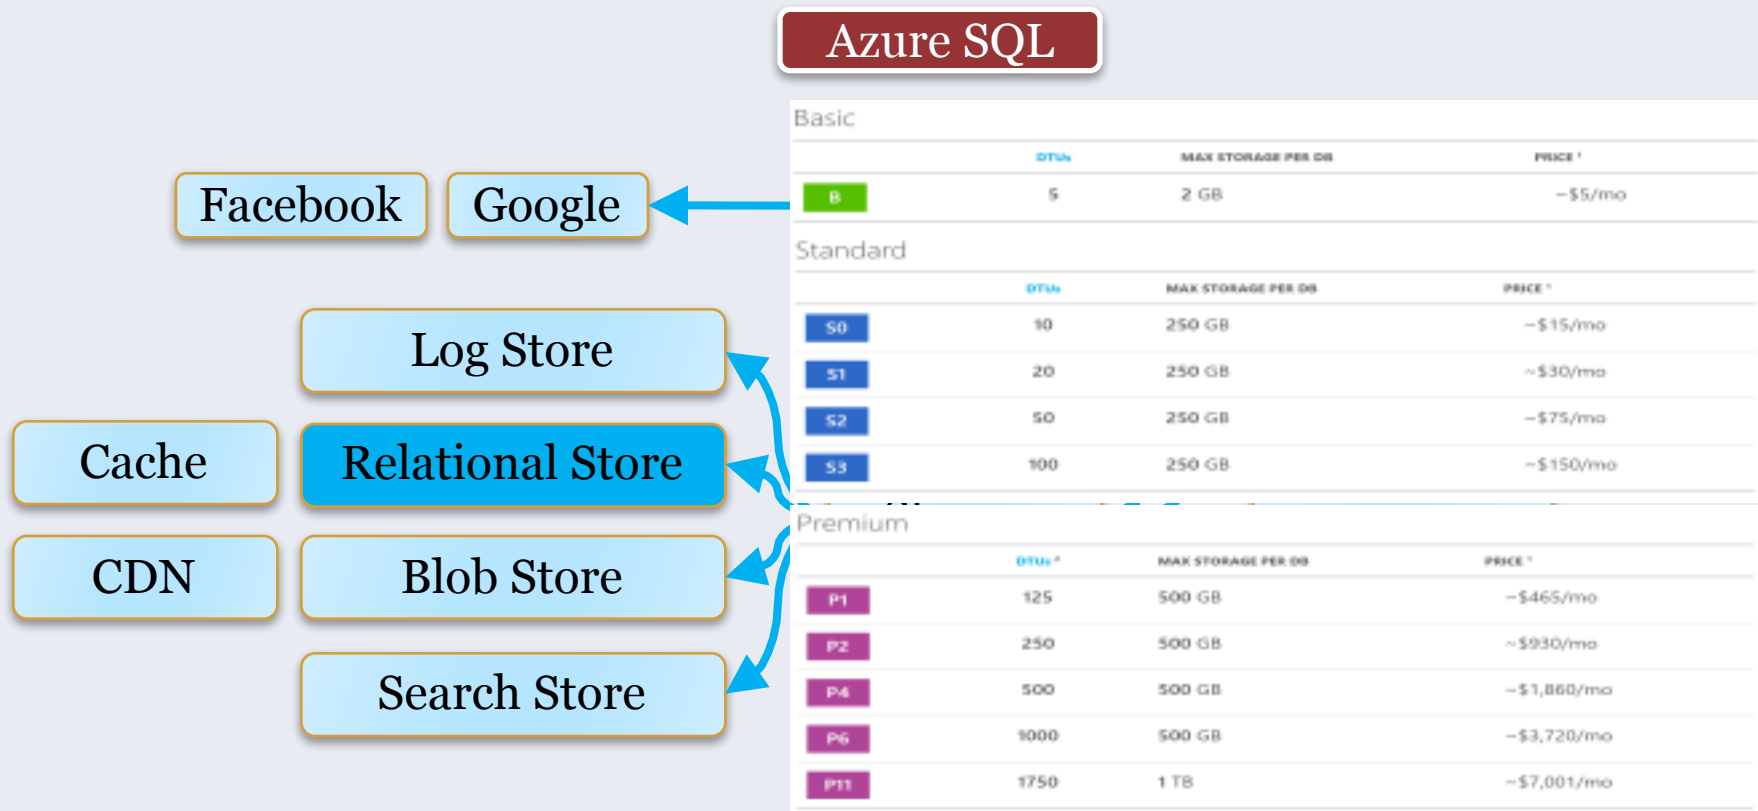
# Large number of configuration choices



| INSTANCE | CORES | RAM     | DISK SIZE | PRICE                      |
|----------|-------|---------|-----------|----------------------------|
| A0       | 1     | 0.75 GB | 19 GB     | \$0.02/hr<br>[~\$15/mo]    |
| A1       | 1     | 1.75 GB | 224 GB    | \$0.06/hr<br>[~\$60/mo]    |
| A2       | 2     | 3.5 GB  | 489 GB    | \$0.16/hr<br>[~\$119/mo]   |
| A3       | 4     | 7 GB    | 999 GB    | \$0.32/hr<br>[~\$238/mo]   |
| A4       | 8     | 14 GB   | 2,039 GB  | \$0.64/hr<br>[~\$476/mo]   |
| A5       | 2     | 14 GB   | 489 GB    | \$0.35/hr<br>[~\$260/mo]   |
| A6       | 4     | 28 GB   | 999 GB    | \$0.71/hr<br>[~\$528/mo]   |
| A7       | 8     | 56 GB   | 2,039 GB  | \$1.41/hr<br>[~\$1,049/mo] |

|     |    |        |        |                             |
|-----|----|--------|--------|-----------------------------|
| D1  | 1  | 3.5 GB | 50 GB  | \$0.16/hr<br>[~\$104/mo]    |
| D2  | 2  | 7 GB   | 100 GB | \$0.28/hr<br>[~\$208/mo]    |
| D3  | 4  | 14 GB  | 200 GB | \$0.56/hr<br>[~\$417/mo]    |
| D4  | 8  | 28 GB  | 400 GB | \$1.12/hr<br>[~\$833/mo]    |
| D11 | 2  | 14 GB  | 100 GB | \$0.33/hr<br>[~\$246/mo]    |
| D12 | 4  | 28 GB  | 200 GB | \$0.652/hr<br>[~\$485/mo]   |
| D13 | 8  | 56 GB  | 400 GB | \$1.173/hr<br>[~\$873/mo]   |
| D14 | 16 | 112 GB | 800 GB | \$2.111/hr<br>[~\$1,571/mo] |

# Large number of configuration choices



# Large number of configuration choices

Facebook

Google

Log Store

Cache

Relational Store

CDN

Blob Store

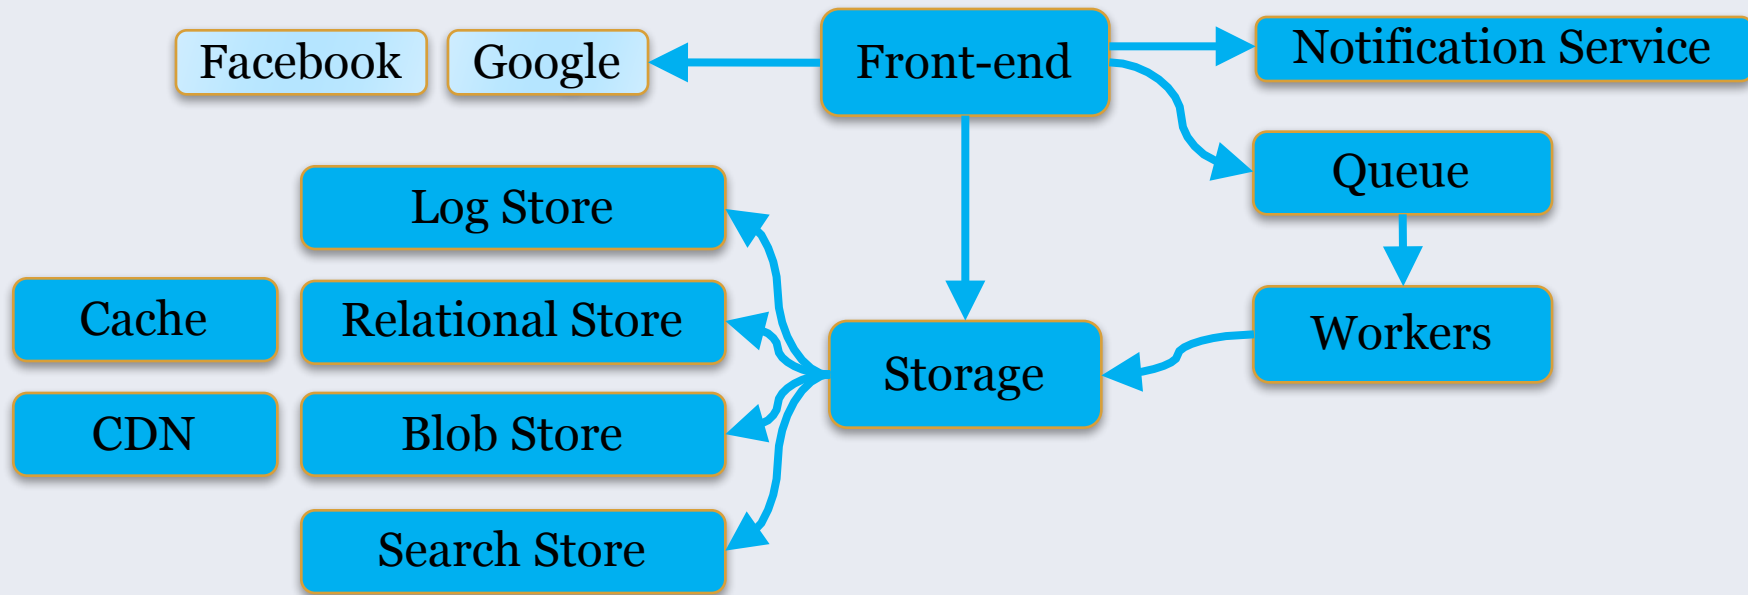
Search Store

Redis Cache

| CACHE NAME | CACHE SIZE | BASIC                    | NETWORK PERFORMANCE |       |
|------------|------------|--------------------------|---------------------|-------|
| C0         | 250 MB     | \$0.002/hr<br>(~\$16/mo) | Low                 |       |
| C1         | 1 GB       | \$0.05/hr<br>(~\$41/mo)  | Low                 | 1000  |
| C2         | 2.5 GB     | \$0.09/hr<br>(~\$67/mo)  | Moderate            | 2000  |
| C3         | 6 GB       | \$0.16/hr<br>(~\$134/mo) | Moderate            | 5000  |
| C4         | 13 GB      | \$0.21/hr<br>(~\$156/mo) | Moderate            | 10000 |
| C5         | 26 GB      | \$0.42/hr<br>(~\$312/mo) | High                | 15000 |
| C6         | 53 GB      | \$0.84/hr<br>(~\$625/mo) | Highest             | 20000 |

| CACHE NAME | CACHE SIZE | STANDARD                   | NETWORK PERFORMANCE | NUMBER OF CLIENT CONNECTIONS |
|------------|------------|----------------------------|---------------------|------------------------------|
| C0         | 250 MB     | \$0.05/hr<br>(~\$41/mo)    | Low                 | 256                          |
| C1         | 1 GB       | \$0.138/hr<br>(~\$103/mo)  | Low                 | 1000                         |
| C2         | 2.5 GB     | \$0.225/hr<br>(~\$167/mo)  | Moderate            | 2000                         |
| C3         | 6 GB       | \$0.45/hr<br>(~\$335/mo)   | Moderate            | 5000                         |
| C4         | 13 GB      | \$0.525/hr<br>(~\$391/mo)  | Moderate            | 10000                        |
| C5         | 26 GB      | \$1.05/hr<br>(~\$791/mo)   | High                | 15000                        |
| C6         | 53 GB      | \$2.10/hr<br>(~\$1,562/mo) | Highest             | 20000                        |
| P2         | 13 GB      | \$1.11/hr<br>(~\$826/mo)   | Moderate            | 15000                        |
| P3         | 26 GB      | \$2.21/hr<br>(~\$1,651/mo) | High                | 30000                        |

# Large number of configuration choices



Reasoning about cost-performance trade-off is hard!

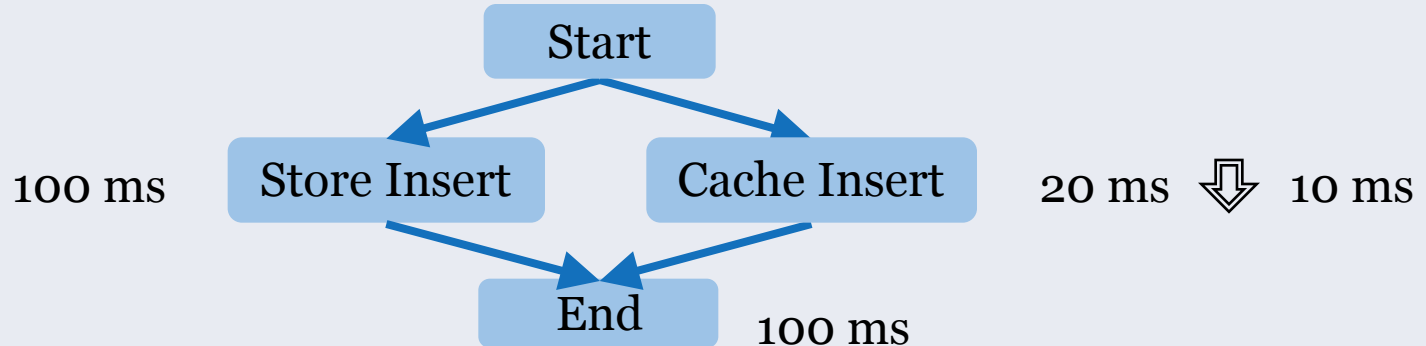
# Cost-Performance Trade-off

- Configuration does not directly map to performance

| INSTANCE | CORES | RAM     | DISK SIZES | PRICE                    |
|----------|-------|---------|------------|--------------------------|
| A0       | 1     | 0.75 GB | 19 GB      | \$0.02/hr<br>(~\$15/mo)  |
| A1       | 1     | 1.75 GB | 224 GB     | \$0.08/hr<br>(~\$60/mo)  |
| A2       | 2     | 3.5 GB  | 489 GB     | \$0.16/hr<br>(~\$115/mo) |
| A3       | 4     | 7 GB    | 999 GB     | \$0.32/hr<br>(~\$236/mo) |
| A4       | 8     | 14 GB   | 2,039 GB   | \$0.64/hr<br>(~\$476/mo) |

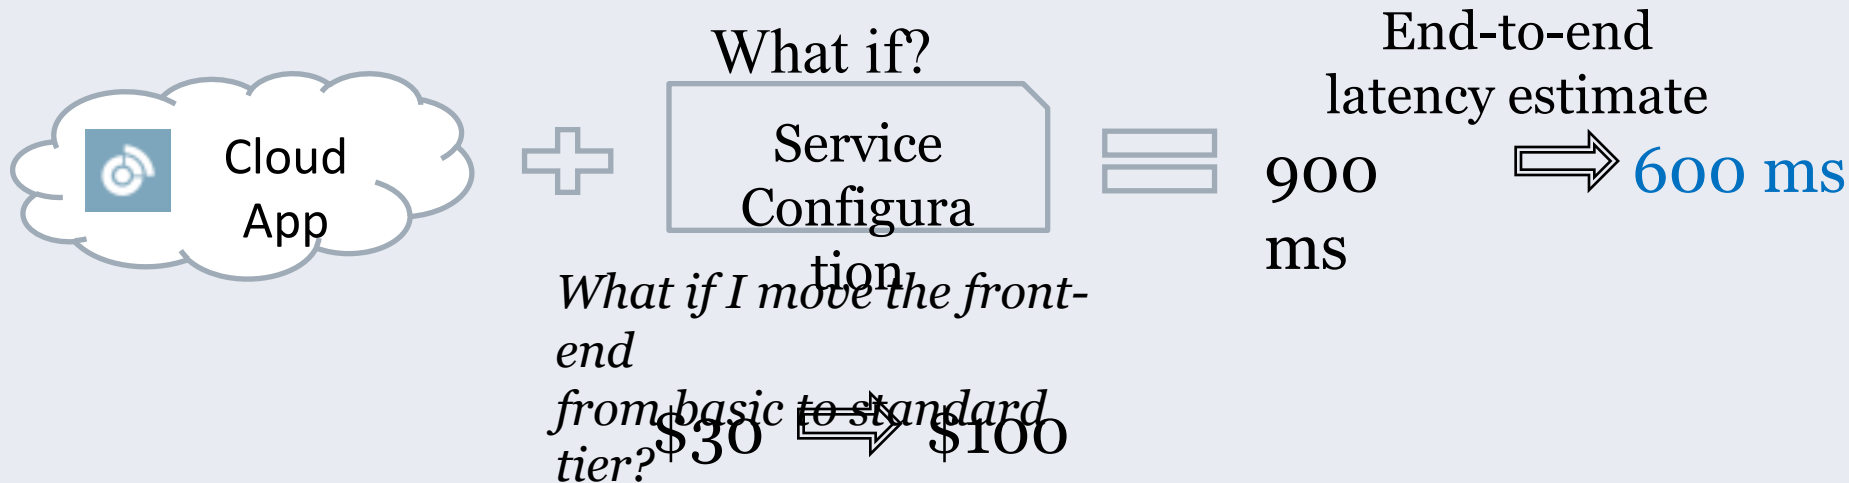


- End-to-end latency depends on application's causal dependency





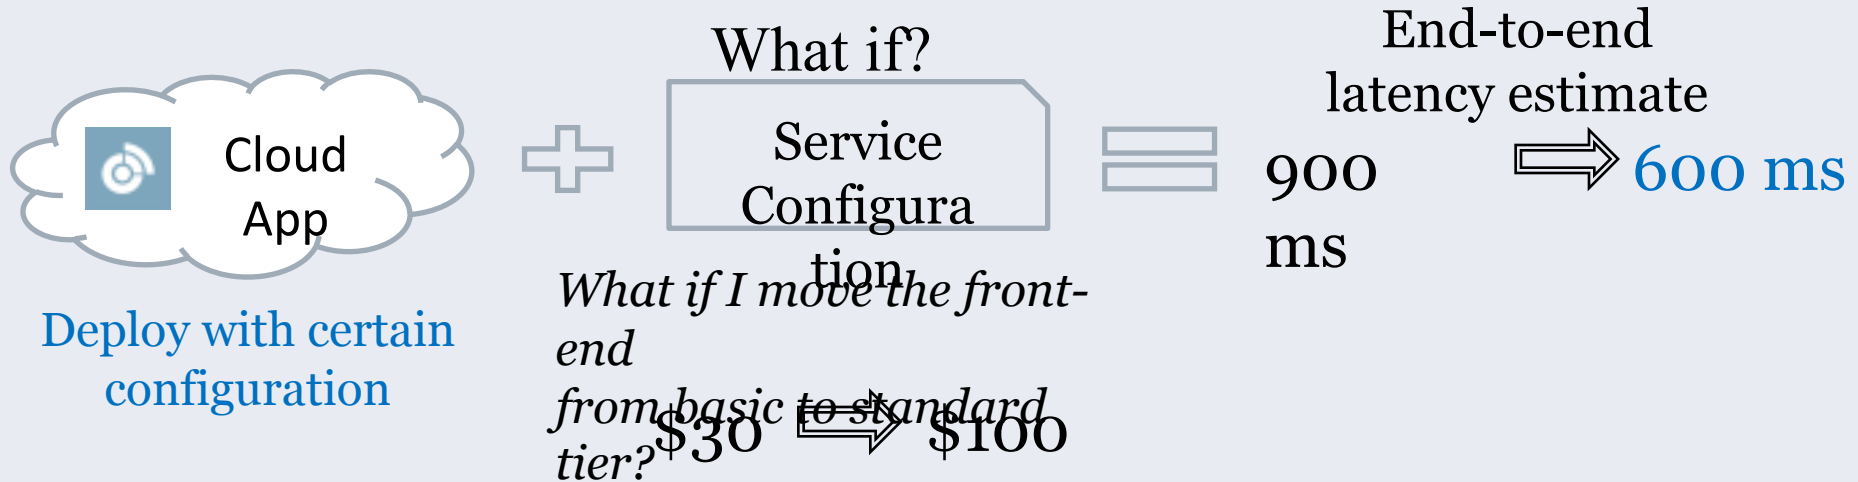
# “What-If” Analysis



Create a new deployment and measure performance

- Expensive
- Time consuming
- High overhead

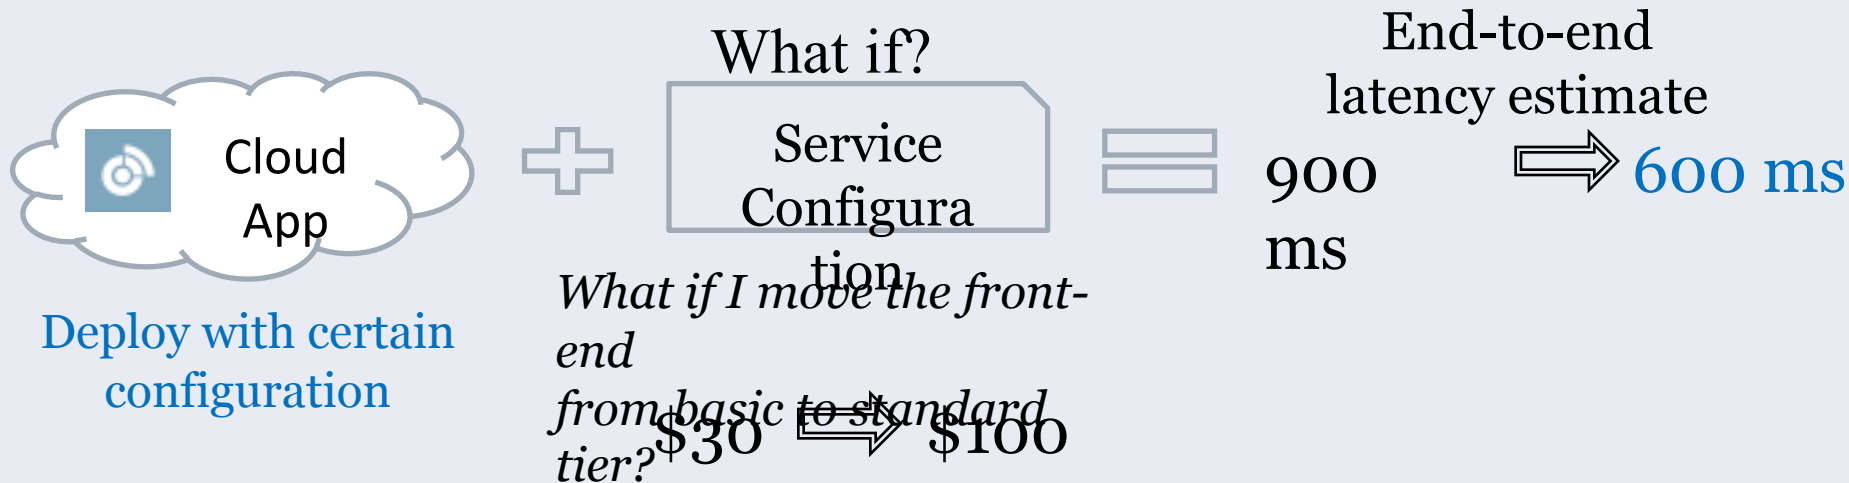
# WebPerf: “What-If” Analysis



Predict performance under hypothetical configurations

- Zero cost
- Near real-time
- Zero developer effort

# WebPerf: “What-If” Analysis



Predict performance under hypothetical configurations

- Zero cost
- Near real-time
- Zero developer effort

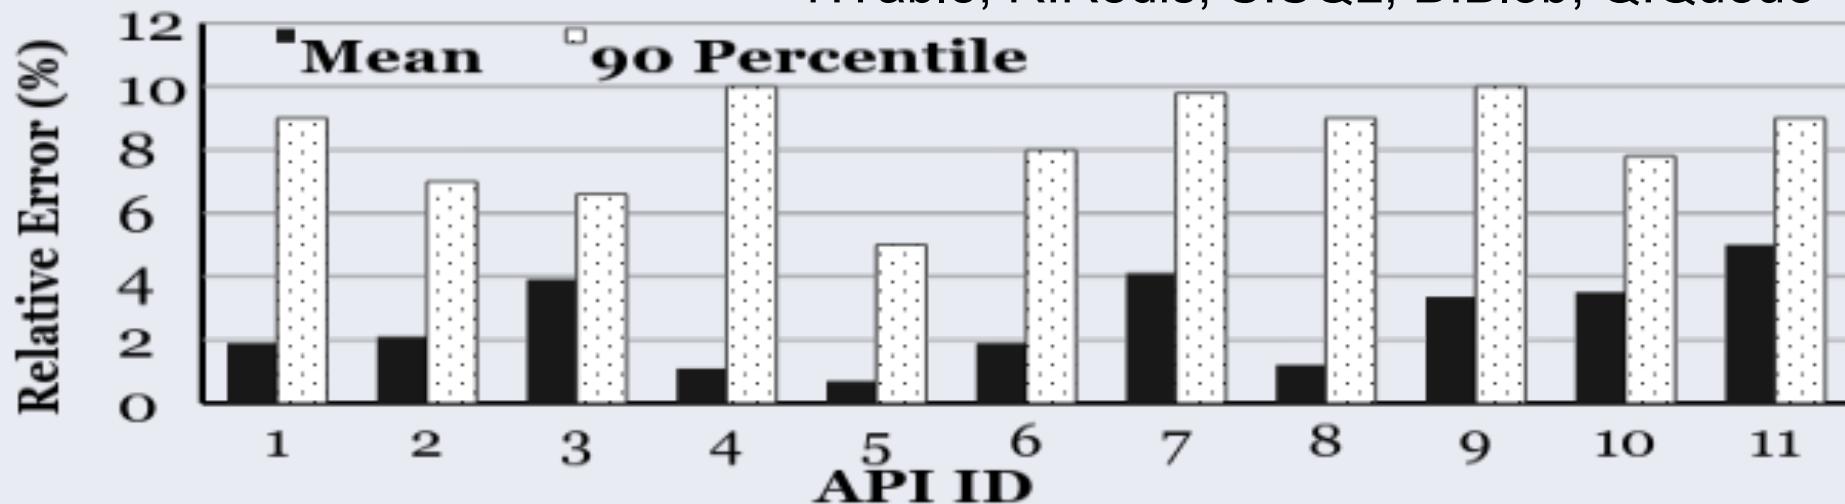
# WebPerf: Key Insights

- Offline, application-independent profiling is useful
  - Modern cloud apps are built using existing services (PaaS)
  - Individual services have predictable performance
    - S3, Azure Table Storage, Dynamo DB, DocumentDB, ...
  - Services are co-located inside the same datacenter
    - Tighter latency distribution
- Causal dependency within application is independent of the what-if scenarios we consider

# Application-Independent Profiling

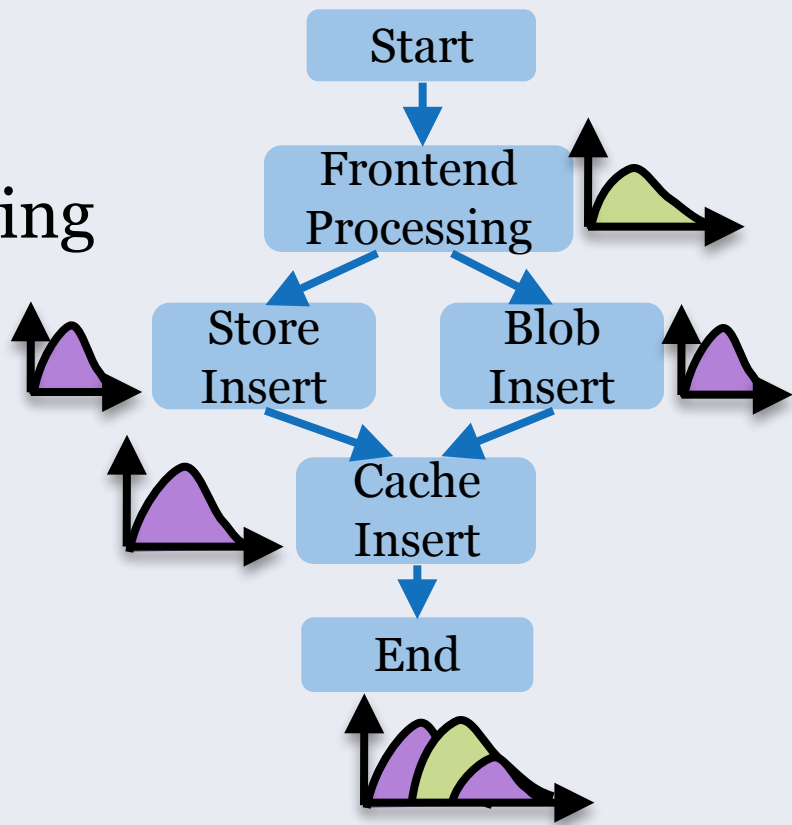
|   |                              |   |                           |    |                  |
|---|------------------------------|---|---------------------------|----|------------------|
| 1 | Delete( <b>Async</b> ) (T) ★ | 5 | ExecuteQuerySegmented (T) | 9  | ToList (S)       |
| 2 | UploadFromStream (B)         | 6 | SortedSetRangeByValue (R) | 10 | Send (R)         |
| 3 | AddMessage(Q)                | 7 | StringGet (R)             | 11 | ReadAsString (B) |
| 4 | Execute (T)                  | 8 | SaveChanges (S)           |    |                  |

★ T:Table, R:Redis, S:SQL, B:Blob, Q:Queue



# WebPerf Design

- Dependency graph extraction
- Application-independent profiling
- Baseline latency estimation
- Latency prediction





# Dependency Graph Extraction

- Design Goals
  - Accurate
  - Real-time with minimal data collection
  - Zero developer effort
  - No modifications to the platform
  - Low overhead
- Automatic Binary Instrumentation
- Modern cloud applications are highly asynchronous
  - Task Asynchronous Programming Pattern

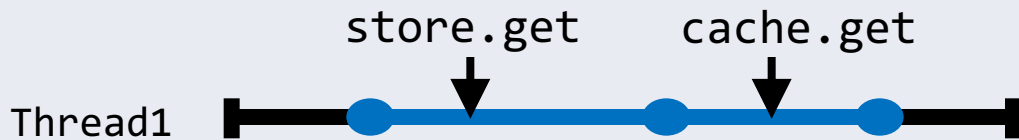


# Task Asynchronous Pattern (TAP)

- Asynchronous operations with a synchronous programming pattern
- Increasingly popular for writing cloud applications
- Supported by many major languages
  - C#, Java, Python, Javascript
- Most Azure services support TAP as the **only** mechanism for doing asynchronous I/O
  - AWS also provides TAP APIs for .NET

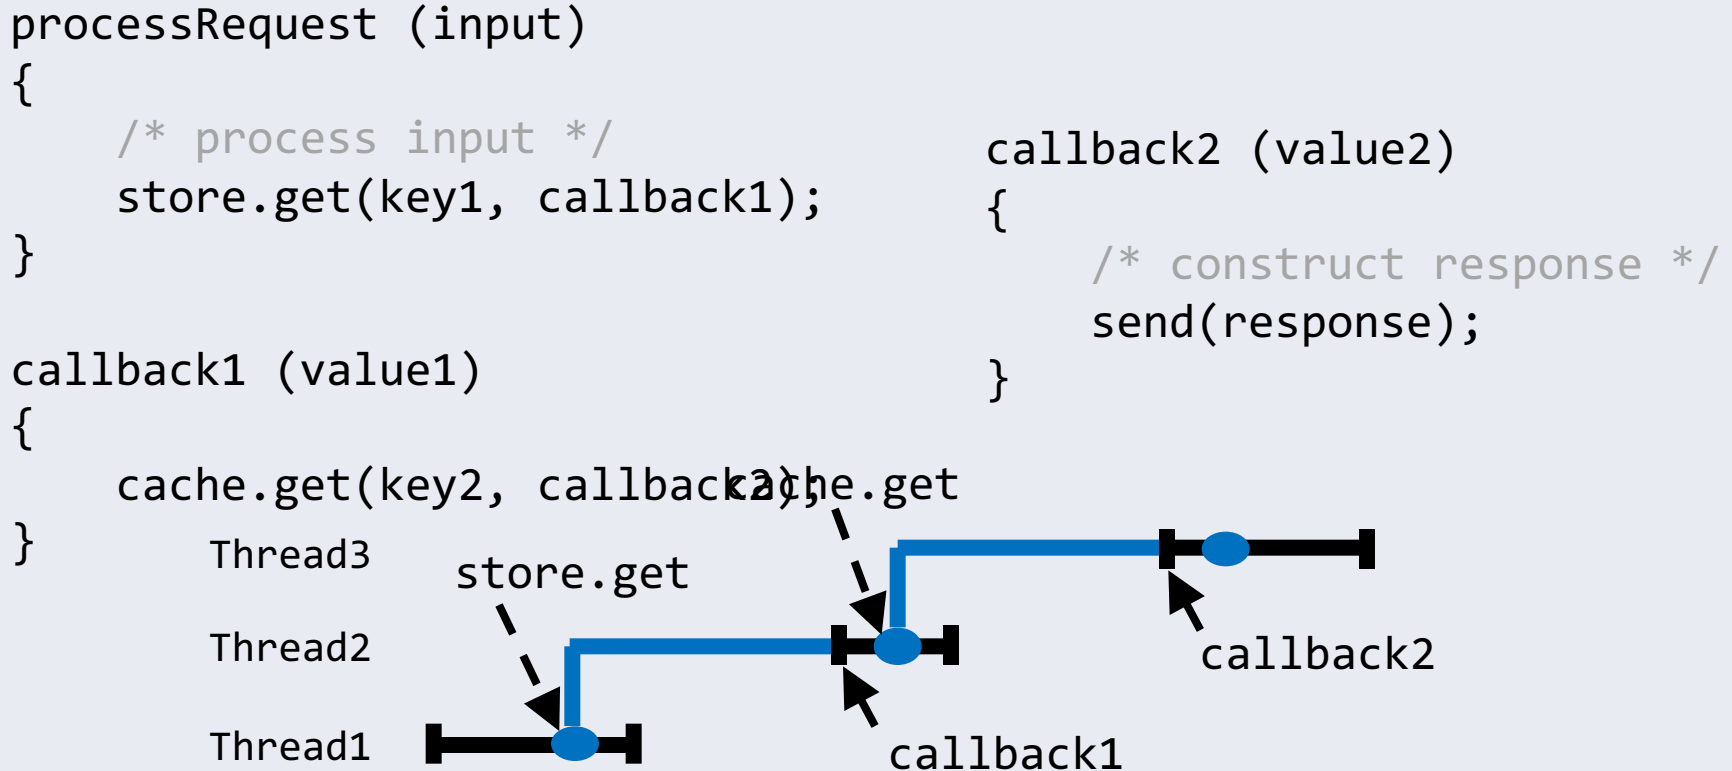
# Synchronous Programming

```
processRequest (input)
{
    /* process input */
    value1 = store.get(key1);
    value2 = cache.get(key2);
    /* construct response */
    return response;
}
```




Blocking I/O limits server throughput

# Asynchronous Programming Model (APM)

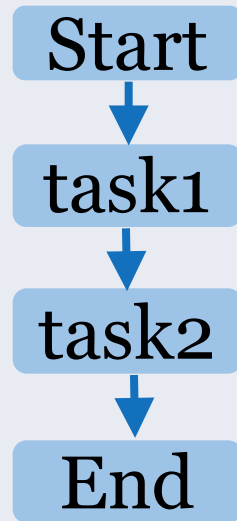


# Task Asynchronous Pattern (TAP)

```
async processRequest (input)
{
    /* process input */
    task1 = store.get(key1);
    value1 = await task1;
    task2 = cache.get(key2);
    value2 = await task2;
    /* construct response */
    return response;
}
```



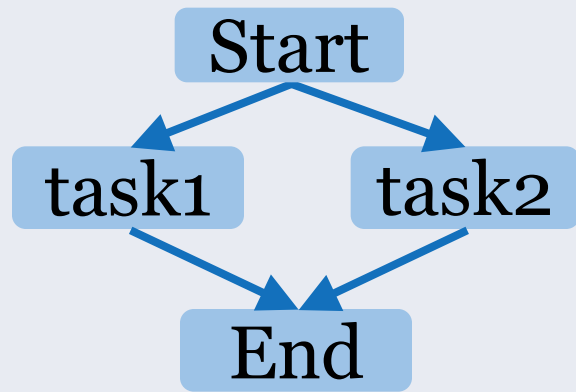
## Dependency Graph



# Task Asynchronous Pattern (TAP)

```
async processRequest (input)
{
    /* process input */
    task1 = store.get(key1);
    task2 = cache.get(key2);
    value1 = await task1;
    value2 = await task2;
    /* construct response */
    return response;
}
```

## Dependency Graph

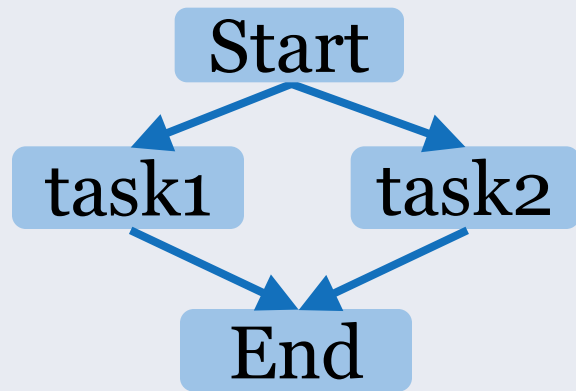


# Task Asynchronous Pattern (TAP)

```
async processRequest (input)
{
    /* process input */
    task1 = store.get(key1);
    task2 = cache.get(key2);
    value1, value2 = await
        Task.WhenAll(task1, task2);
    /* construct response */
    return response;
}
```

**WhenAll:** Continue only when all tasks finish

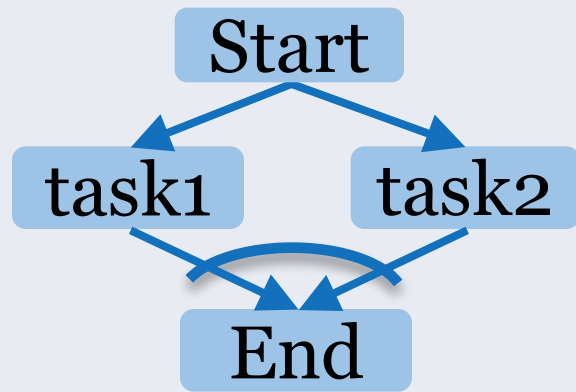
## Dependency Graph



# Task Asynchronous Pattern (TAP)

```
async processRequest (input)
{
    /* process input */
    task1 = store.get(key1);
    task2 = cache.get(key2);
    value = await
        Task.WhenAny(task1, task2);
    /* construct response */
    return response;
}
```

## Dependency Graph



**WhenAny:** Continue after *any one* of tasks finishes

# Automatic Binary Instrumentation

```
async processRequest (input)
```

```
{
```

```
-1
```

```
/* process input */
```

```
0
```

```
task1 = store.get(key1);  
value1 = await task1;
```

```
1
```

```
task2 = cache.get(key2);  
value2 = await task2;  
/* construct response */  
return response;
```

```
} Instrument state machine
```

Track tasks and continuations



```
-1
```

continuation

```
0
```

continuation

```
1
```

```
class processRequest__
```

```
{
```

```
string input;
```

```
AsyncTaskMethodBuilder builder;
```

```
string key1, key2, response;
```

```
int asyncId = -1;
```

```
public void MoveNext()
```

```
{
```

```
asyncId = Tracker.AsyncStart(asyncId);
```

```
Tracker.StateStart(asyncId);
```

```
switch (state)
```

```
{
```

```
case -1:
```

```
state = 0;
```

```
/* process input */
```

```
var task1 = store.get(key1);
```

```
Tracker.TaskStart(task1, asyncId);
```

```
builder.Completed(task1.Awaiter, this);
```

```
Tracker.Await(task1, asyncId);
```

```
case 0:
```

```
state = 1;
```

```
var task2 = cache.get(key1);
```

```
Tracker.TaskStart(task2, asyncId);
```

```
builder.Completed(task2.Awaiter, this);
```

```
Tracker.Await(task2, asyncId);
```

```
case 1:
```

```
/* construct response */
```



# Automatic Binary Instrumentation

```
async processRequest (input)
```

```
{
```

```
-1
```

```
/* process input */
```

```
0
```

```
task1 = store.get(key1);
```

```
value1 = await task1;
```

```
1
```

```
task2 = cache.get(key2);
```

```
value2 = await task2;
```

```
/* construct response */
```

```
return response;
```

```
} Instrument state machine
```



## Dependency Graph

Start



task1



task2



End

Track tasks and continuations

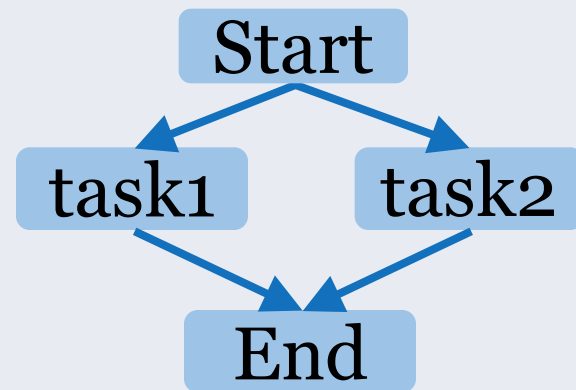
# Automatic Binary Instrumentation

- **Tracking async state machines**
  - Monitor task start and completion
  - Track state machine transitions
- **Tracking pull-based continuations**
  - Link tasks to corresponding awaits
  - Link awaits to continuations
- **Tracking synchronization points**
  - Track WhenAll, WhenAny, cascaded task dependencies
- **Keeping the overhead low**
  - Instrument APIs with known signatures
  - Instrument only leaf tasks

# Dependency graph extraction

- Highly accurate
- Real-time
- Zero developer effort
- Extremely low overhead

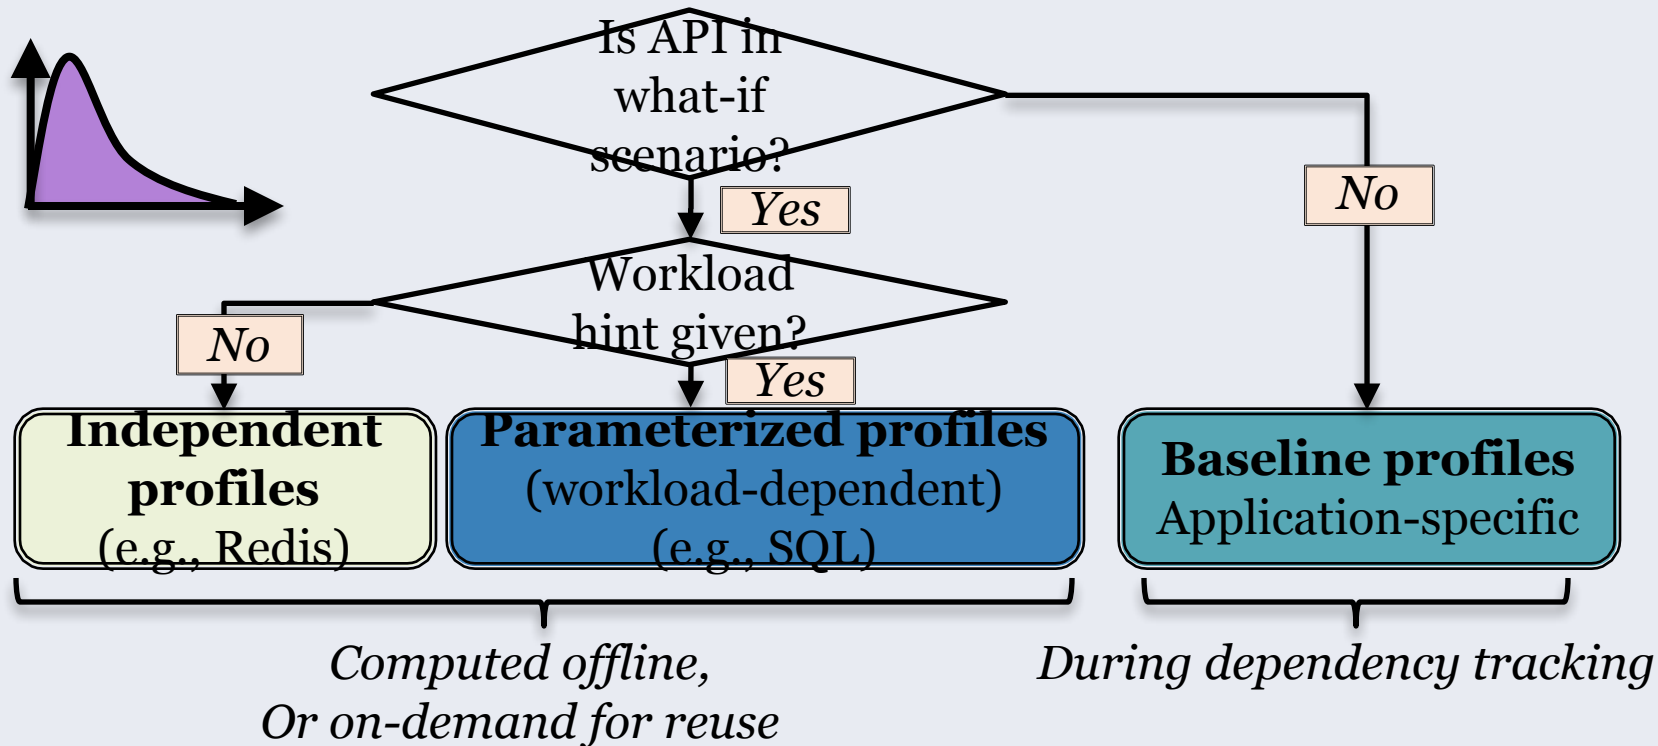
## Dependency Graph



[Some Result]

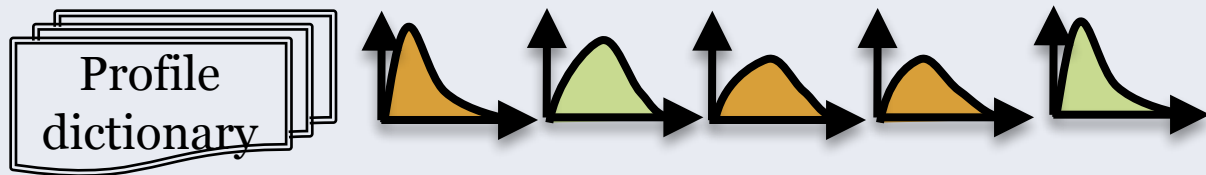
# API Profiling

- A profile of a cloud API is a distribution of its latency



# API Profiling

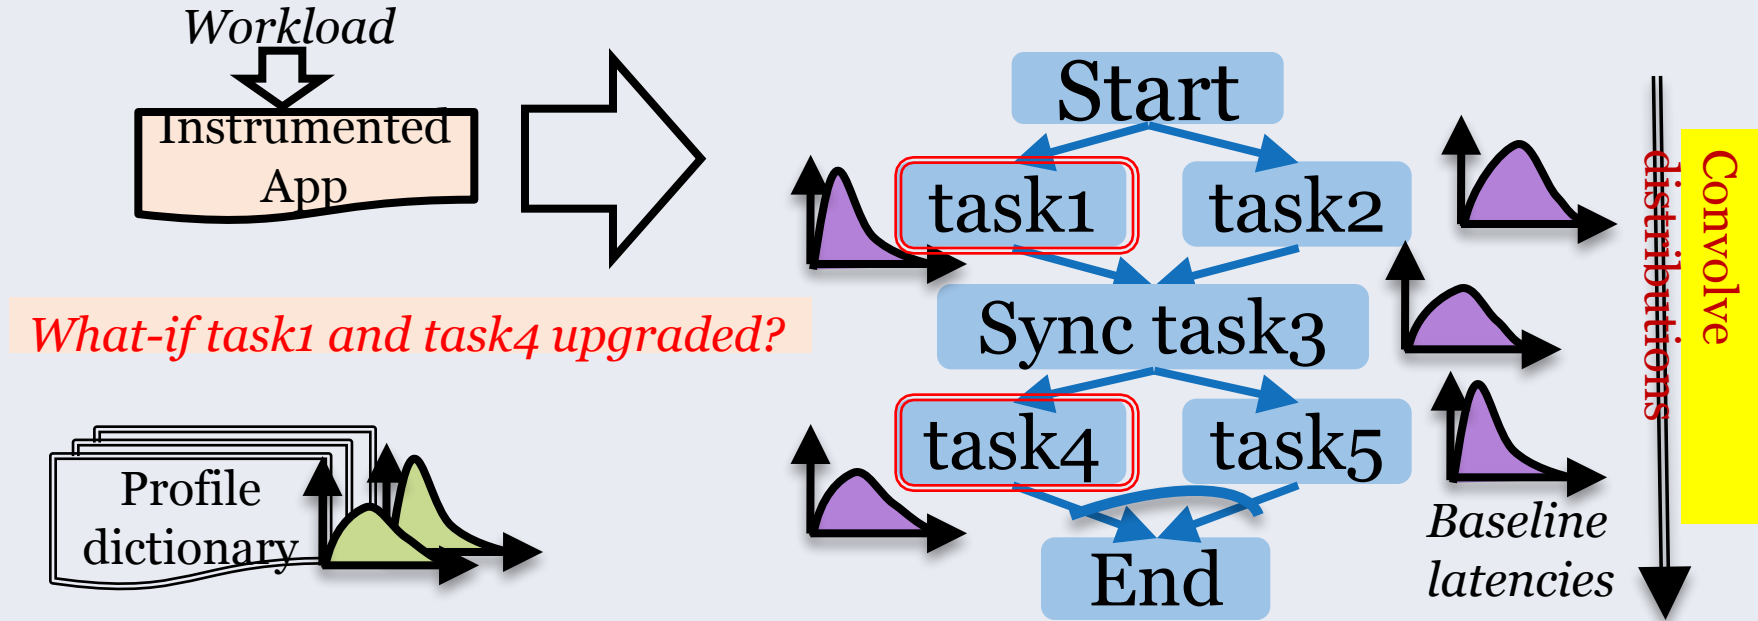
- WebPerf builds profiles offline and maintains in a dictionary



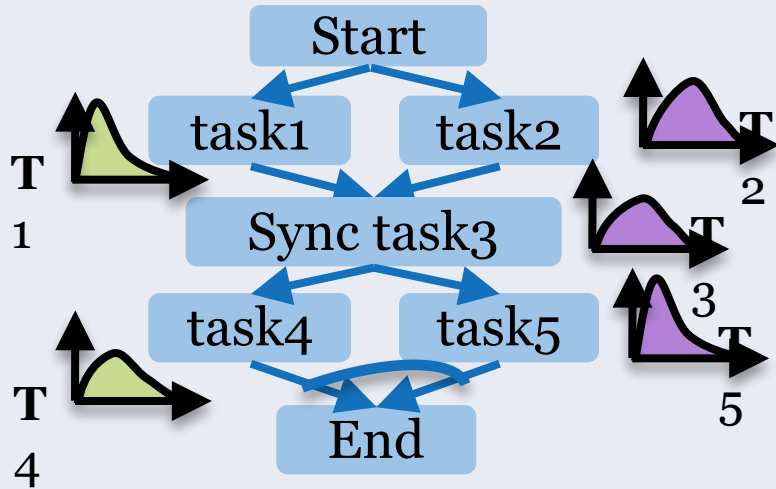
- Starts with common profiles, and builds additional profiles on-demand and reuses them
- **Optimal profiling**: to minimize measurement costs (details in paper)

# What-If Engine

- Predicts cloud latency under a given what-if scenario



# Convolving distributions

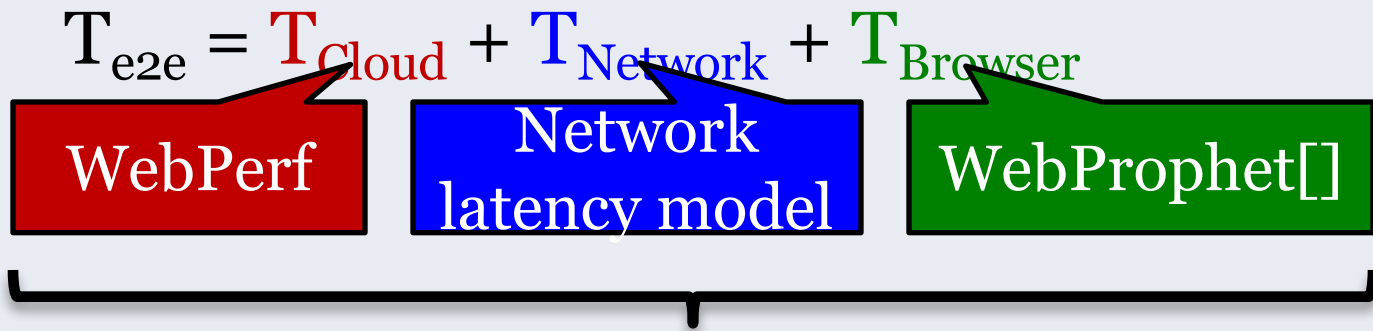


## Bottom-up evaluation:

- When **All**: Prob**Max**(t1, t2, ...)
- When **Any**: Prob**Min**(t1, t2, ...)
- When **Done**: Prob**Add**(t1, t2, ...)

$$\mathbf{T}_{\{s2e\}} = \text{ProbAdd}(\text{ProbMax}(\mathbf{T}_1, \mathbf{T}_2), \mathbf{T}_3, \text{ProbMin}(\mathbf{T}_4, \mathbf{T}_5))$$

# End-to-end Latency Prediction



Combine using Monte-Carlo simulation

- Details in paper

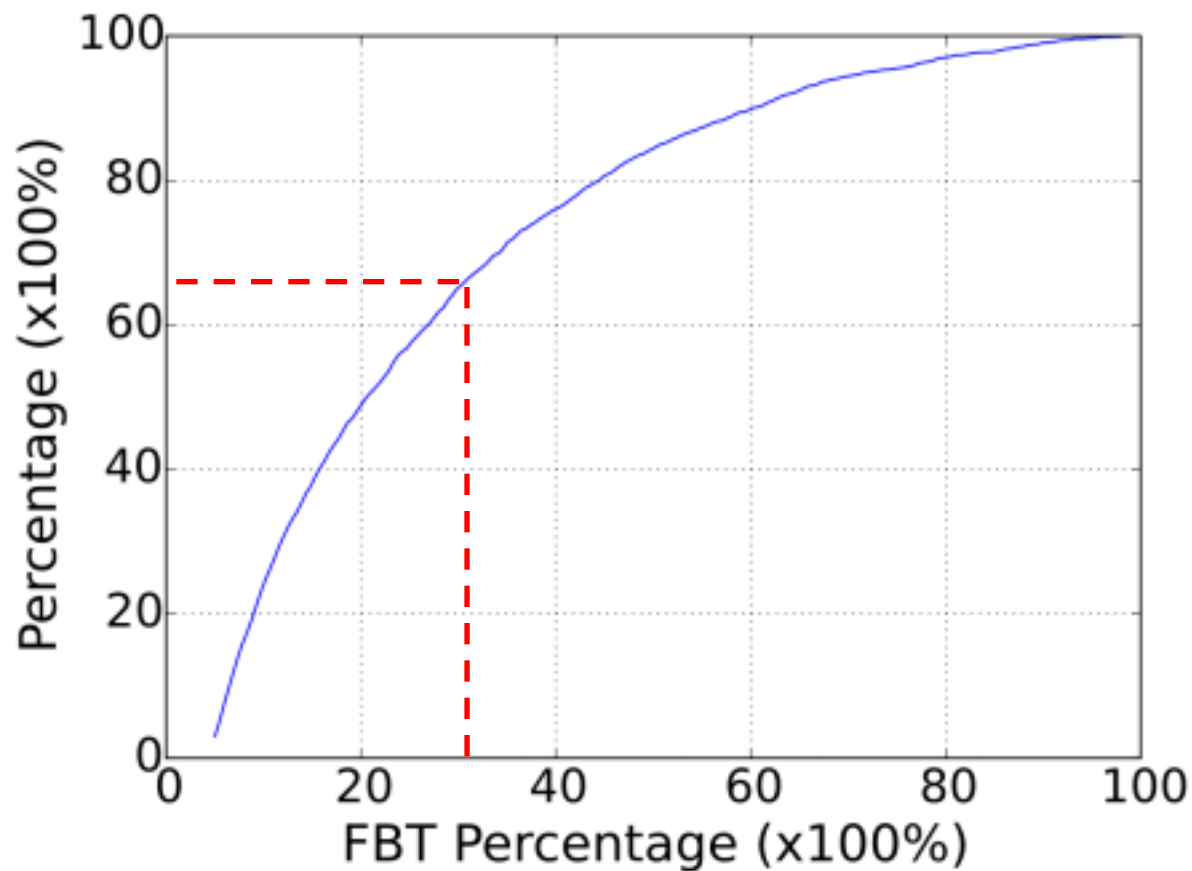


# WebPerf Evaluation

- Six 3<sup>rd</sup> party applications and six scenarios

| Application     | Azure services used                                   | Average I/O Calls |
|-----------------|---|-------------------|
| SocialForum     | Blob storage, Redis cache, Service bus, Search, Table | 116               |
| SmartStore.Net  | SQL   | 41                |
| ContosoAds      | Blob storage, Queue, SQL, Search                      | 56                |
| EmailSubscriber | Blob storage, Queue, Table                            | 26                |
| ContactManager  | Blob storage, SQL                                     | 8                 |
| CourseManager   | Blob storage, SQL                                     | 44                |

# CDF of First Byte Time Percentage

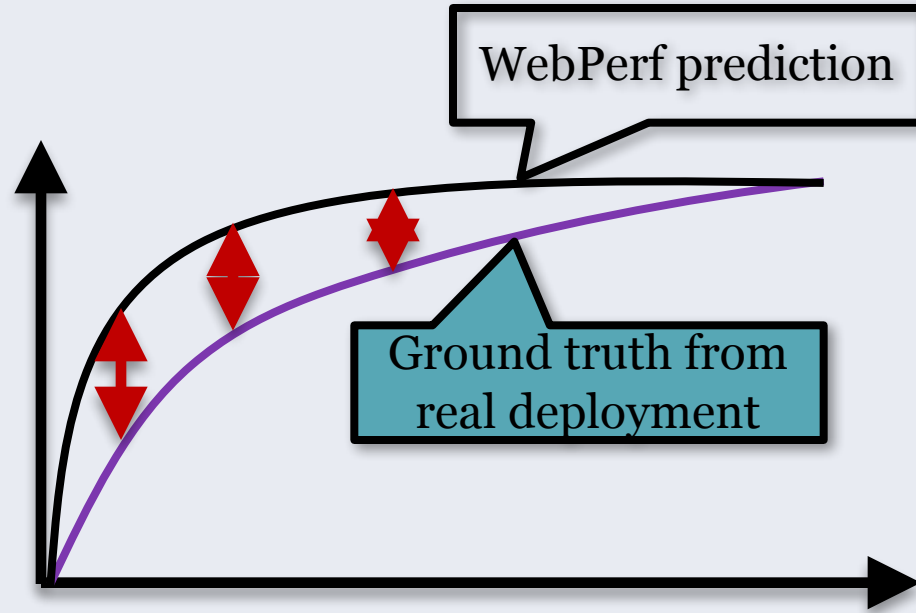


# WebPerf Evaluation

- Six 3<sup>rd</sup> party applications and **six scenarios**

| What-if scenario   | Example  |
|--|--|
| <b>Tier:</b> A resource X is upgraded to tier Y                                      | X = A Redis cache, Y = a standard tier (from a basic tier) |
| <b>Load:</b> X concurrent requests to resource Y                                     | X = 100 , Y = the application or a SQL database            |
| <b>Interference:</b> CPU and/or memory pressure, from collocated applications, of X% | X = 50% CPU, 80% memory                                    |
| <b>Location:</b> A resource X is deployed at location Y                              | X = A Redis Cache or a front end, Y = Singapore            |
| <b>Failure:</b> An instance of a replicated resource X fails                         | X = A replicated front-end or SQL database                 |

# WebPerf Evaluation



- Metric: distribution of **relative errors**

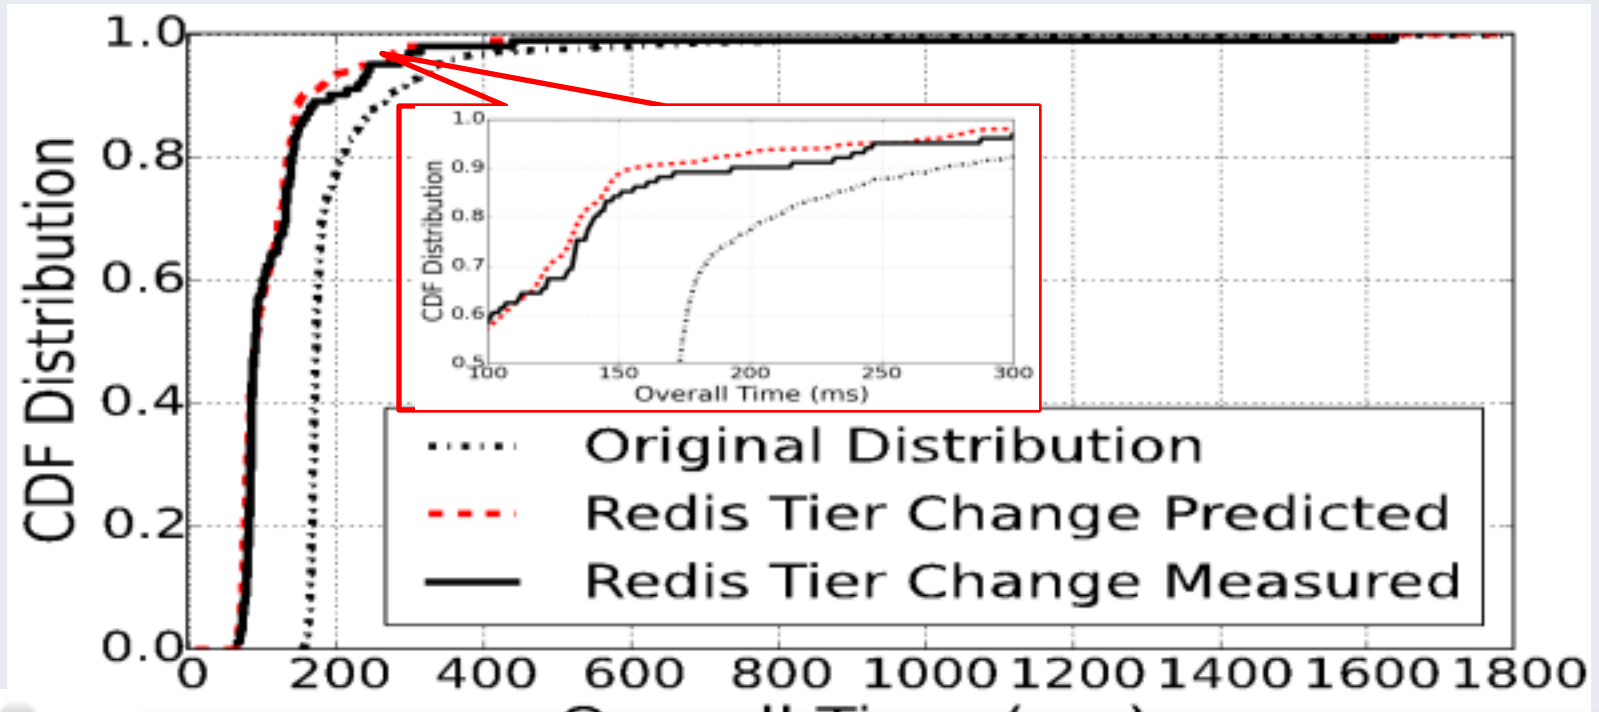
## Cache

## Redis Cache

# Underspecified Configuration Dimensions

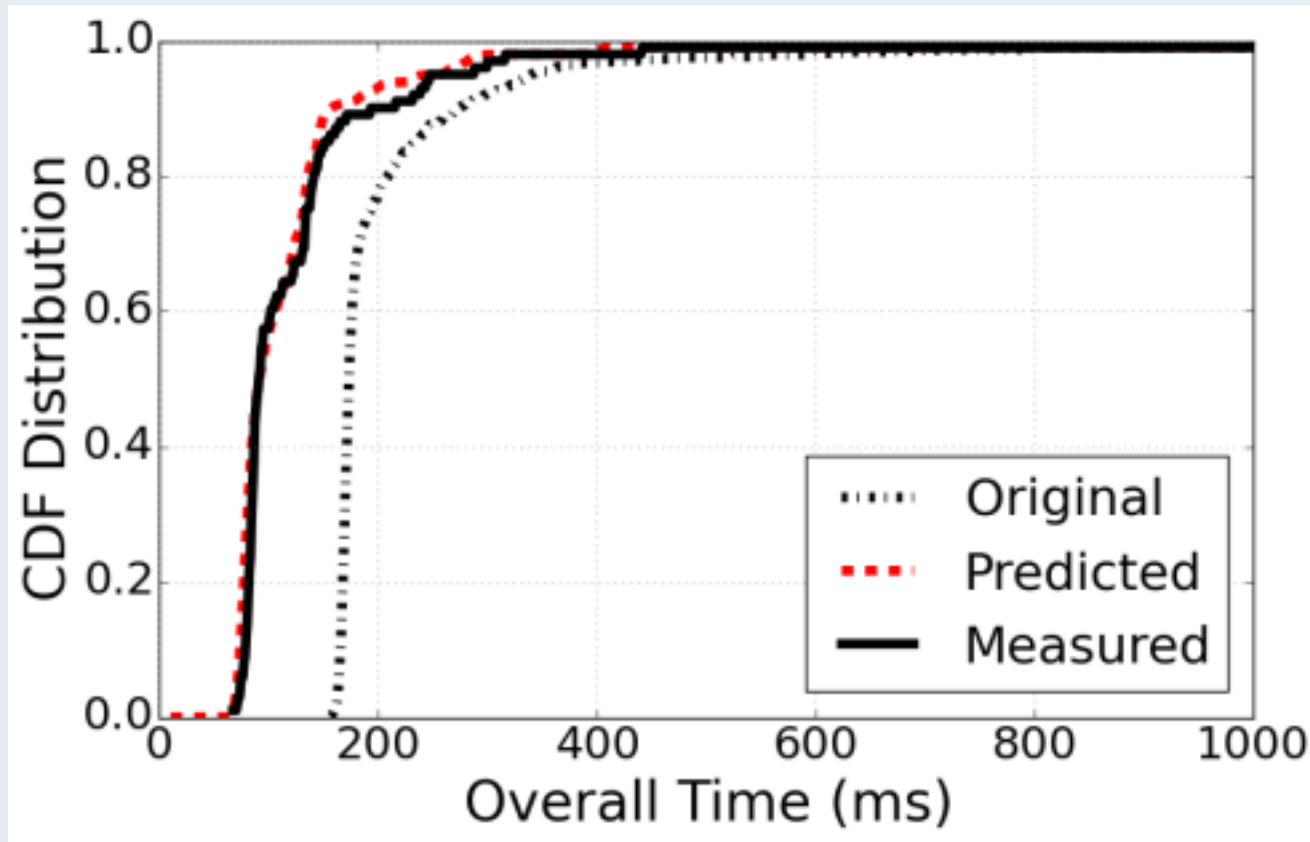
| CACHE NAME | CACHE SIZE | BASIC                    | NETWORK PERFORMANCE | NUMBER OF CLIENT CONNECTIONS |
|------------|------------|--------------------------|---------------------|------------------------------|
| C0         | 250 MB     | \$0.022/hr<br>(~\$16/mo) | Low                 | 256                          |
| C1         | 1 GB       | \$0.055/hr<br>(~\$41/mo) | Low                 | 1000                         |
| C2         | 2.5 GB     | \$0.09/hr<br>(~\$67/mo)  | Moderate            | 2000                         |
| C3         | 6 GB       | \$0.18/hr<br>(~\$134/mo) | Moderate            | 5000                         |
| C4         | 13 GB      | \$0.21/hr<br>(~\$156/mo) | Moderate            | 10000                        |
| C5         | 26 GB      | \$0.42/hr<br>(~\$312/mo) | High                | 15000                        |
| C6         | 53 GB      | \$0.84/hr<br>(~\$625/mo) | Highest             | 20000                        |

# What-if the Redis cache is upgraded from the original standard C0 to Standard C2 tier?

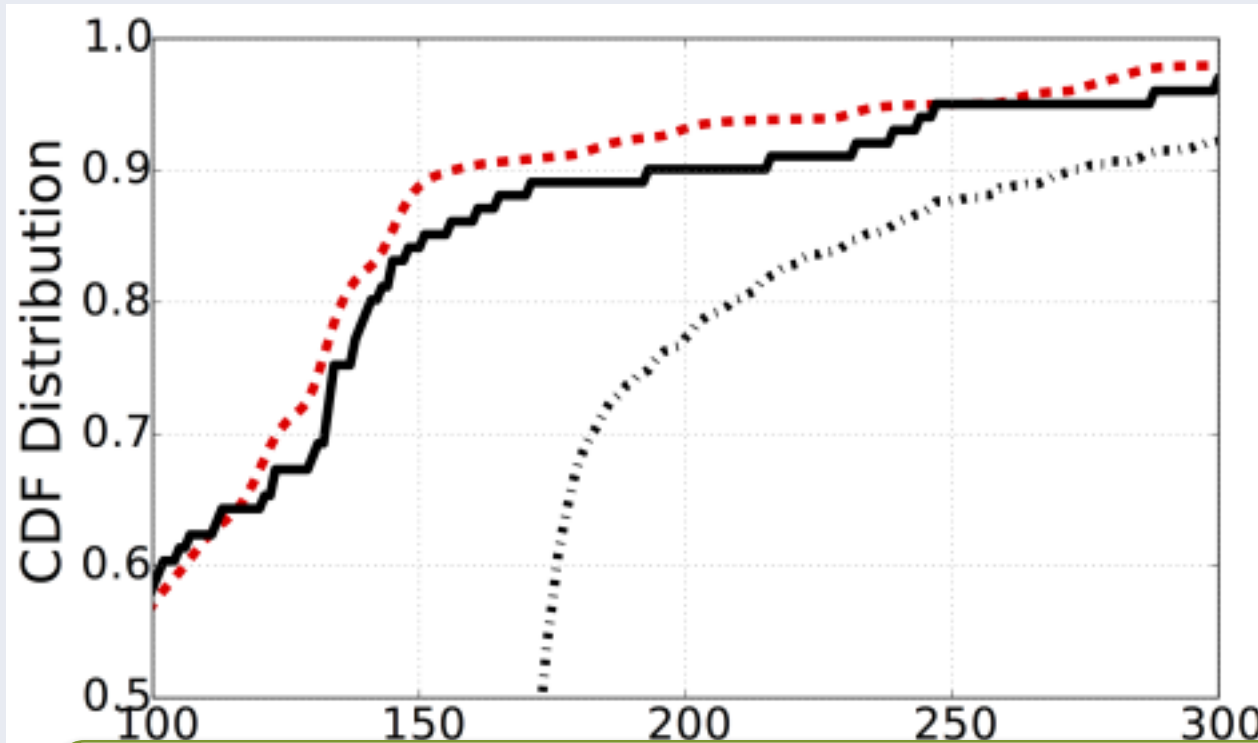


Maximum cloud side latency prediction error is only 5%

# What-if the Redis cache is upgraded from the original standard C0 to Standard C2 tier?



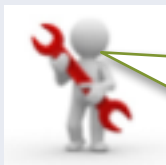
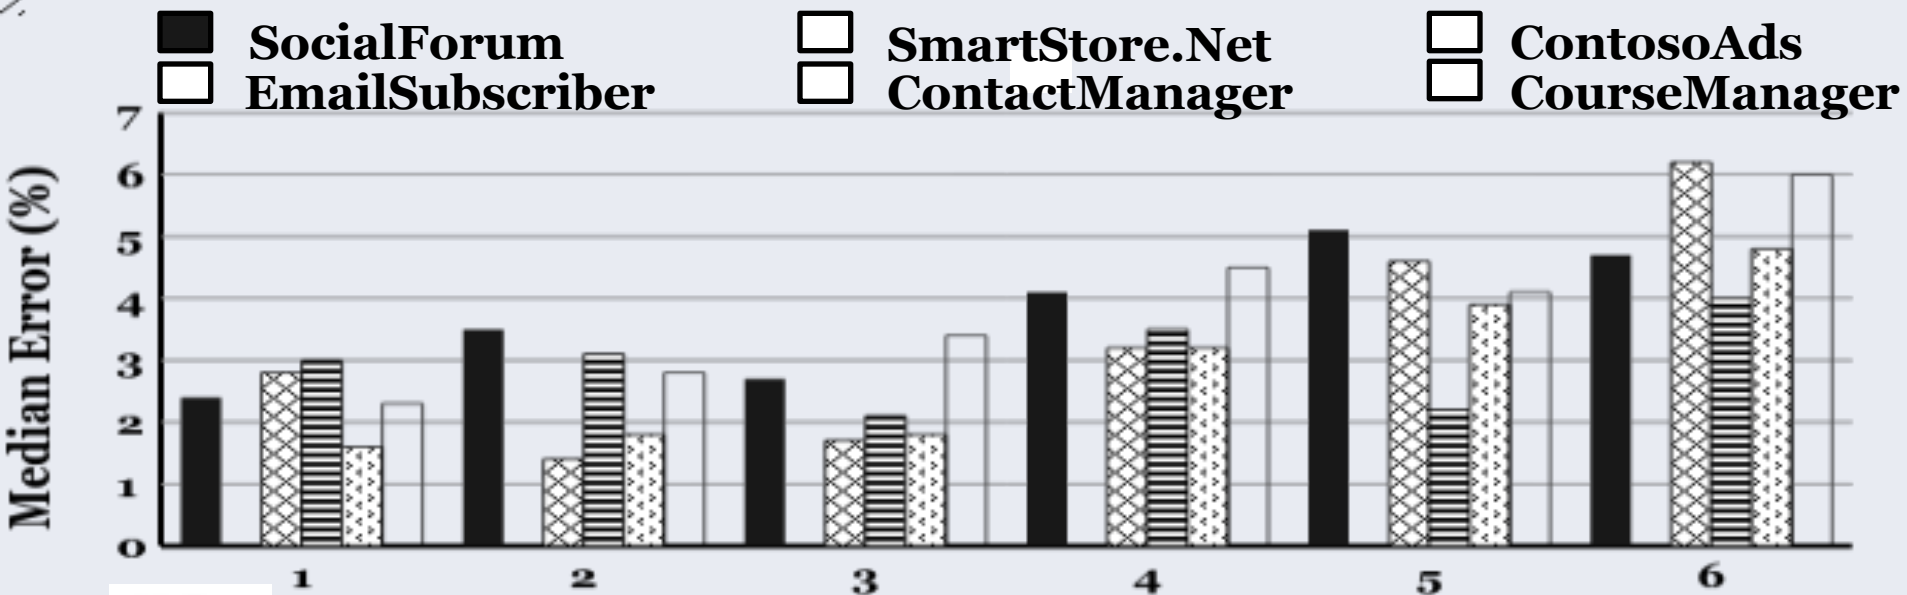
# What-if the Redis cache is upgraded from the original standard C0 to Standard C2 tier?



Maximum cloud side latency prediction error is only 5%

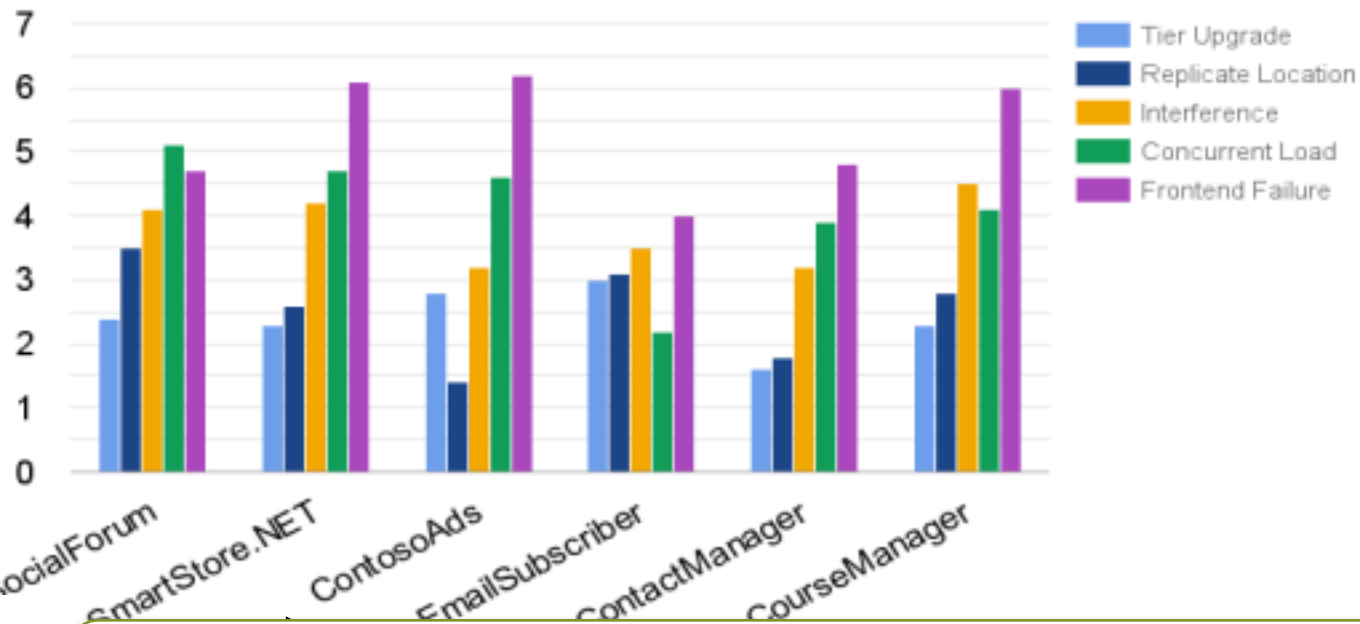


# Performance for Six Applications



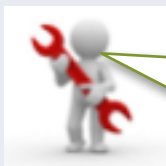
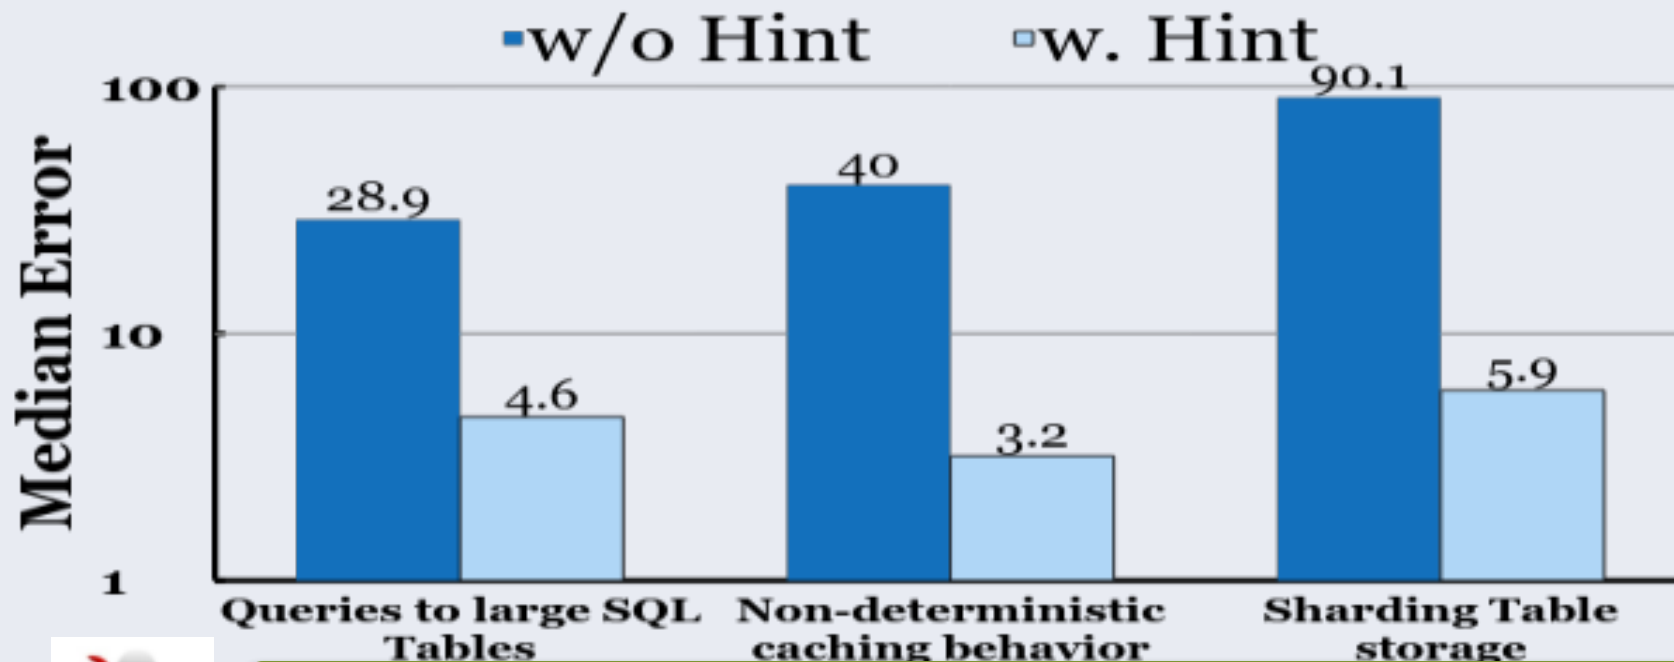
Median prediction error for cloud side latency is  $< 7\%$

# Performance for Six Applications



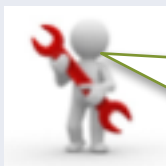
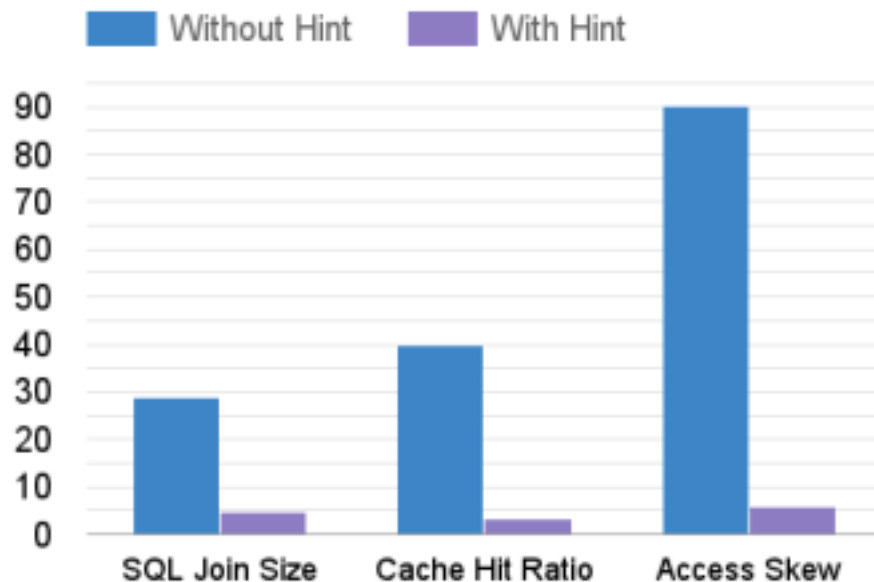
Median prediction error for cloud side latency is  $< 7\%$

# Workload Hints



Workload hints can bring order of magnitude accuracy improvement

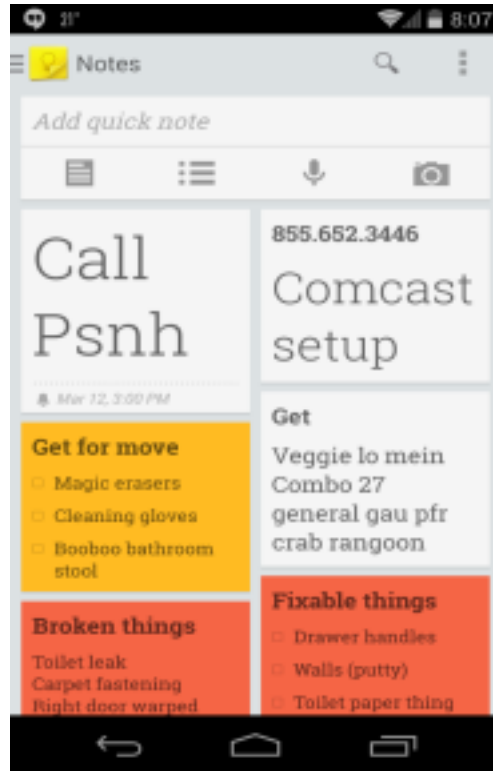
# Workload Hints



Workload hints can bring order of magnitude accuracy improvement

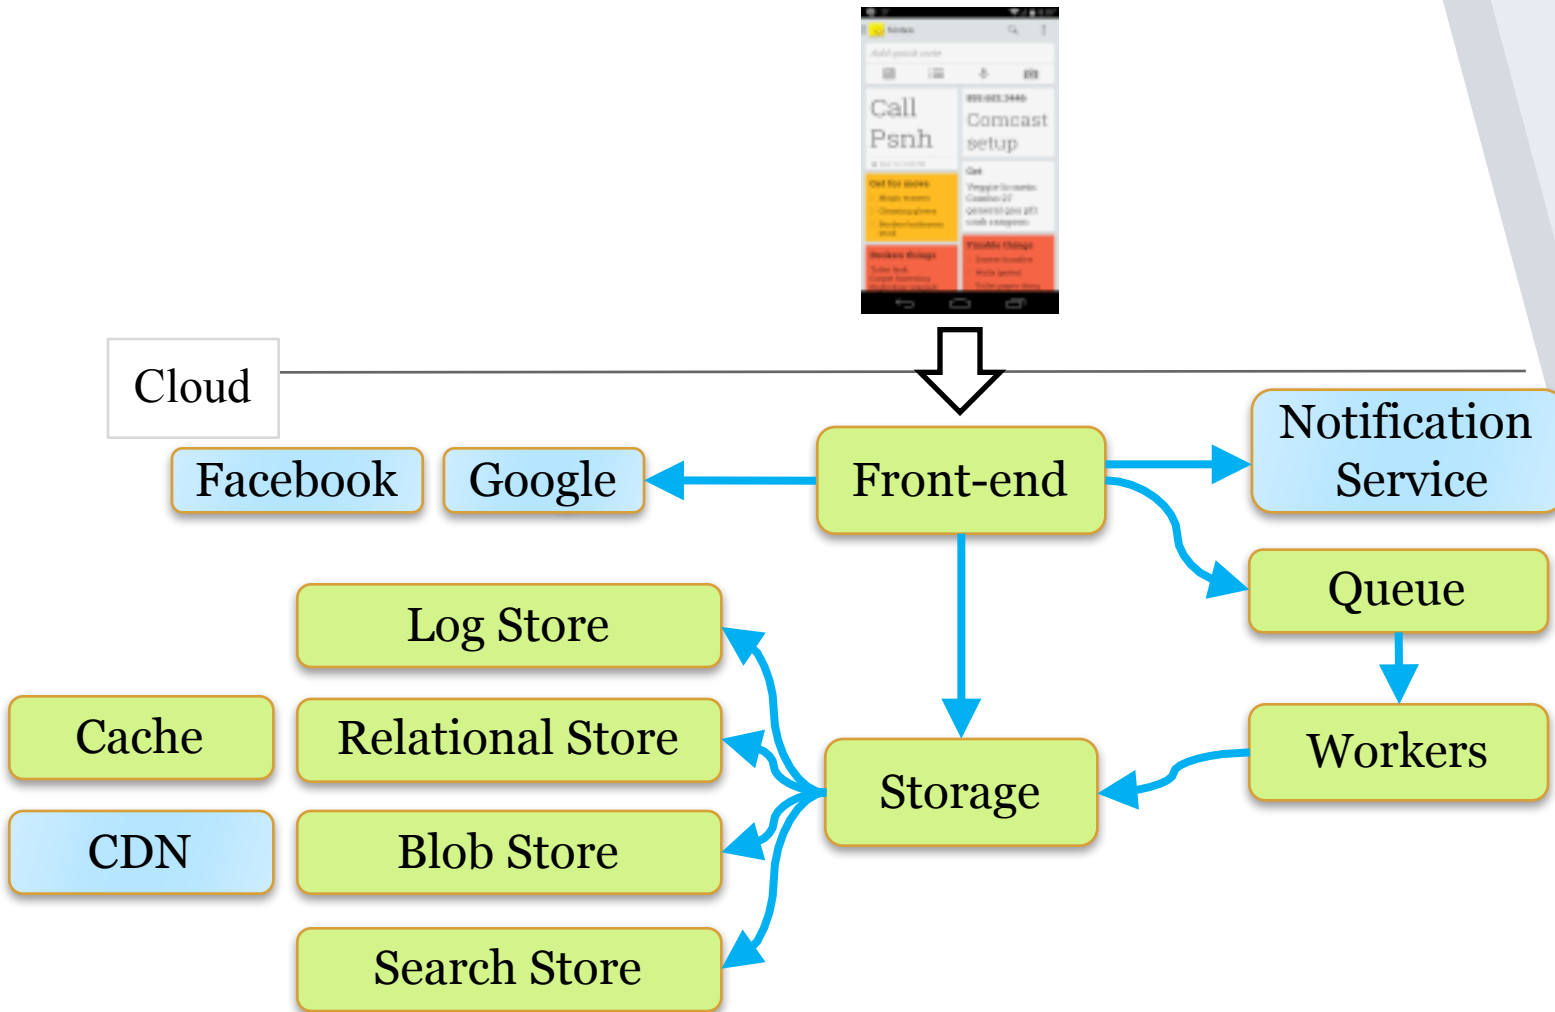
# Other findings

- Performance of many applications and scenarios can be predicted reasonably well
  - Thanks to cloud provider's SLAs
- Harder cases
  - Workload-dependent performance: hints help
  - High-variance profiles: prediction has high variance
  - Non-deterministic control flow (e.g., cache hit/miss):
    - Separate prediction for each control flow
  - Hard-to-profile APIs (e.g., SQL query with join)
    - Poor prediction



Behind apps, several  
distributed,  
asynchronous  
components

# Modern Cloud Applications are Complex





Hard to reason about the  
— ~~performance~~ of cloud-hosted

Cloud-side  
Latency

Web apps





Hard to reason about the  
cost/performance tradeoffs of  
different configurations

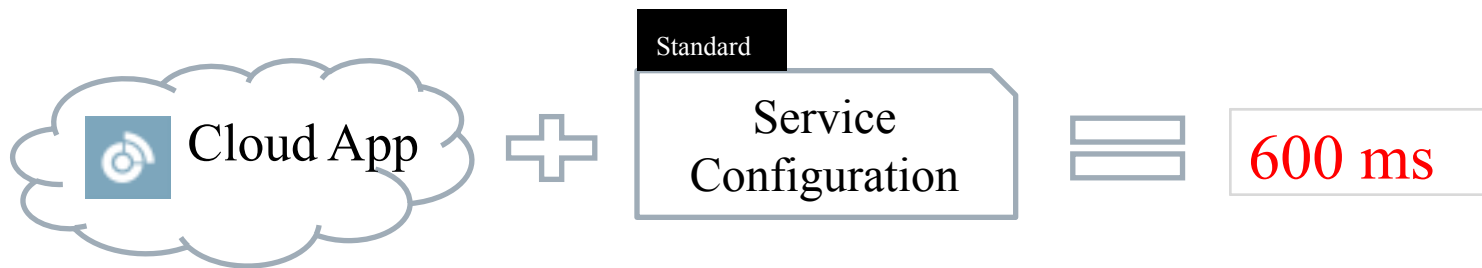
# Workload Dependence

| INSTANCE | CORES | RAM     | DISK SIZES | PRICE                    |
|----------|-------|---------|------------|--------------------------|
| A0       | 1     | 0.75 GB | 19 GB      | \$0.02/hr<br>(~\$15/mo)  |
| A1       | 1     | 1.75 GB | 224 GB     | \$0.08/hr<br>(~\$60/mo)  |
| A2       | 2     | 3.5 GB  | 489 GB     | \$0.16/hr<br>(~\$119/mo) |
| A3       | 4     | 7 GB    | 999 GB     | \$0.32/hr<br>(~\$238/mo) |
| A4       | 8     | 14 GB   | 2,039 GB   | \$0.64/hr<br>(~\$476/mo) |



What if?

*What if I move the front-end  
from basic to standard tier?*



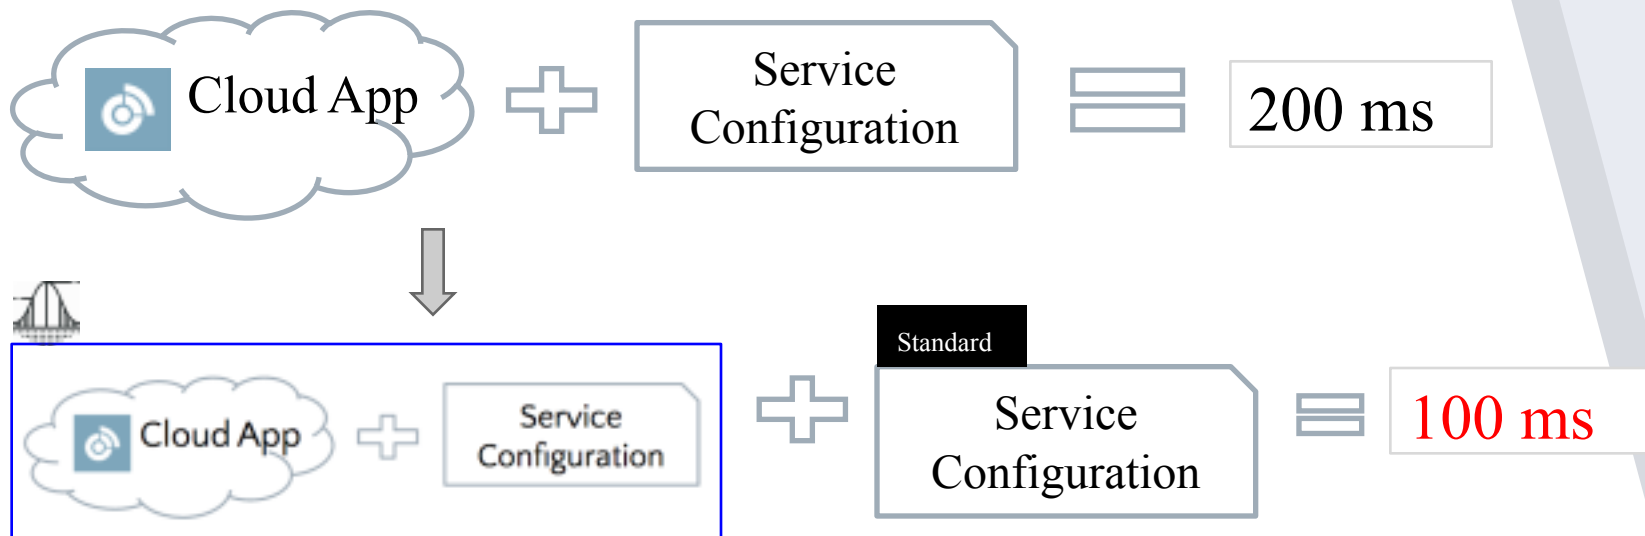
➤ Expensive

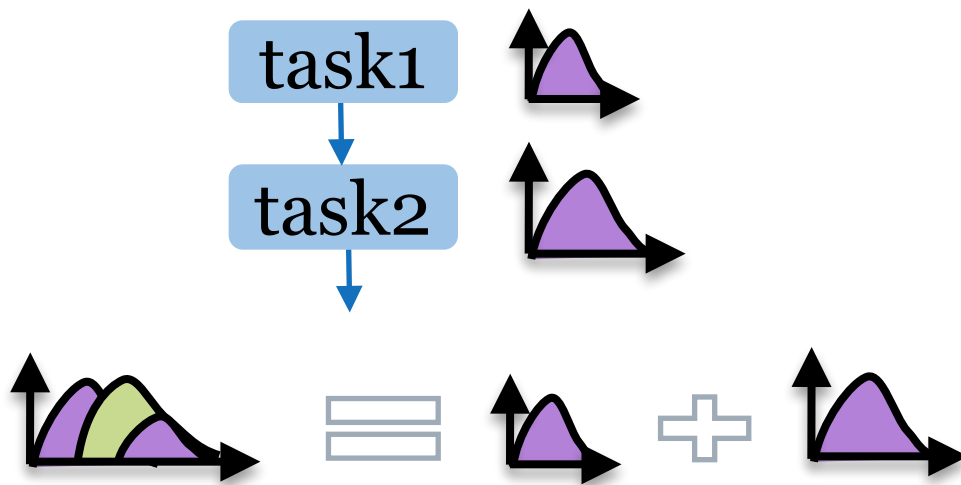
➤ Slow

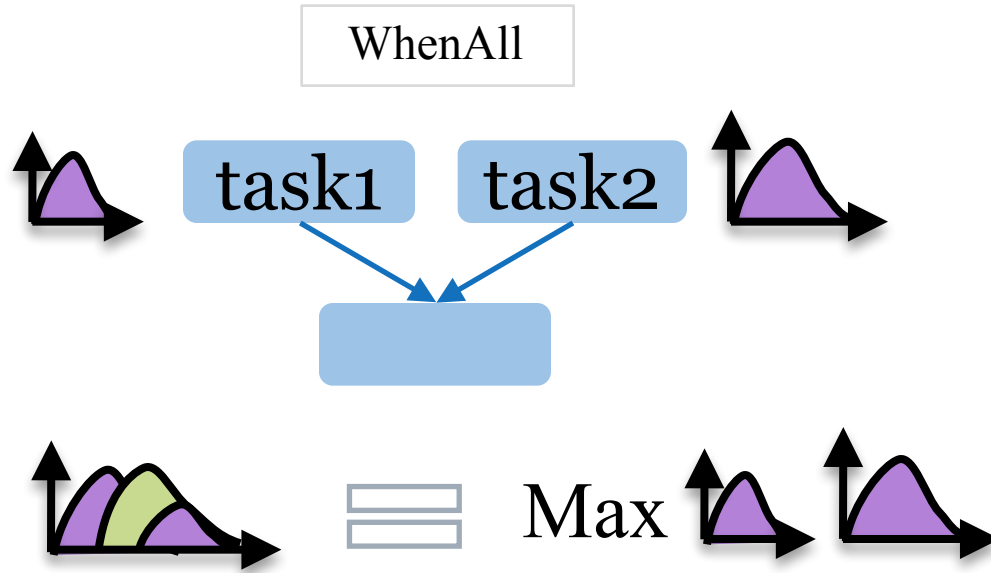
➤ High Effort

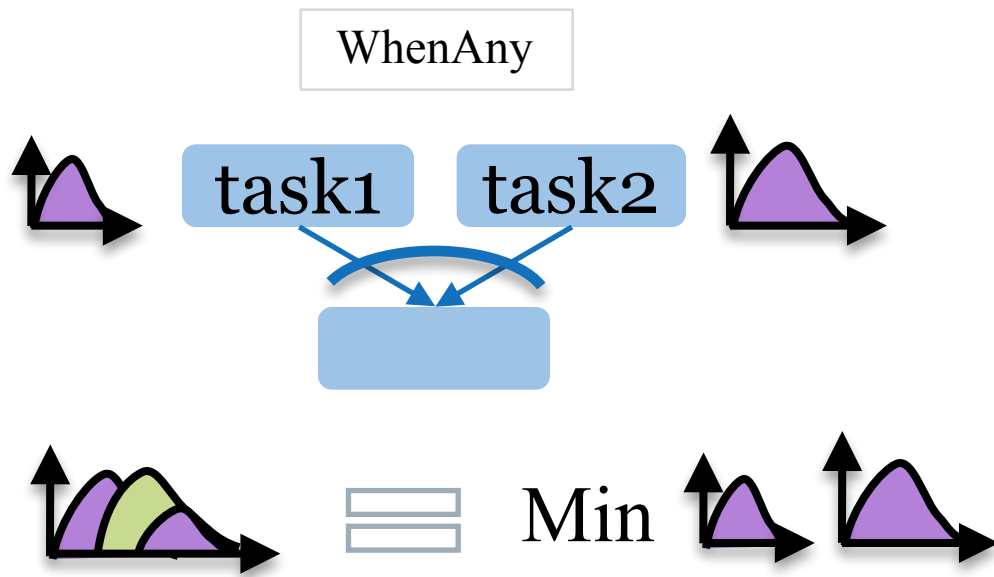
## What if?

*What if I move the front-end from basic to standard tier?*



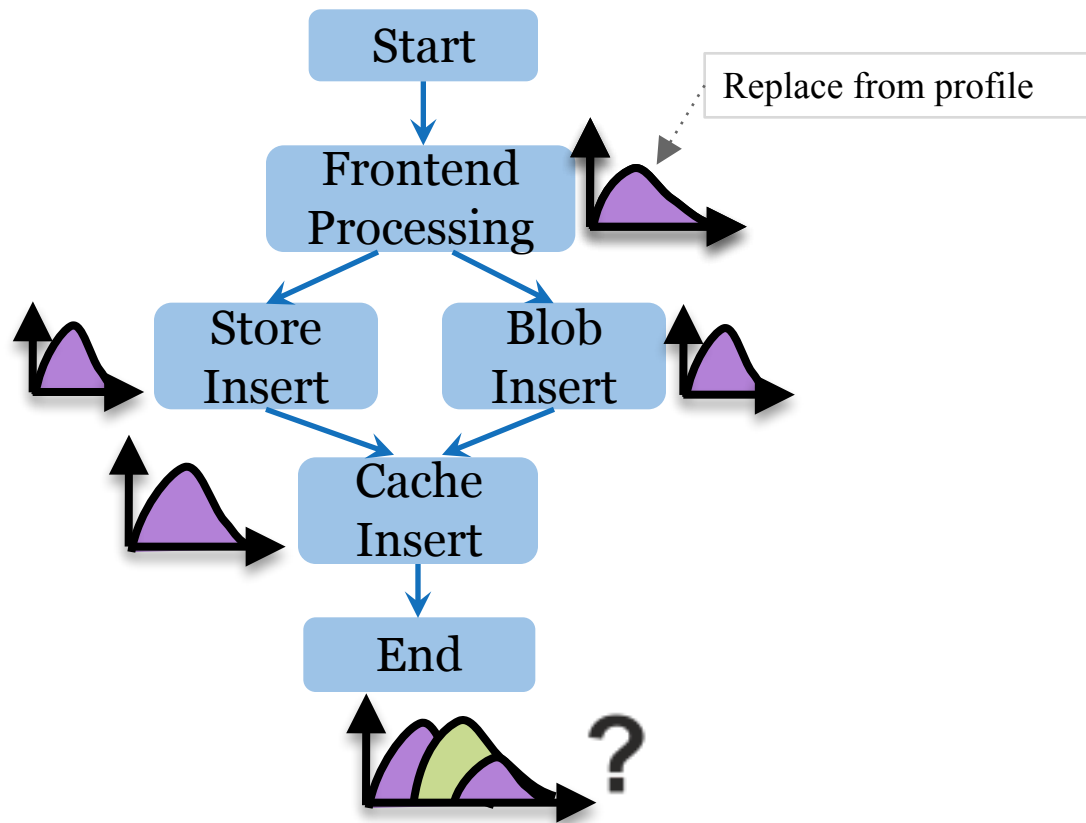






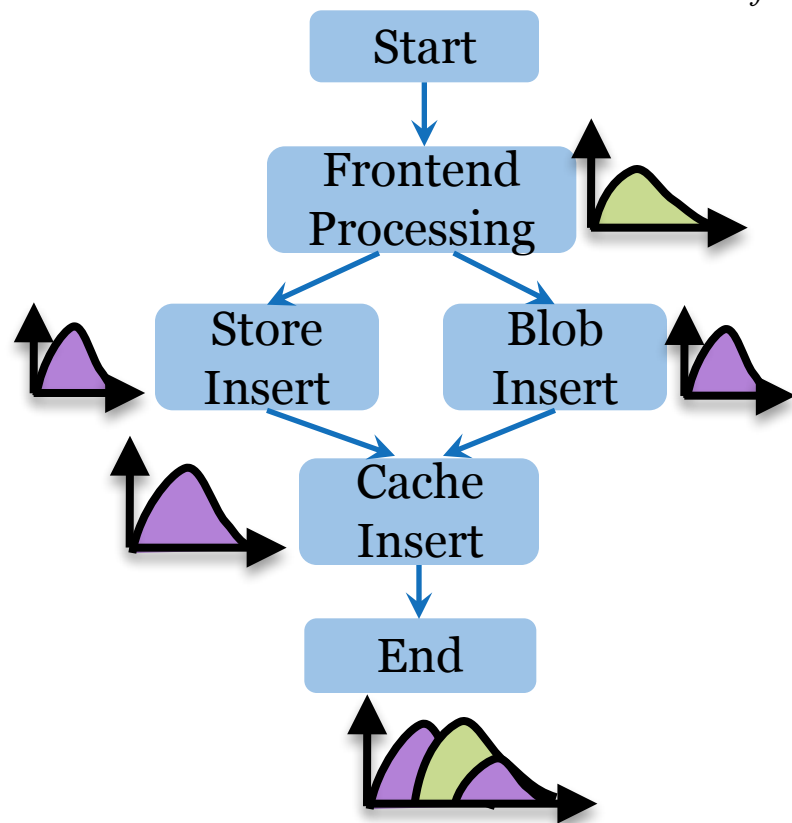
*What if I move the front-end  
from basic to standard tier?*

## Cloud-Side Latency Estimation





*What if I move the front-end  
from basic to standard tier?*



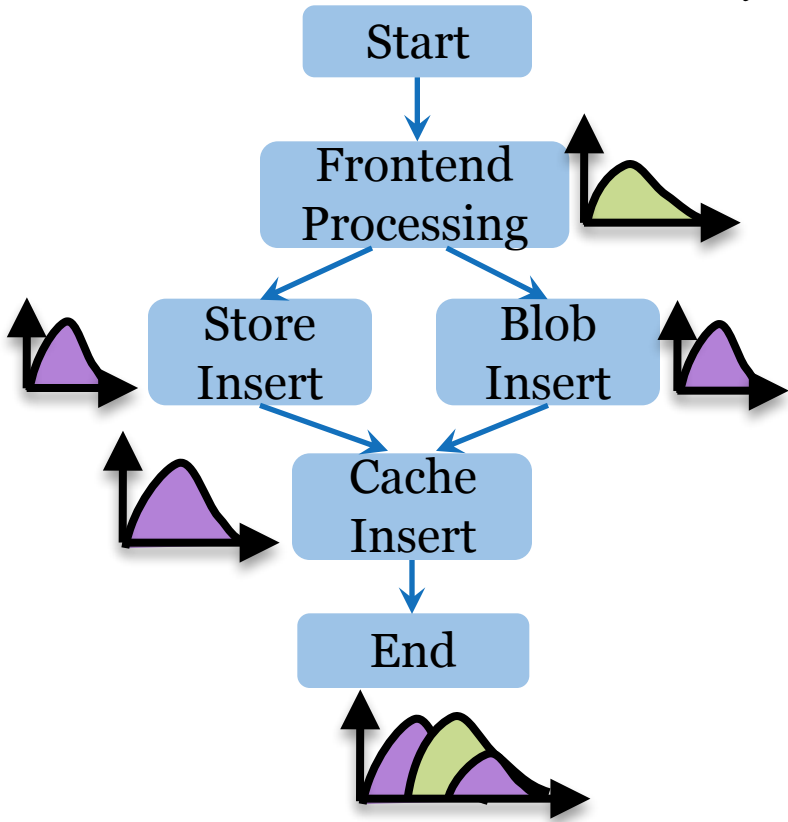
Dependency graph  
extraction

Baseline latency  
estimation

Component Profiling

Cloud-side latency  
estimation

*What if I move the front-end  
from basic to standard tier?*



Dependency graph  
extraction

Baseline latency  
estimation

Component Profiling

Cloud-side latency  
estimation



How accurate is WebPerf? ❄️

What is WebPerf's overhead?

What are the primary sources of prediction error?

Are workload hints necessary?



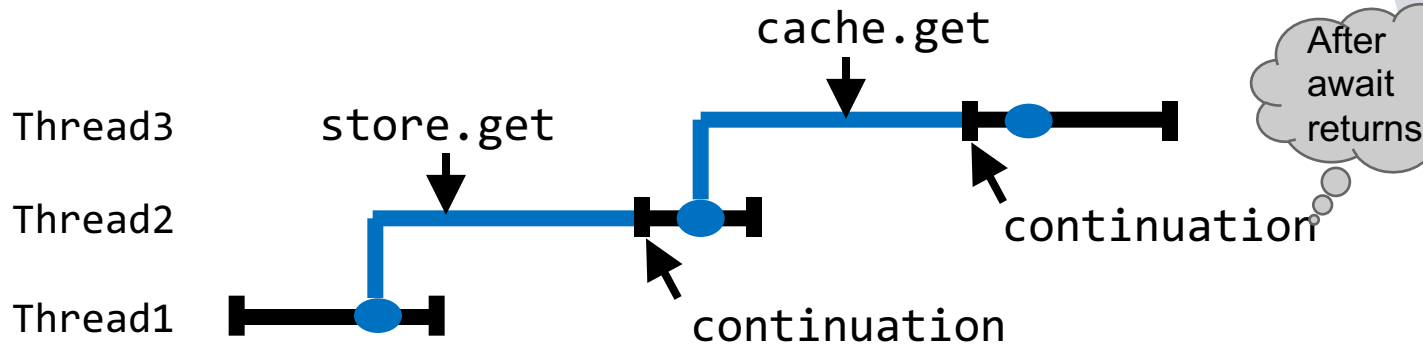
Can WebPerf predict end-to-end latency?

```
async processRequest (input)
{
  /* process input */
  task1 = store.get(key1);
  value1 = await task1;
  task2 = cache.get(key2);
  value2 = await task2;
  /* construct response */
  return response;
}
```

On front-  
end

Start task

Continue



# Distributional Difference

Measure of  
Prediction  
Accuracy

