

# VIA: Improving Internet Telephony Quality Using Predictive Relay Selection

***Junchen Jiang***, Rajdeep Das, Ganesh Ananthanarayanan, Philip A. Chou,  
Venkata N. Padmanabhan, Vyas Sekar, Esbjorn Dominique, Marcin  
Goliszewski, Dalibor Kukoleca, Renat Vafin, Hui Zhang



# Key takeaways in one minute

We studied **430 million** Skype calls and found...

**One fifth** of calls use paths with **poor** network performance



**“Managed Overlay”** could alleviate **over half** of calls on these paths




**VIA: Data-driven relaying** can realize most Managed Overlay’s potential

# Internet telephony is everywhere!




More apps focusing on Internet telephony

WeChat




Rapid growth over the last decade

Peak # of users concurrent online on Skype



# Call quality is sensitive to bad network performance

**Datasets:** Avg RTT, loss rate, jitter for each call of 430 million Skype calls  
Small fraction has user-provided quality scores ★★★★★



Thresholds of poor network performance

# Many calls have poor network performance

**Datasets:** Avg RTT, loss rate, jitter for each call of 430 million Skype calls  
Small fraction has user-provided quality scores ★★★★★

**One fifth** of calls have **poor** network performance

Our goal: **Alleviate poor network performance for Skype**



# Outline

- Problem: Network performance of Skype is bad
- ➔ **Opportunity: Managed Overlay has huge potential**
- Solution: VIA for optimal relay selection
- Evaluation: VIA is close-to-optimal

# Revisiting Overlay Networks by Managed Overlay

## Managed Overlay has new benefits

- *World-wide* distributed *DCs* as relays
- *Well connected*
- *Deployed* by many providers
- *Single* administrative entity




How much can Internet telephony benefit from it?



ExpressRoute National Clouds Peering Locations  
ExpressRoute Peering Locations

📦 AWS Regions  
⬤ AWS Edge Locations

# Selecting the best relay option



Key is to select the best **relay option** (direct, one-relay, or multi-relay)

Q1: Does picking best relay option have significant impact?

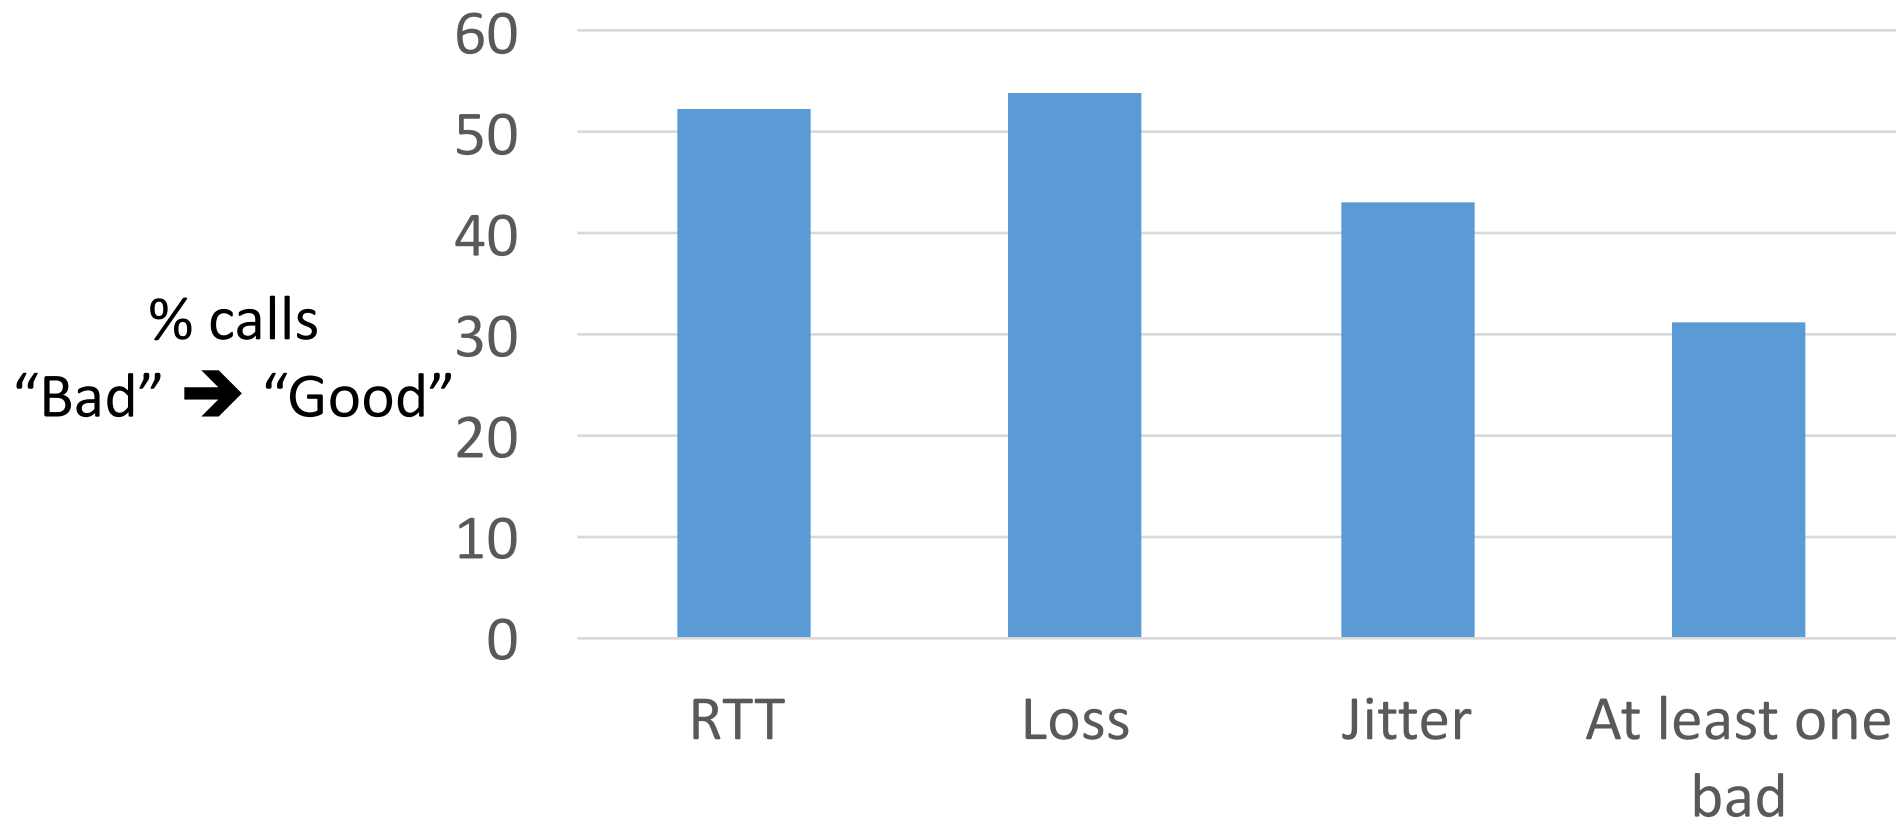
Q2: If so, how to pick the best relay option?



# Managed Overlay has huge potential benefit

Consider an oracle that picks the best relay option for each src-dst AS pair in 24 hours

Poor performance:  $\text{RTT} > 320\text{ms}$ ;  $\text{Loss rate} > 1.2\%$ ;  $\text{Jitter} > 12\text{ms}$




Substantial fraction of bad-performance calls ***could*** be alleviated

# Outline

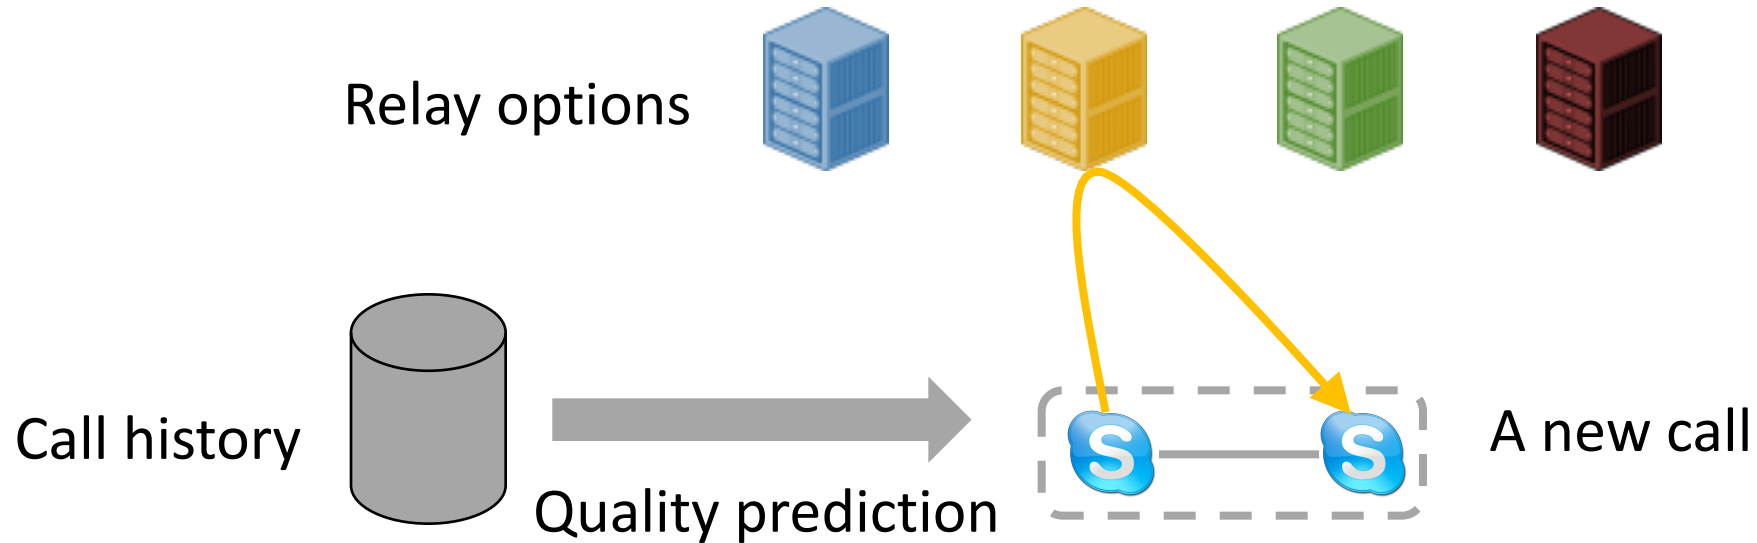
- Problem: Network performance of Skype is bad
- Opportunity: Managed Overlay has huge potential
- ➔ **Solution: VIA for optimal relay selection**
- Evaluation: VIA is close-to-optimal

# VIA: Realizing the benefit of Managed Overlay using Centralized Predictive Control



# Strawman 1: Pure prediction-based

Use long-term history to predict performance




**Problem: Call performance has great inherent **variance****

E.g., predicting “next day” using “last week” leads to over 30% error on latency.

# Strawman 2: Pure exploration-based

In a short time window, explore relay options, then exploit the best one



**Problem: Call distribution is highly **skewed****

*“# of actual calls  $\gg$  # of relay options”* NOT true for most AS pairs

# Key idea: Guided exploration


**VIA**



Prediction-based



Exploration-based

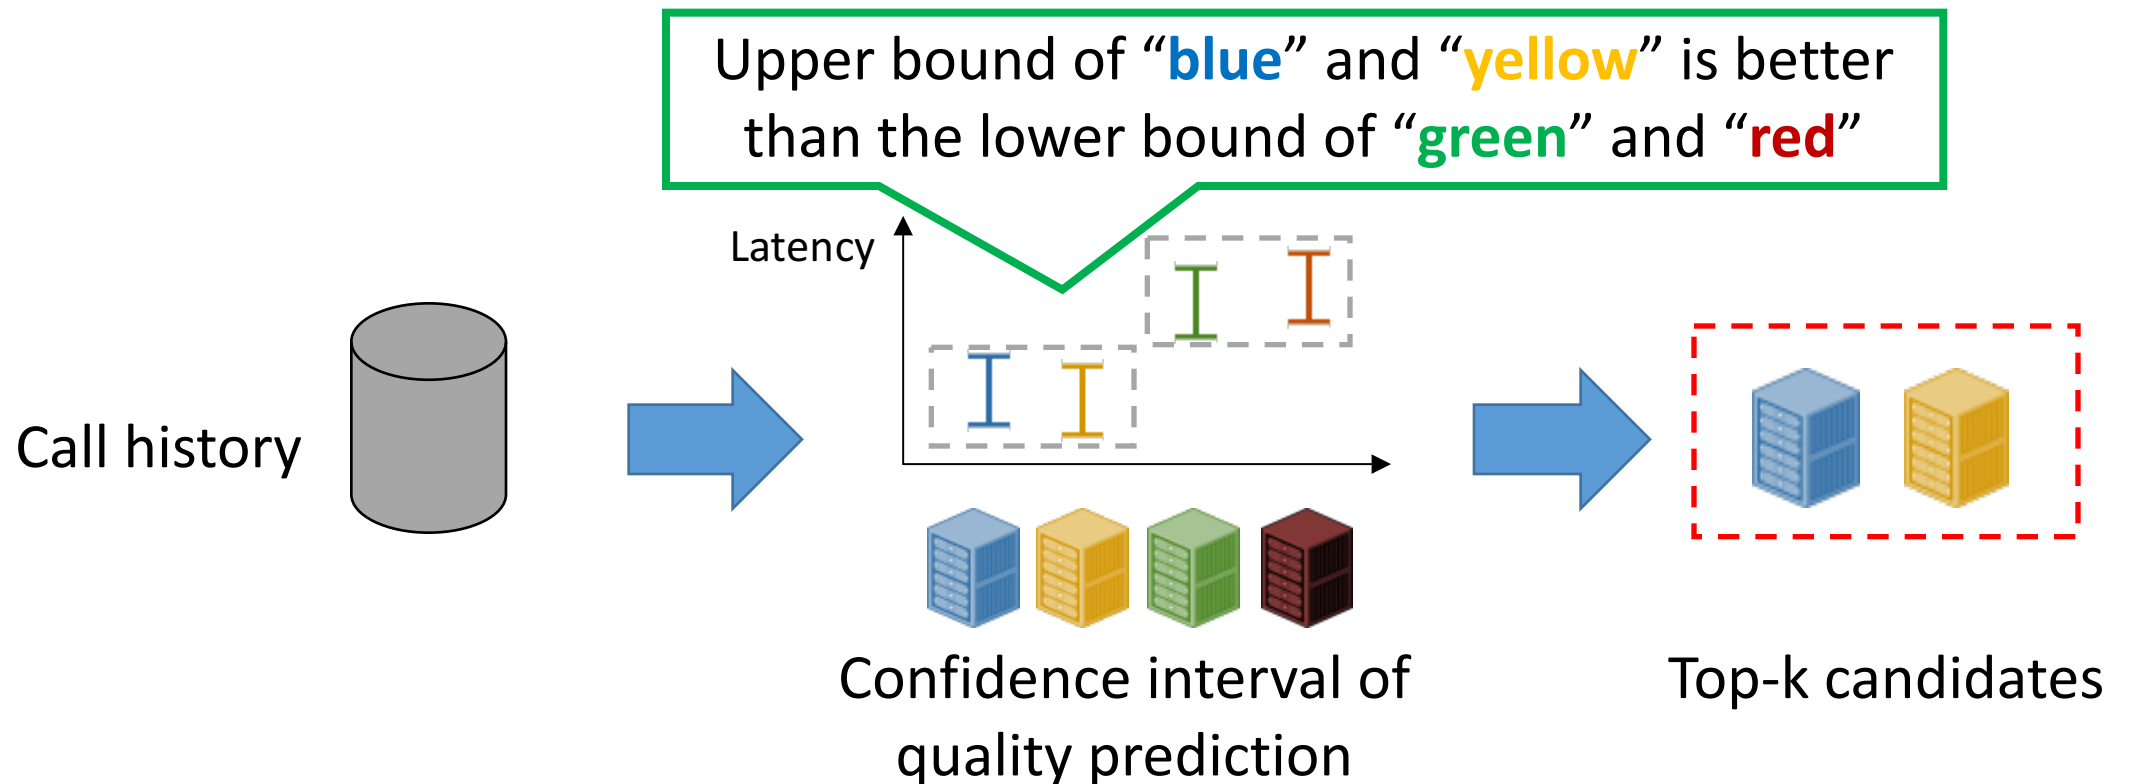


VIA's idea: **Guided Exploration**

Rough prediction can still identify top  $k$  candidates, which can be explored efficiently.


# Step 1: Prediction-based pruning

Focus on relay options whose *confidence intervals* are better than those of others.




# Step 2: Exploring top-k candidates


**Multi-armed bandit process**  
Upper Confidence Bounds (UCB1)



UCB1: Always pick the one with the highest UCB

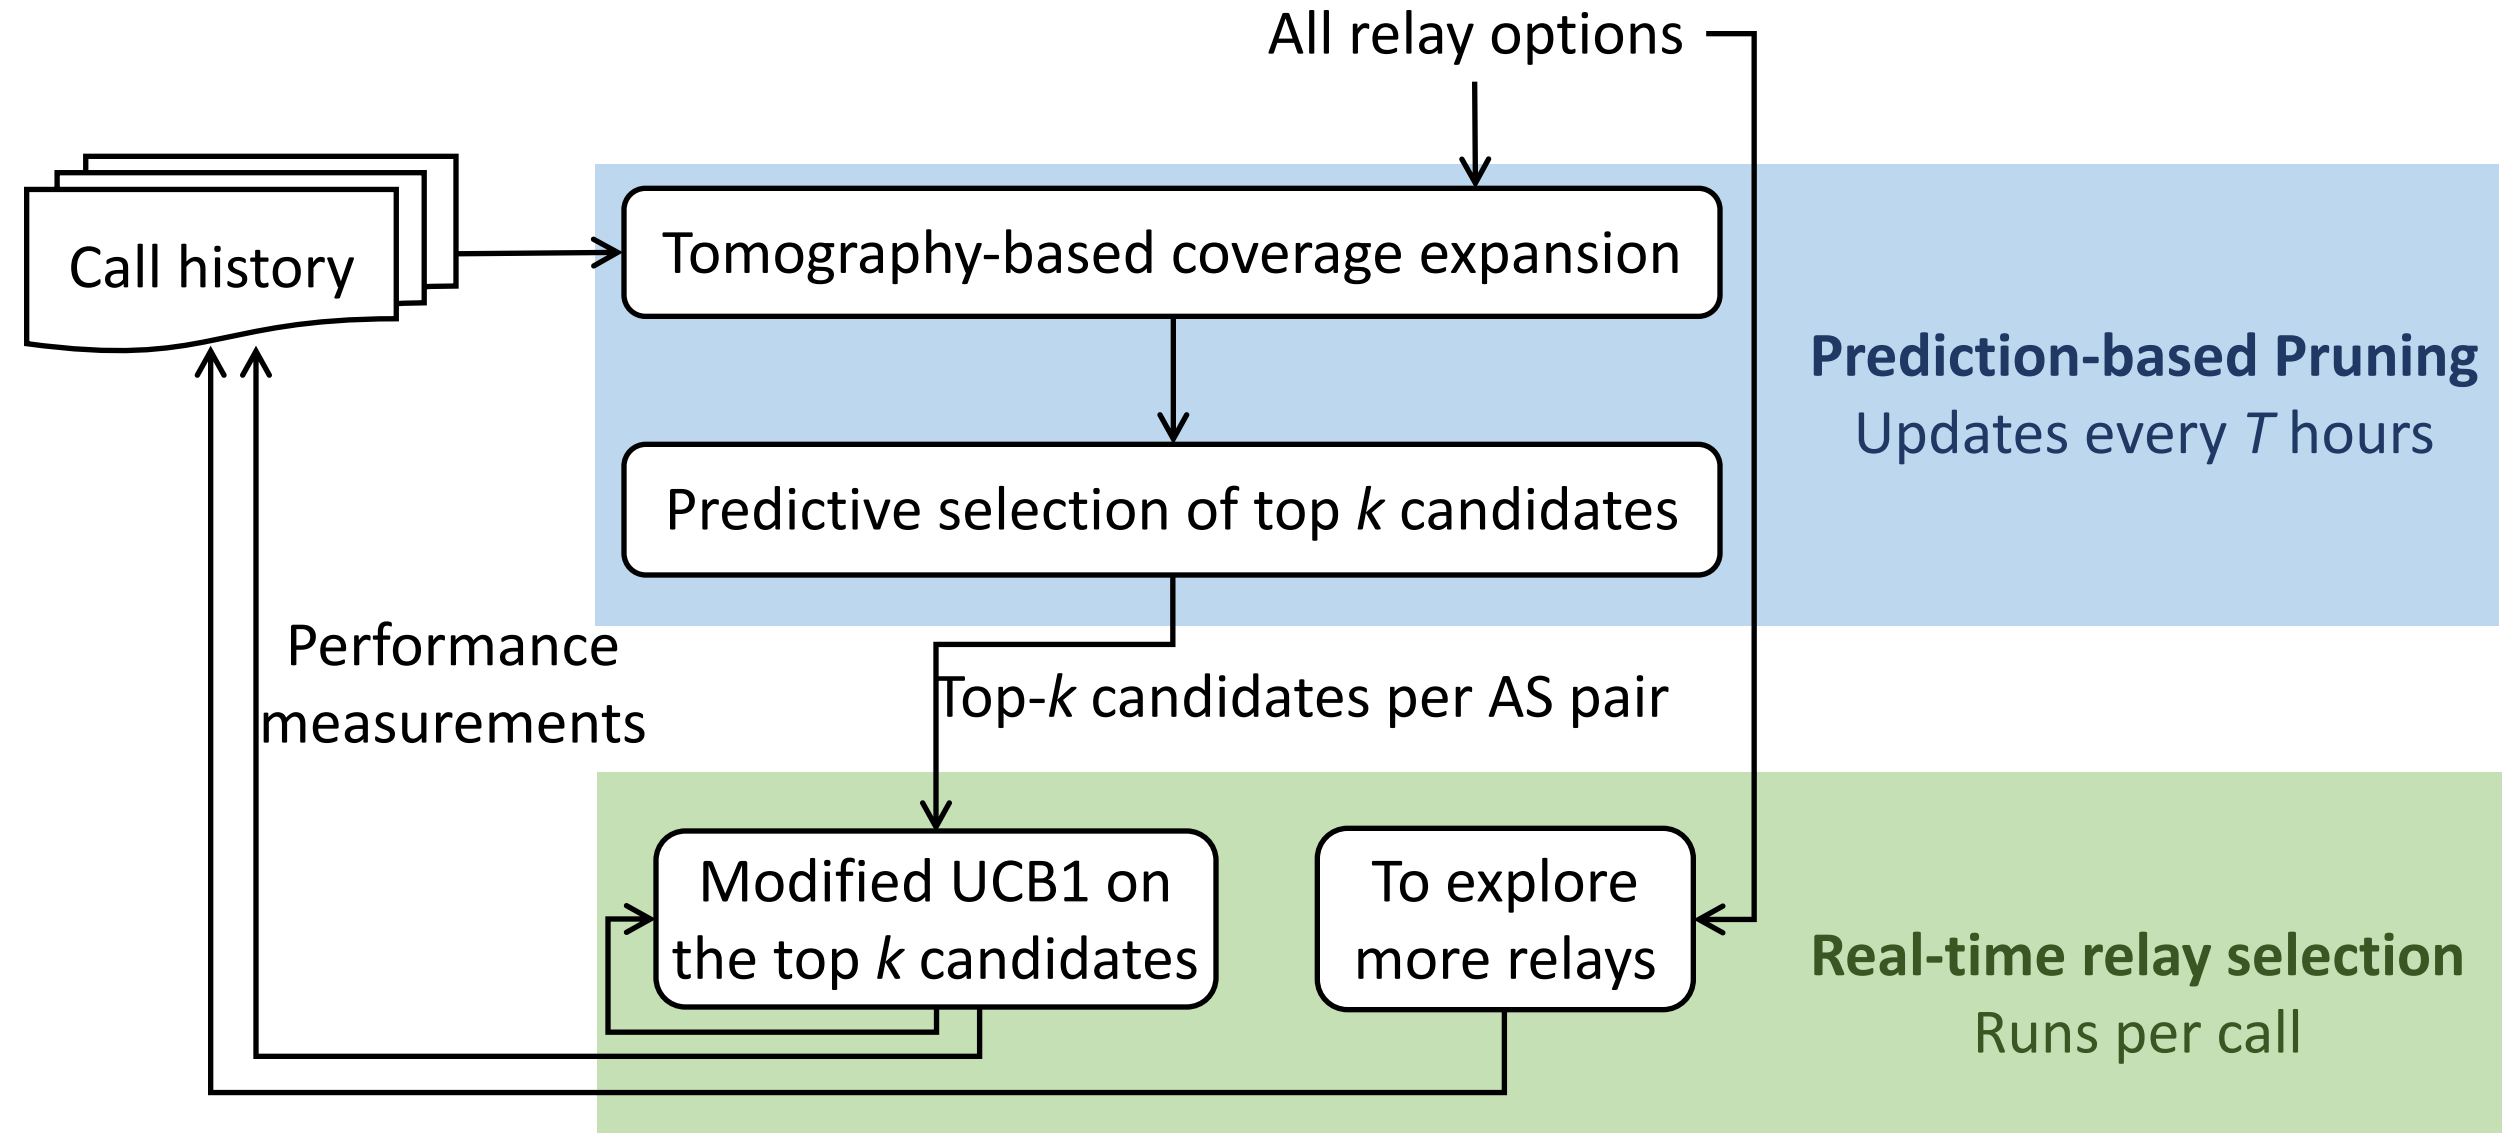


**Our problem looks like MAB**  
UCB1 with domain-specific twists





# Putting them together: *Guided Exploration in action*



# More in our paper

- Budgeted relaying
- Network tomography
- Granularity of prediction
- International vs. domestic calls

## VIA: Improving Internet Telephony Call Quality Using Predictive Relay Selection

Junchen Jiang<sup>1,2</sup>, Rajdeep Das<sup>1</sup>, Ganesh Ananthanarayanan<sup>1</sup>, Philip A. Chou<sup>1</sup>, Venkata N. Padmanabhan<sup>1</sup>, Vyas Sekar<sup>2</sup>, Esbjorn Dominique<sup>1</sup>, Marcin Goliszewski<sup>1</sup>, Dalibor Kukoleca<sup>1</sup>, Renat Vafin<sup>1</sup>, Hui Zhang<sup>2,3</sup>

<sup>1</sup> Microsoft, <sup>2</sup> CMU, <sup>3</sup> Conviva

### ABSTRACT

The use of the Internet for voice calls is here to stay. In spite of the volume and importance of Internet telephony, we have little understanding of (1) how network performance impacts user-perceived call quality, and (2) why and where such quality problems occur in the wild. To bridge this gap, we analyze a data set of 430 million calls from Skype, with clients across 1900 ASes and 126 countries. We observe that call quality problems are quite pervasive. More importantly, these problems are *significantly spread out geographically* and over time, thereby making simple fixes targeted at specific “pockets” of poor performance largely ineffective.

To alleviate call quality problems, we present an architecture called VIA that revisits the use of classical overlay techniques to relay calls. We argue that this approach is both timely and pragmatic given the emergence of private back-

### CCS Concepts

• **Applied computing** → Internet telephony; • **Networks** → Overlay and other logical network structures; Network performance analysis; Network measurement;

### Keywords

Internet telephony; Quality of experience; Predictive relay selection; Managed overlay networks

### 1 Introduction


Over the last several years, we have seen a dramatic rise in Internet-based telephony, especially for long-distance international calling [5, 6]. The importance of audio calling is evident as almost all major content and social networking platforms today offer some form of Internet calling capability (e.g., Skype, Google Hangouts, Facebook Messenger,

# Outline


- Problem: Network performance of Skype is bad
  - Opportunity: Managed Overlay has huge potential
  - Solution: VIA for optimal relay selection
- ➔ **Evaluation: VIA is close-to-optimal**

# VIA achieves close-to-optimal performance

Poor performance:  $\text{RTT} > 320\text{ms}$ ;  $\text{Loss rate} > 1.2\%$ ;  $\text{Jitter} > 12\text{ms}$



# Benefit varies across ASes



# Conclusion

- **Internet telephony** is here to stay
- Problem: **One fifth** of calls have bad network performance.
- Opportunity: **Managed Overlay** *could* significantly reduce bad-performance calls.
- Challenges: Optimal relay selection  
Pure prediction and pure exploration won't work!
- Solution: **VIA** can realize most of the benefit of Managed Overlay.  
Key idea: **Guided exploration**:  
Predictive pruning of relay options + Exploration over top k candidates