

# An Internet-Wide Analysis of Traffic Policing


Tobias Flach, Pavlos Papageorge, Andreas Terzis, Luis Pedrosa,  
Yuchung Cheng, Tayeb Karim, Ethan Katz-Bassett, Ramesh Govindan

[policing-paper@google.com](mailto:policing-paper@google.com)



USC University of  
Southern California






amazon.com®

bing

You**Tube**

**NETFLIX**

Want to accommodate multitude  
of services/policies  
→ Traffic Engineering



Exponential  
growth of  
video traffic



Account for ~  
50% of traffic in  
North America

Want to maximize quality  
of experience (QoE) for  
their users

Often need high bitrate  
with low tolerance for  
latency and packet loss

# Traffic Engineering: Policing vs. Shaping

Goal: Enforce a rate limit (maximum throughput)

Solutions:

- a. **Drop** packets once the limit is reached  
→ Traffic Policing

***Focus of  
this talk***

- b. **Queue** packets (and send them out at the maximum rate)  
→ Traffic Shaping


# Contribution

Analyze the **prevalence** and **impact** of **traffic policing** on a **global scale**, as well as explore ways to **mitigate** the **impact** of policers.


# Outline

1. How Policing Works
2. Detecting the Effects of Policing in Packet Captures
3. A Global-Scale Analysis of Policing in the Internet
4. Mitigating the Impact of Policers

# How Policing Works




# How Policing Works






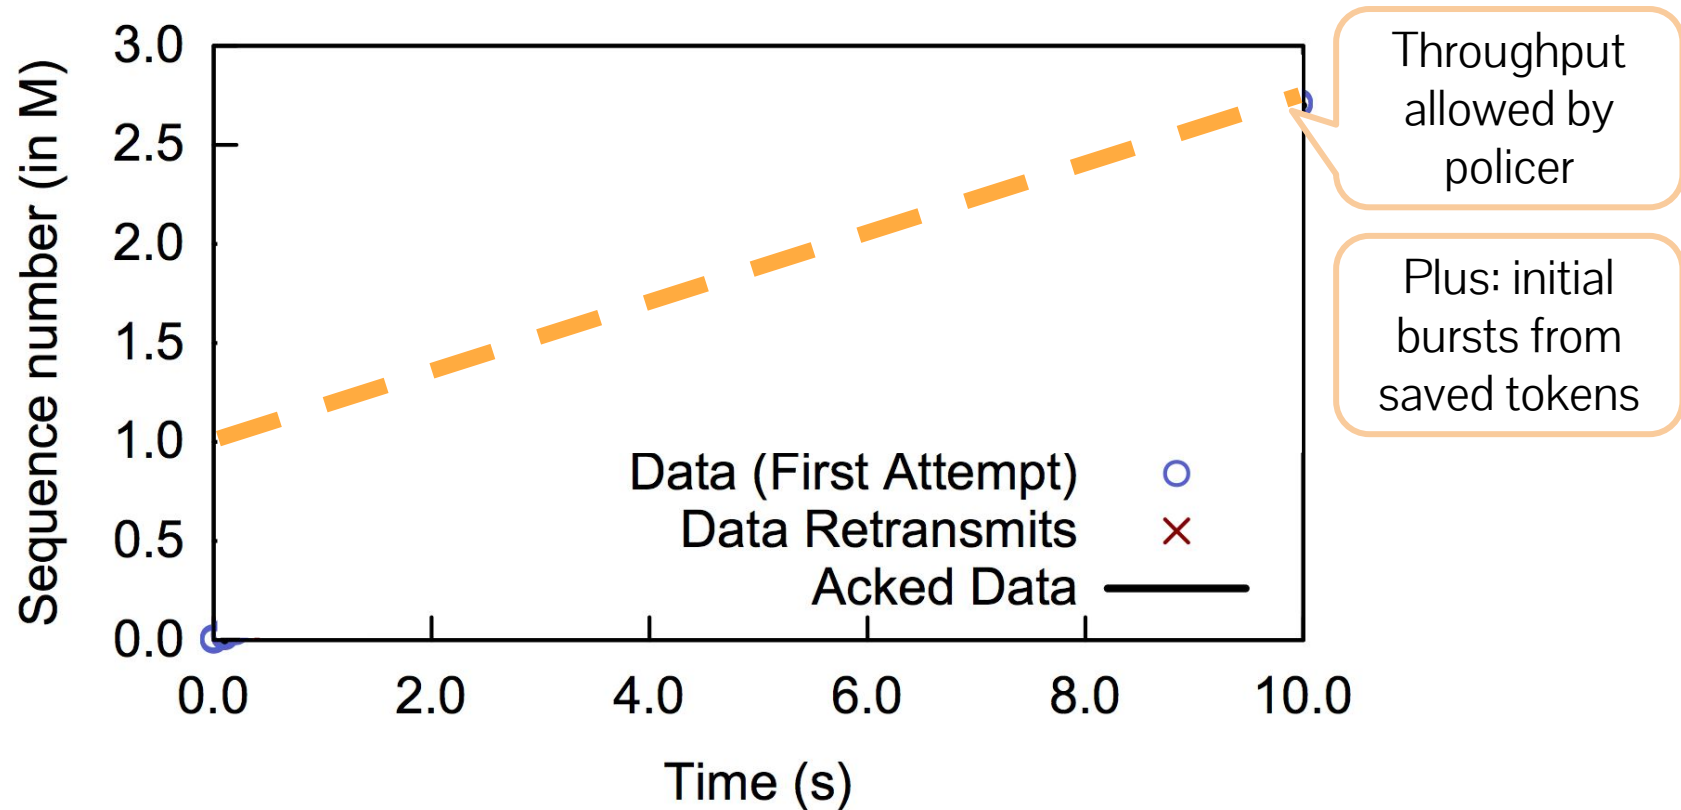
# Policing in Action




# Policing in Action



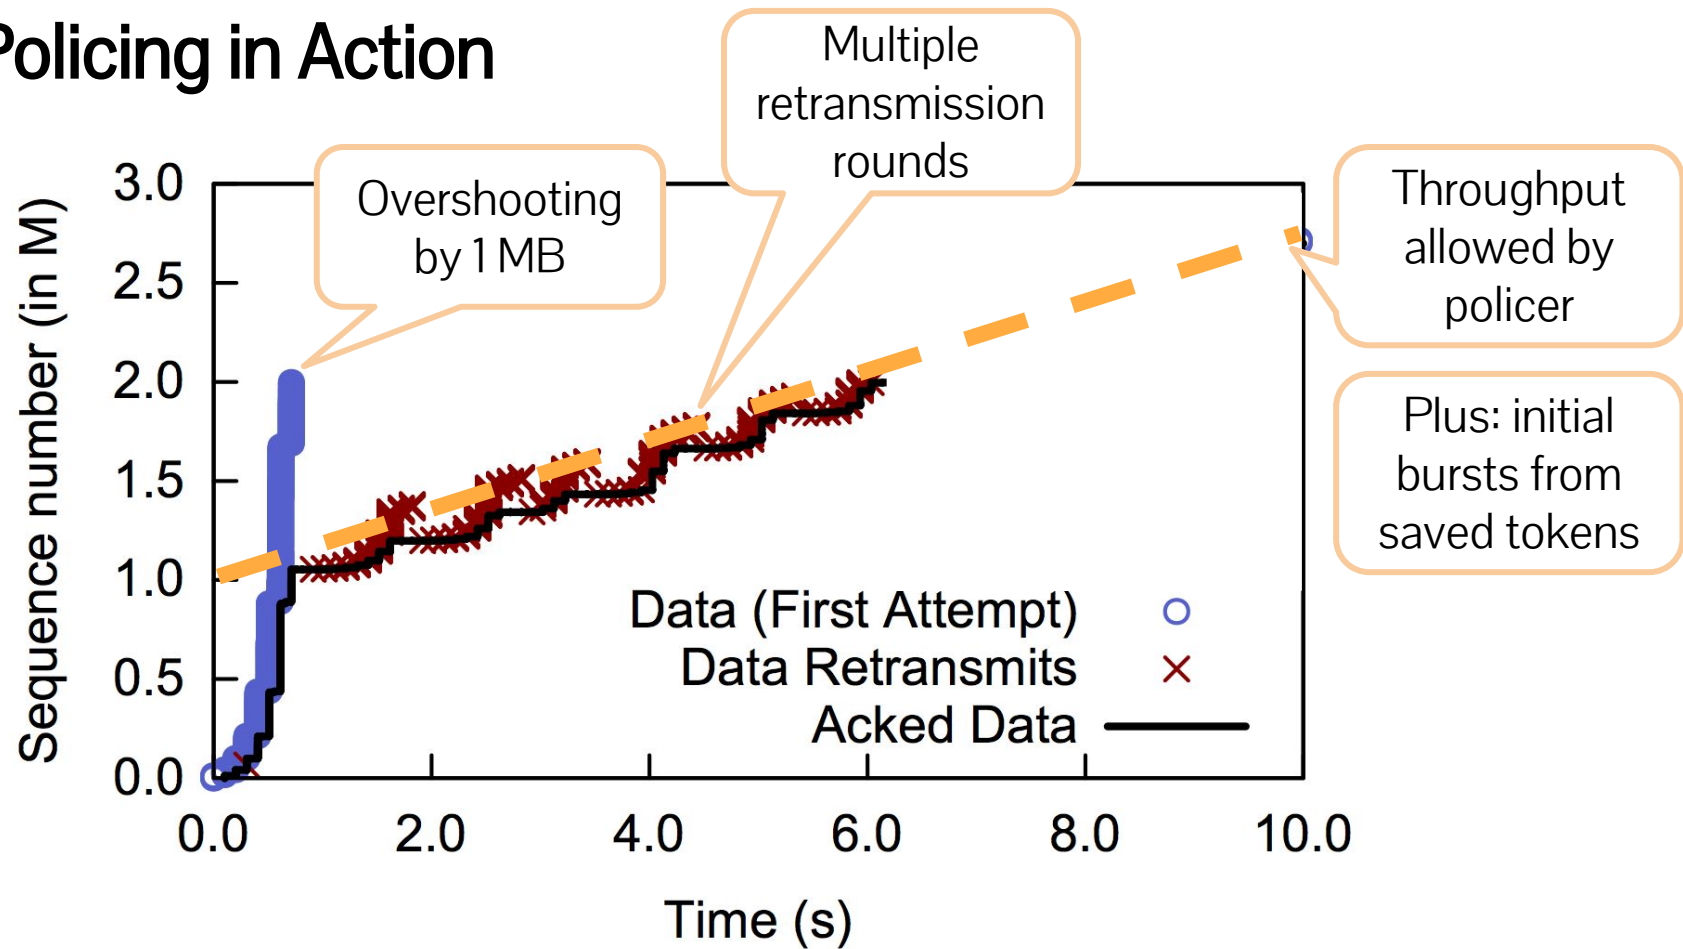
# Policing in Action




# Policing in Action



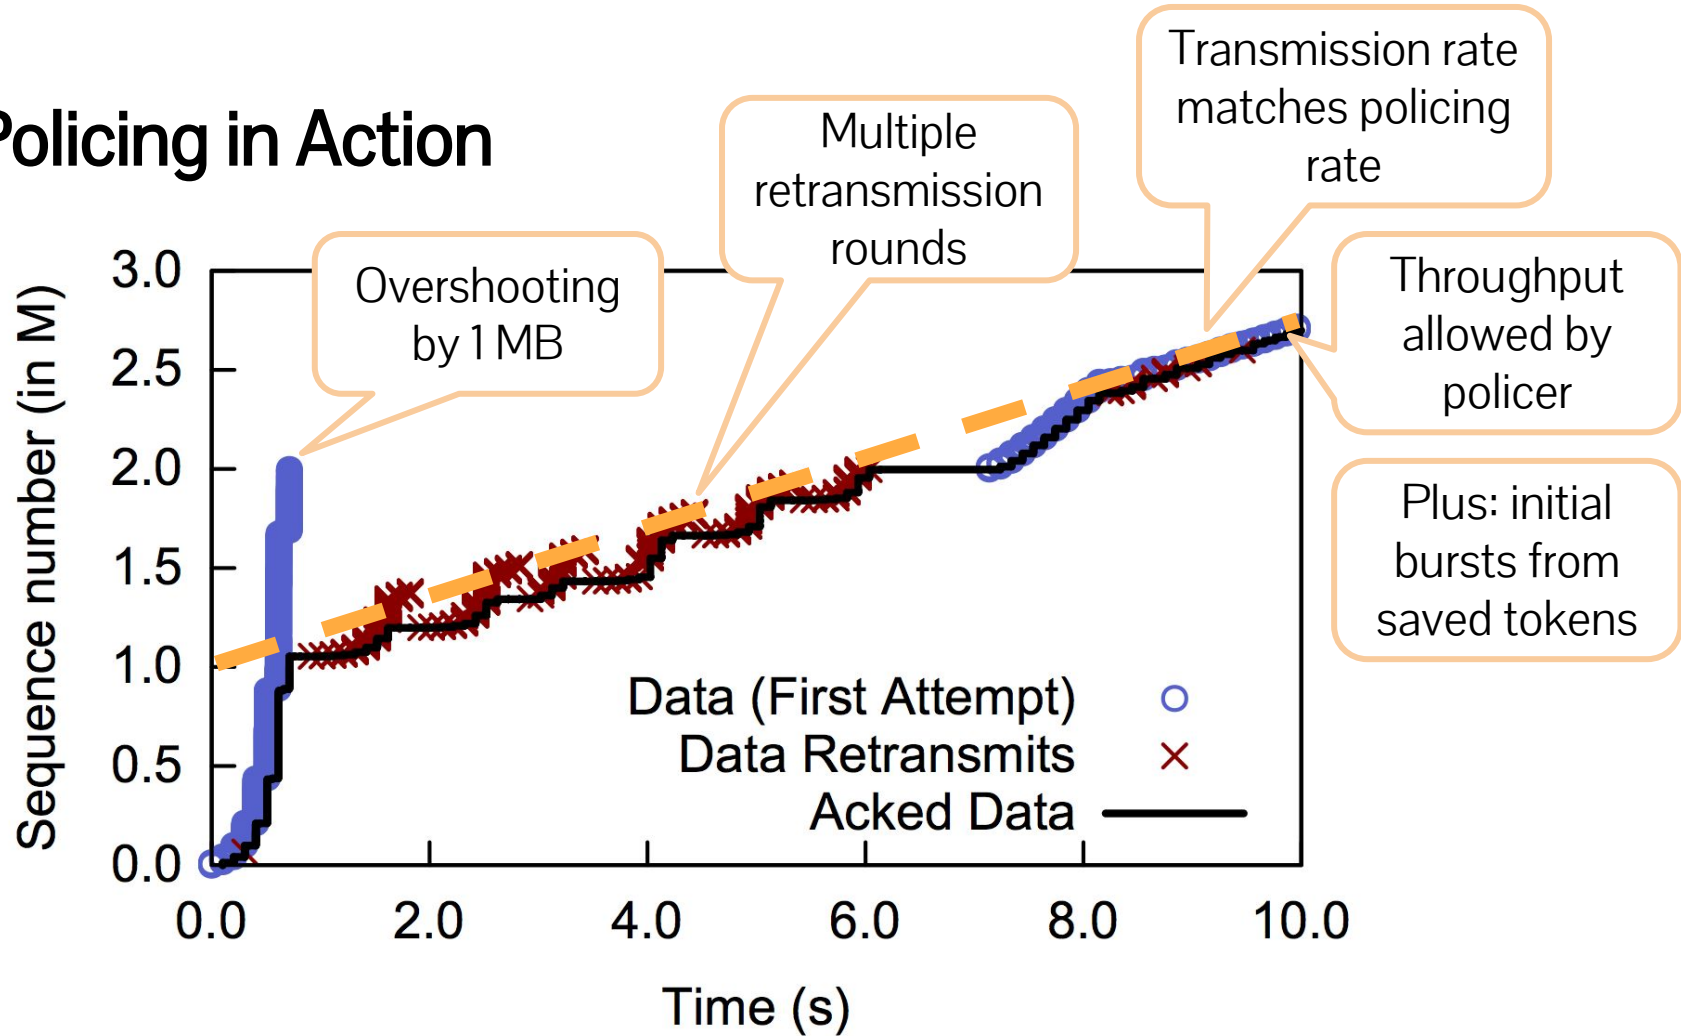
# Policing in Action



# Policing in Action



# Policing in Action



# Policing can have negative side effects for all parties

- Content providers
  - Excess load on servers forced to retransmit dropped packets  
(global average: 20% retransmissions vs. 2% when not policed)
- ISPs
  - Transport traffic across the Internet only for it to be dropped by the policer
  - Incurs avoidable transit costs
- Users
  - Can interact badly with TCP-based applications
  - We measured degraded video quality of experience (QoE) → user dissatisfaction



Analyze the


**prevalence** and **impact** of policing on a **global scale**

Develop a  
mechanism to detect  
policing in packet  
captures

Tie connection  
performance back to  
already collected  
application metrics

Collect packet traces  
for sampled client  
connections at most  
Google frontends

# Analysis Pipeline



# Detection Algorithm

Progress

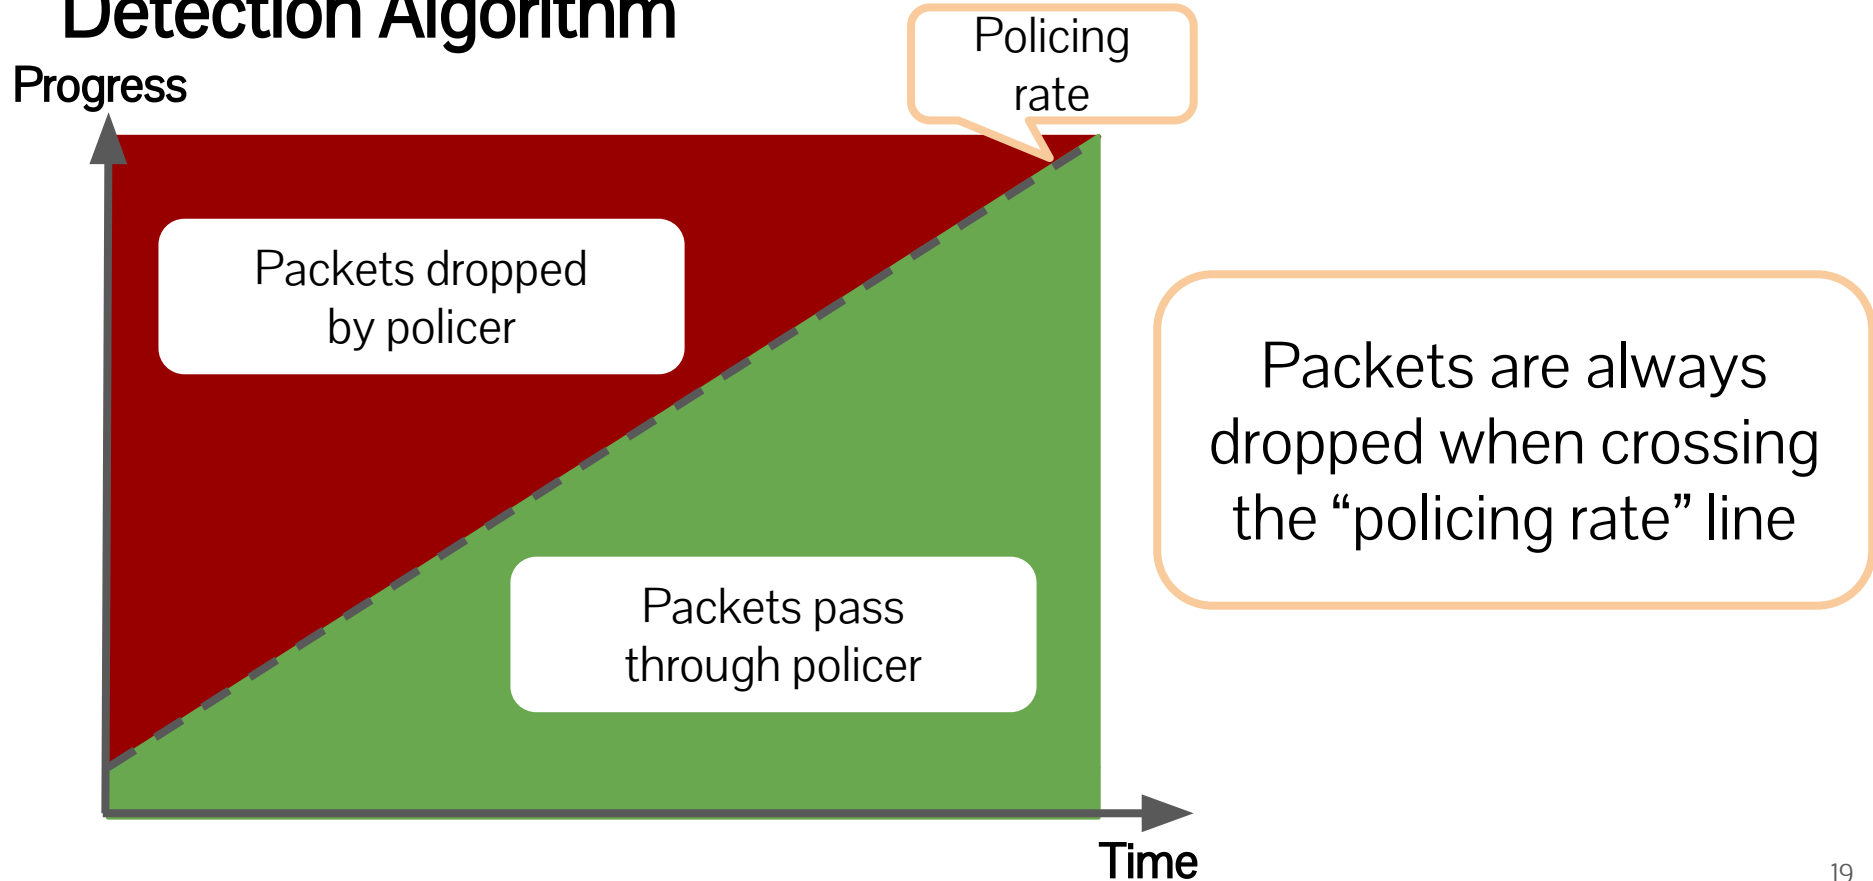
Policing  
rate

Packets dropped  
by policer

Packets pass  
through policer


Packets are always  
dropped when crossing  
the “policing rate” line

Time



# Detection Algorithm

Progress



1

Find the  
policing rate


- Use measured throughput between an early and late loss as estimate

2


Match performance  
to expected policing  
behavior

- Everything above the policing rate gets dropped
- (Almost) nothing below the policing rate gets dropped

# Avoiding Falsely Labeling Loss as Policing




But: Traffic below policing rate  
should go through



But: Traffic above policing rate  
should be dropped

# Congestion Looks Similar to Policing!



Packets are usually dropped when a router's buffer is already full


Buffer fills → queuing delay increases

Use inflated latency as signal that loss is not caused by a policer

# Validation 1: Lab Setting

- Goal: Approximate the accuracy of our heuristic
- Generated test traces covering common reasons for dropped packets
  - Policing (used a router with support for policing)
  - Congestion
  - Random loss
  - Shaping
- High accuracy for almost all configurations (see paper for details)
  - Policing: 93%
  - All other reasons for loss: > 99%

## Validation 2: Live Traffic



- Observed only few policing rates in ISP deep dives
  - ISPs enforce a limited set of data plans
- Confirmed that per ISP policing rates cluster around a few values across the whole dataset
- And: Observed no consistency across flows without policing



# Outline

1. How Policing Works
2. Detecting the Effects of Policing in Packet Captures
3. A Global-Scale Analysis of Policing in the Internet
4. Mitigating the Impact of Policers

# Internet-Wide Analysis of Policing

- Sampled flows collected from most of Google's CDN servers
  - 7-day sampling period (in September 2015)
  - 277 billion TCP packets
  - 270 TB of data
  - 800 million HTTP queries
  - Clients in over 28,400 ASes
- To tie TCP performance to application performance, we analyzed flows at HTTP request/response (“segment”) granularity

# #1: Prevalence of Policing

Region	Policed segments (overall)
Africa	1.3%
Asia	1.3%
Australia	0.4%
Europe	0.7%
N. America	0.2%
S. America	0.7%

# #1: Prevalence of Policing

Up to 7% of lossy segments  
are policed

Lossy:  
15 losses or more  
per segment

Region	Policed segments (overall)	Policed (among lossy)
Africa	1.3%	6.2%
Asia	1.3%	6.6%
Australia	0.4%	2.0%
Europe	0.7%	5.0%
N. America	0.2%	2.6%
S. America	0.7%	4.1%

## #2: Policer-induced Losses

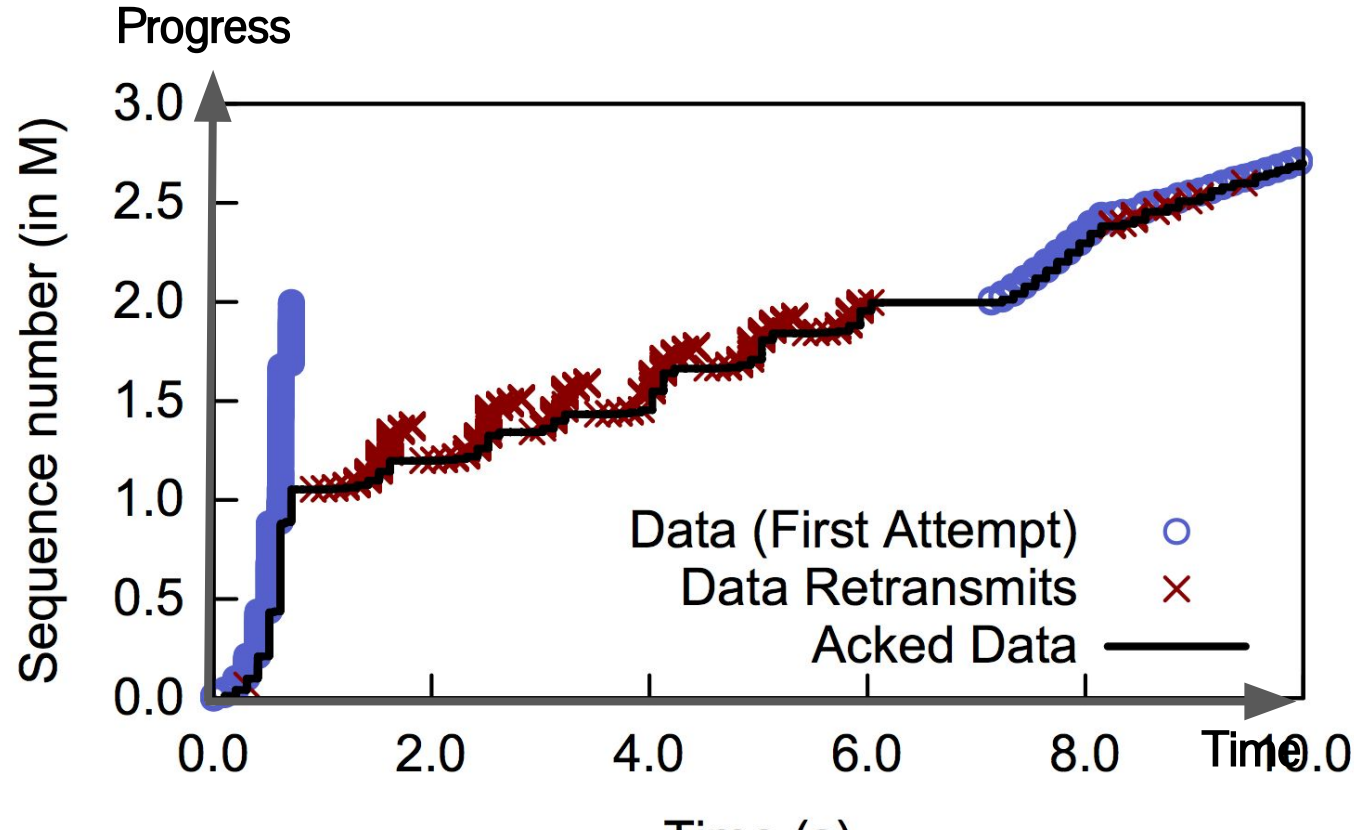
Up to 7% of lossy segments  
are policed

Average loss rate increases from  
2% to over 20% when policed


Lossy:  
15 losses or more  
per segment

Region	Policed segments (overall)	Policed (among lossy)	Loss (policed)	Loss (non-policed)
Africa	1.3%	6.2%	27.5%	4.1%
Asia	1.3%	6.6%	24.9%	2.9%
Australia	0.4%	2.0%	21.0%	1.8%
Europe	0.7%	5.0%	20.4%	1.3%
N. America	0.2%	2.6%	22.5%	1.0%
S. America	0.7%	4.1%	22.8%	2.3%

# Sudden Bandwidth Change Induces Heavy Loss



# Sudden Bandwidth Change Induces Heavy Loss

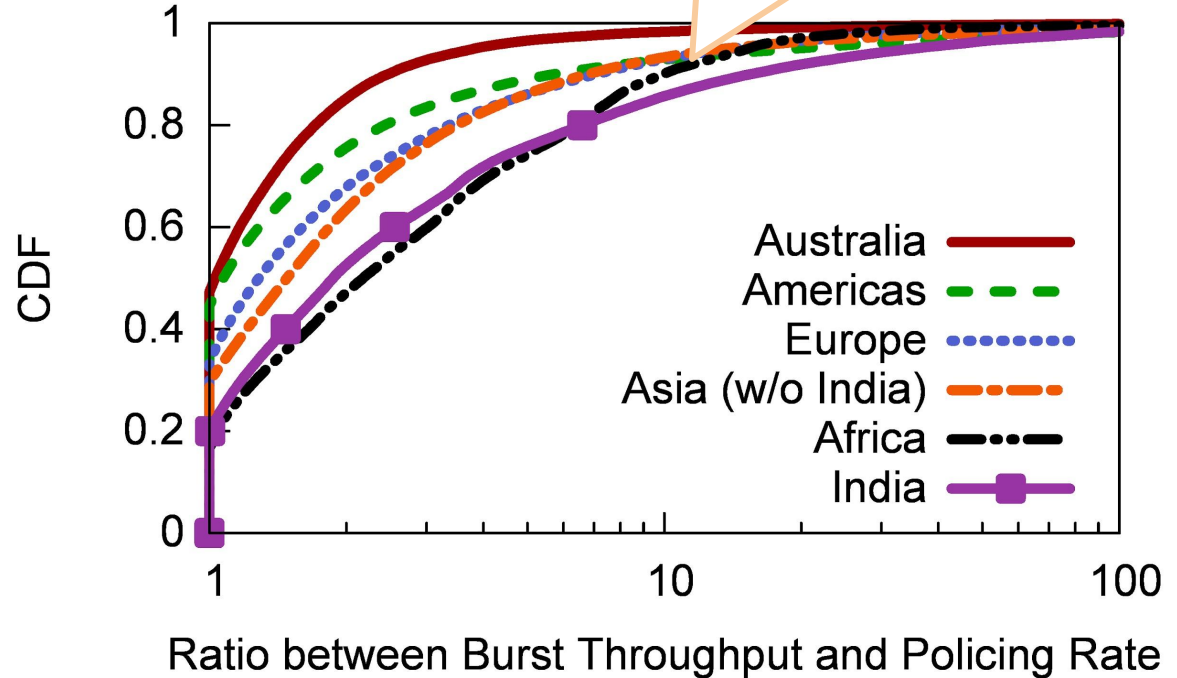


# #3: Burst Throughput vs. Policing Rate

Up to 7% of lossy segments are policed

Average loss rate increases from 2% to over 20% when policed

Policing rate often over 50% lower than burst throughput





# Quality of Experience Metrics

## Rebuffer Time:

Time that a video is paused *after playback started* due to insufficient stream data buffered

## Watch Time:

Fraction of the video watched by the user

## Rebuffer to Watch Time Ratio:

Goal is *zero* (no rebuffering delays after playback started).


## #4: Impact on Quality of Experience

Up to 7% of lossy segments are policed

Average loss rate increases from 2% to over 20% when policed

Policing rate often over 50% lower than burst throughput

In the tail, policed segments can have up to 200% higher rebuffering times



# Mitigating Policer Impact

## For content providers

No access to policers and their configurations

But can control transmission patterns to minimize risk of hitting an empty token bucket

## For policing ISPs

Access to policers and their configurations

Can deploy alternative traffic management techniques

# Mitigating Policer Impact

## For content providers

Rate limiting

Pacing

Reducing losses  
during recovery in Linux

## For policing ISPs

Policer optimization

Shaping

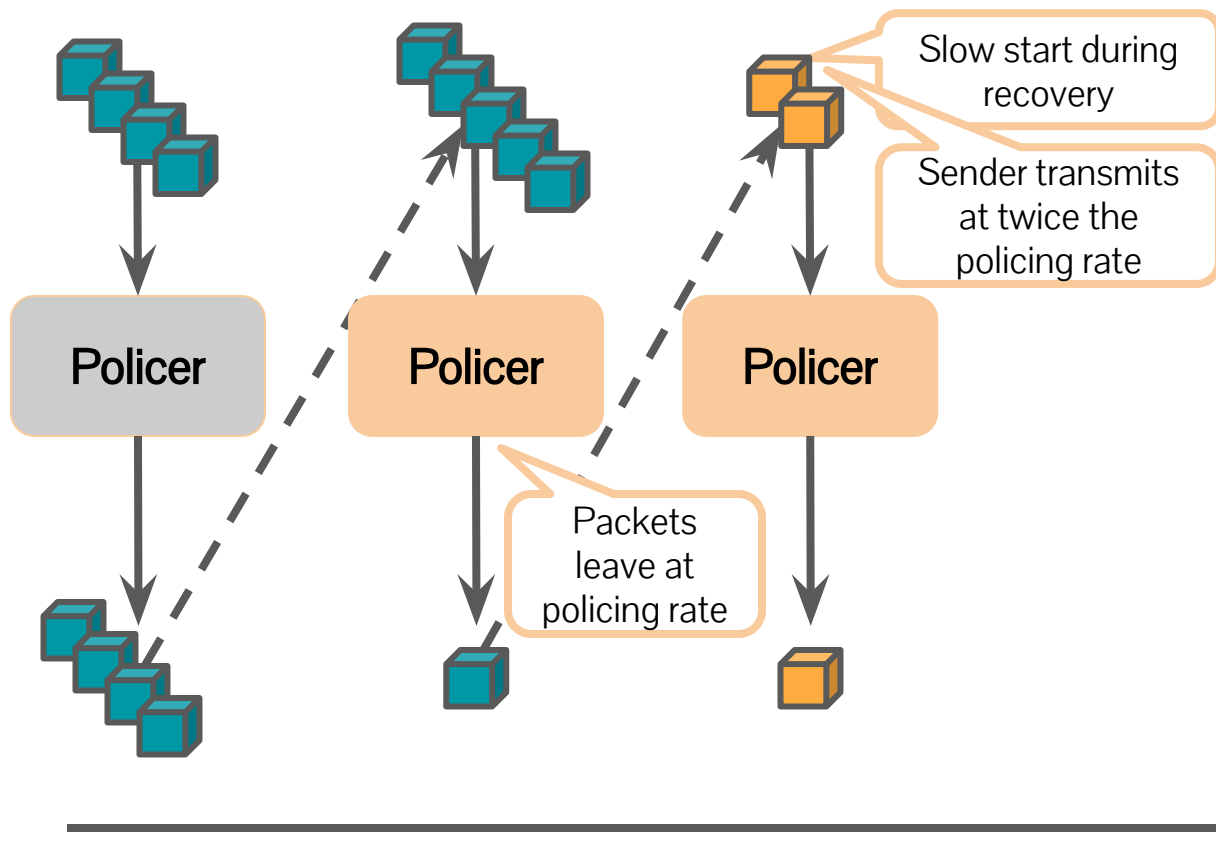
# Mitigating Policer Impact

**For content providers**

**For policing ISPs**

Reducing losses  
during recovery in Linux

# Reducing Losses During Recovery in Linux



Send only one packet per ACK

Solution:  
Packet conservation until  
ACKs indicate no further  
losses

- Reduces median loss rates by 10 to 20%
- Upstreamed to Linux kernel 4.2

# Conclusion

- ISPs need ways to deal with increasing traffic demands and want to enforce plans → traffic policing is one option
- On a global scale up to 7% of lossy segments are affected by traffic policing
- Policed connections see ...
  - Much higher loss rates
  - Long recovery times when policers allow initial bursts
  - Worse video rebuffering times (QoE)
- Negative effects can be mitigated
  - Content providers: Rate limiting, pacing, prevention of loss during recovery
  - ISPs: Better policing configurations, shaping

Questions? Email us: [policing-paper@google.com](mailto:policing-paper@google.com)

Data: <http://usc-nsl.github.io/policing-detection/>