# PISCES: A Programmable, Protocol-Independent Software Switch

Muhammad Shahbaz, **Sean Choi**, Ben Pfaff, Changhoon Kim,
Nick Feamster, Nick McKeown, and Jennifer Rexford



vmware
BAREFOOT NETWORKS

**Fixed Set of Protocols**

TCP

Ethernet

IPv4

UDP

HTTP

IPv6

BGP

TLS

**Fixed-Function Switch Chip**

- Ease of **Adding** new protocols
- Ease of **Removing** unused protocols
- Gain greater **Visibility** into the network
- **Perform** network functions at the switch

**Custom Protocols**

**Programmable Switching Chip**

TCP

Ethernet

IPv4

CUSTOM_P

HTTP

IPv6

BGP

TLS

intel®

pliant

BAREFOOT NETWORKS

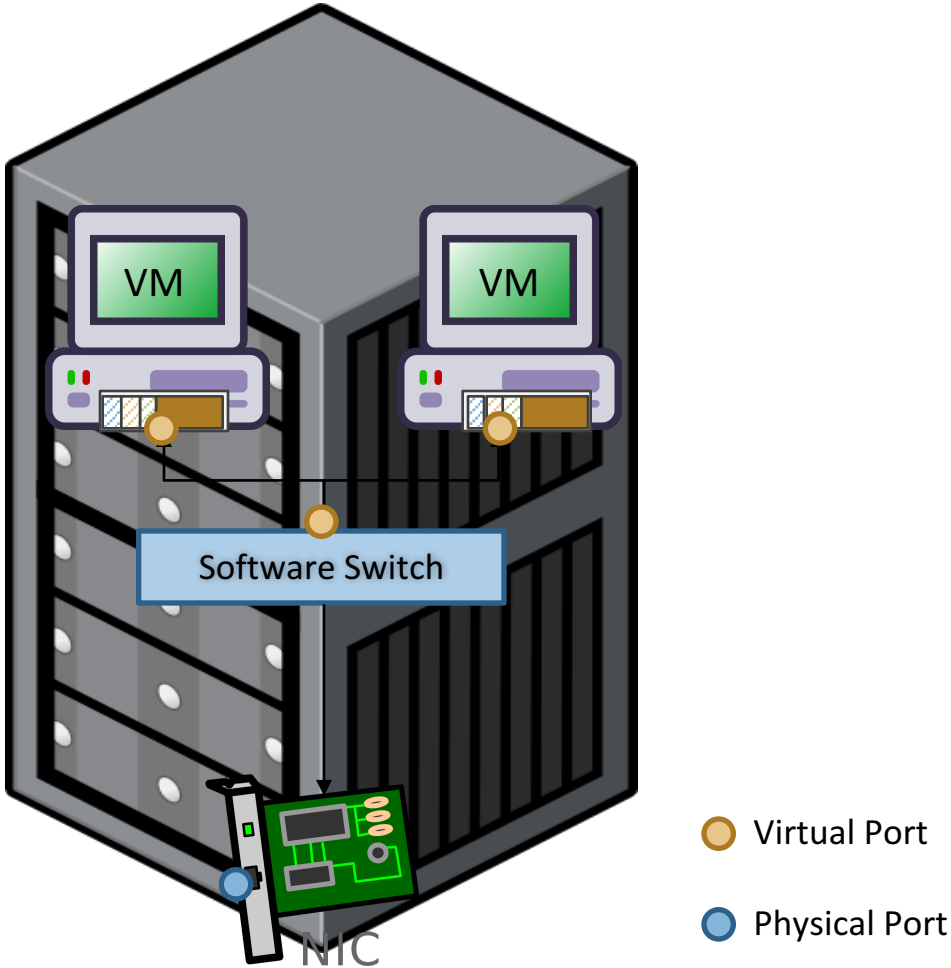
VM
VM
Software Switch
NIC
Virtual Port
Physical Port

# Approx. Number of Physical Ports vs. Virtual Ports in Global Data Centers [1]

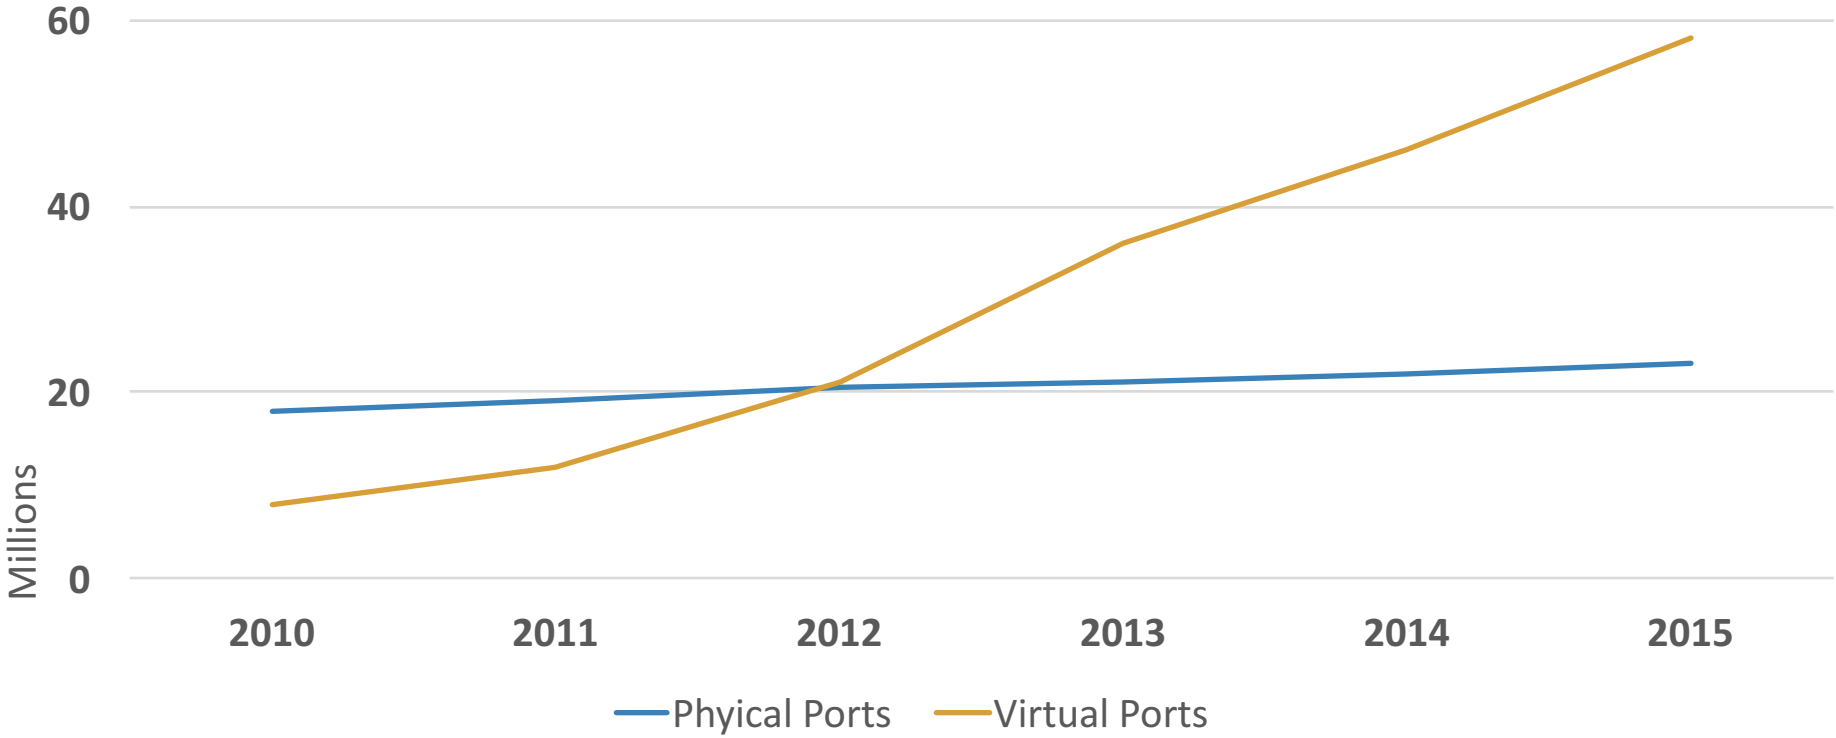


[1] Martin Casado, VMWorld 2013

It should be **EASY** to **program** software switches!

**Not really…**

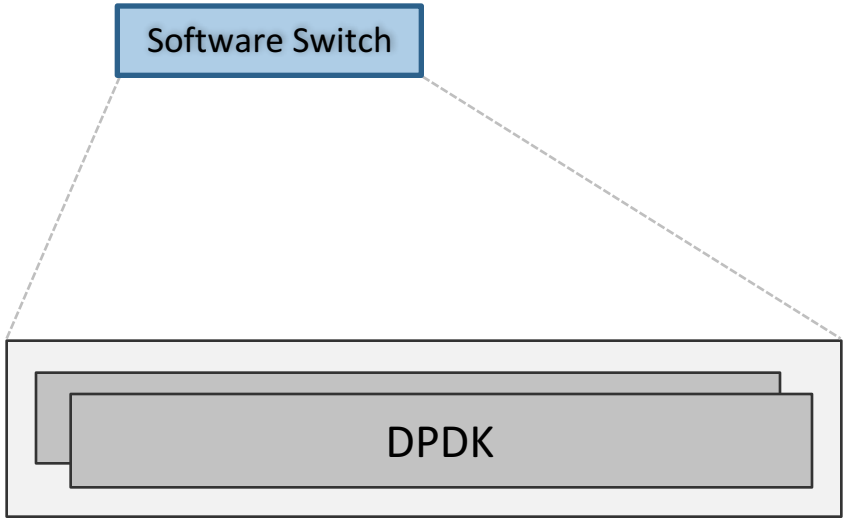
Software Switch
DPDK

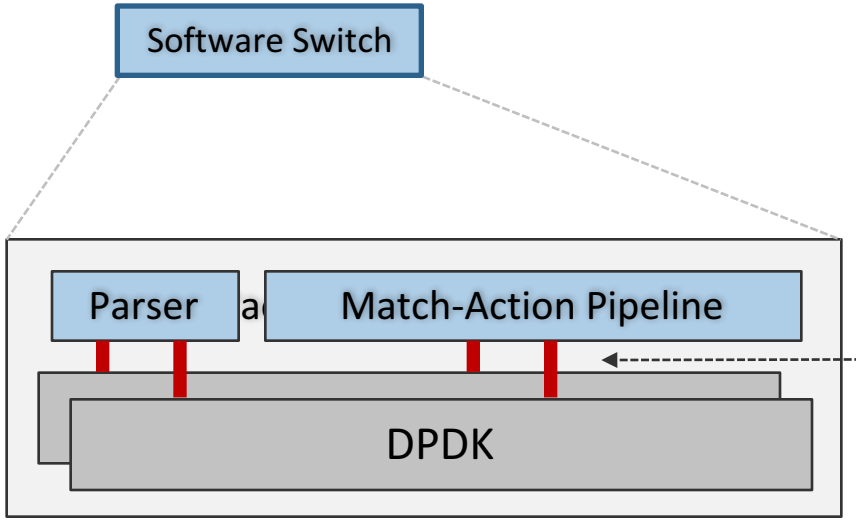


**Requires domain expertise in:**

- Network protocol design
- Kernel development

**Slow to release changes**

**Specialized APIs**

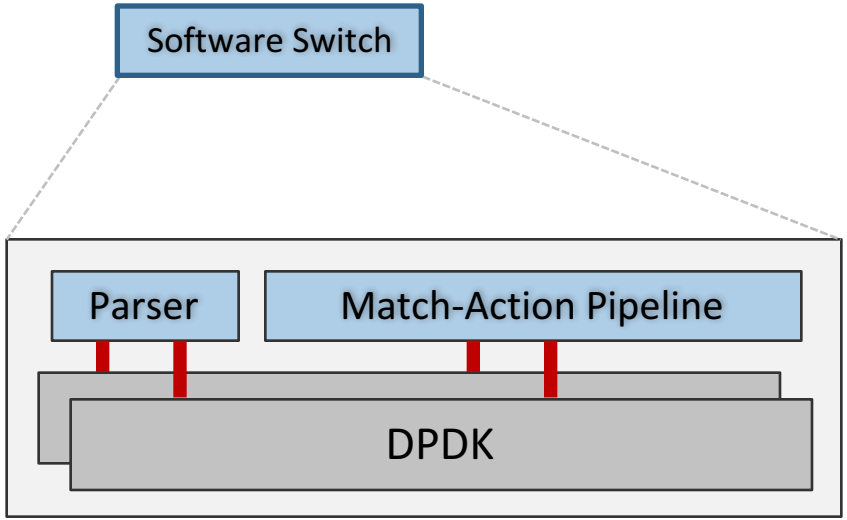


To add **TCP Flag** in Open vSwitch…

changed **20 files** and **370 lines of code**[1]

**Weeks** of development and Test

[1] https://github.com/openvswitch/ovs/commit/dc235f7fbcff

# We can do this in **4 lines** and **within minutes** with **PISCES!**

```
header_type tcp_t {
  fields {
    srcPort : 16;
    dstPort : 16;
    seqNo : 32;
    ackNo : 32;
    dataOffset : 4;
    res : 4;
    tcp_flags : 12;
    window : 16;
    checksum : 16;
    urgentPtr : 16;
  }
}
```

```
header_type flow_t {
  fields {
    ...
    tcp_flags_pad : 4;
    tcp_flags : 12;
    ...
  }
}
```

```
parser tcp {
  extract(tcp);
  set_metadata(flow.tcp_flags,
               tcp.tcp_flags);
  return ingress;
}
```

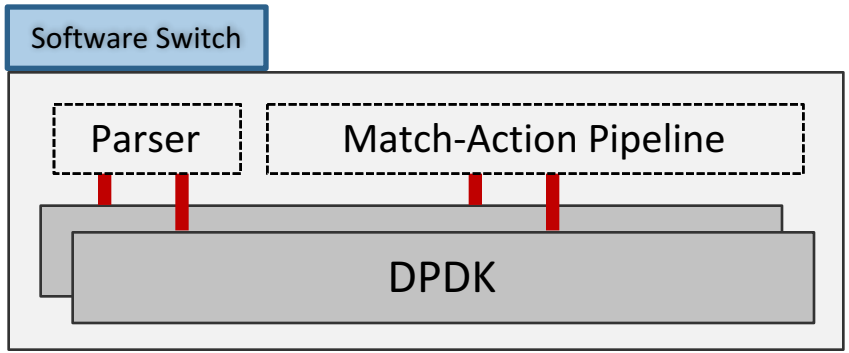
Software Switch
Parser
Match-Action Pipeline
DPDK

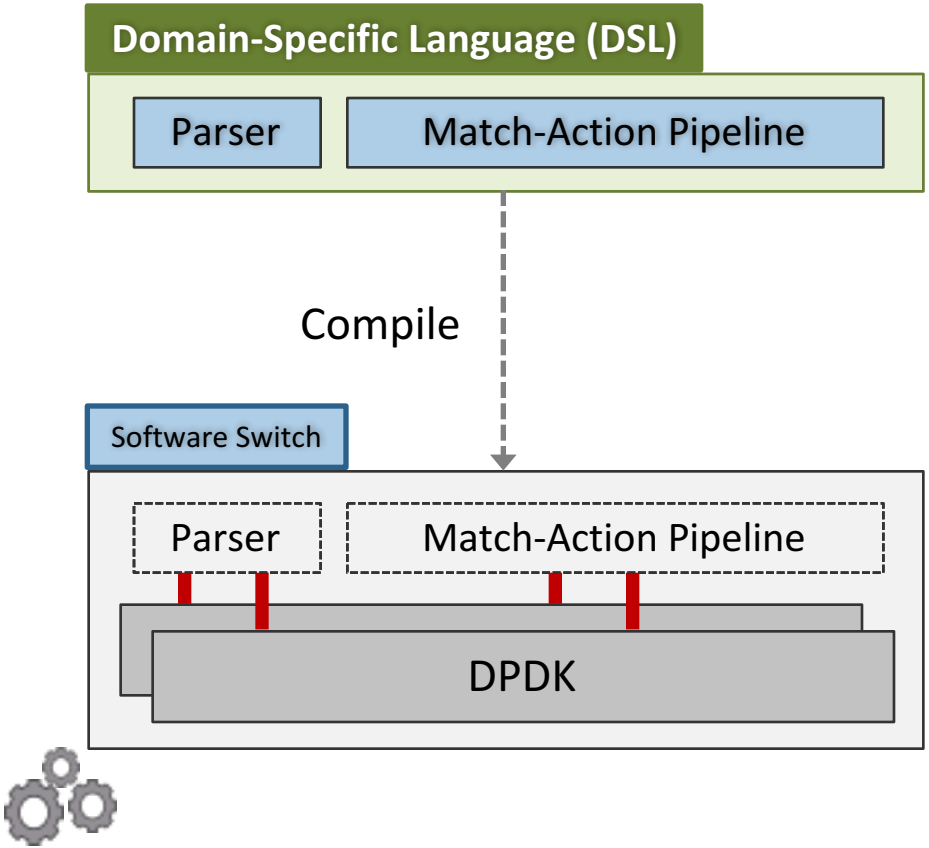


## TCP Header

```
header_type tcp_t {
  fields {
    srcPort : 16;
    dstPort : 16;
    seqNo : 32;
    ackNo : 32;
    dataOffset : 4;
    res : 4;
    window : 16;
    checksum : 16;
    urgentPtr : 16;
  }
}
parser tcp {
  extract(tcp);
  return ingress;
}
...
```

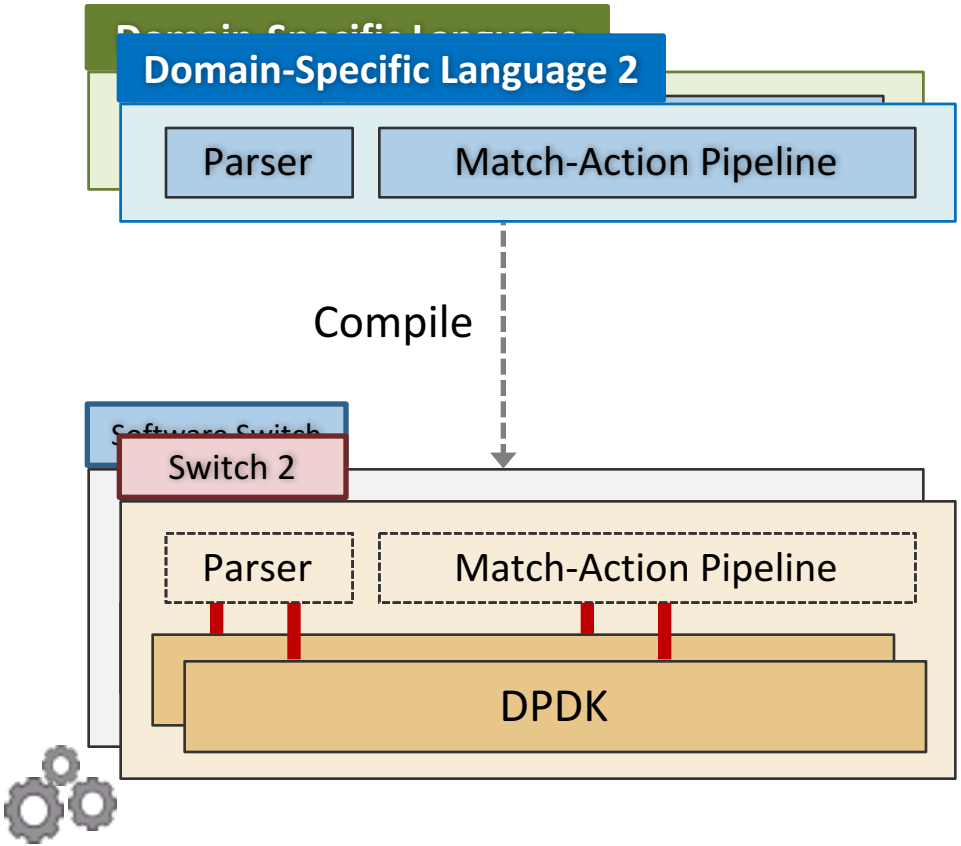


**PISCES** is a **software switch** that takes

- a **Domain-Specific Language** input

  and outputs a customized

- a **Software Switch Target**

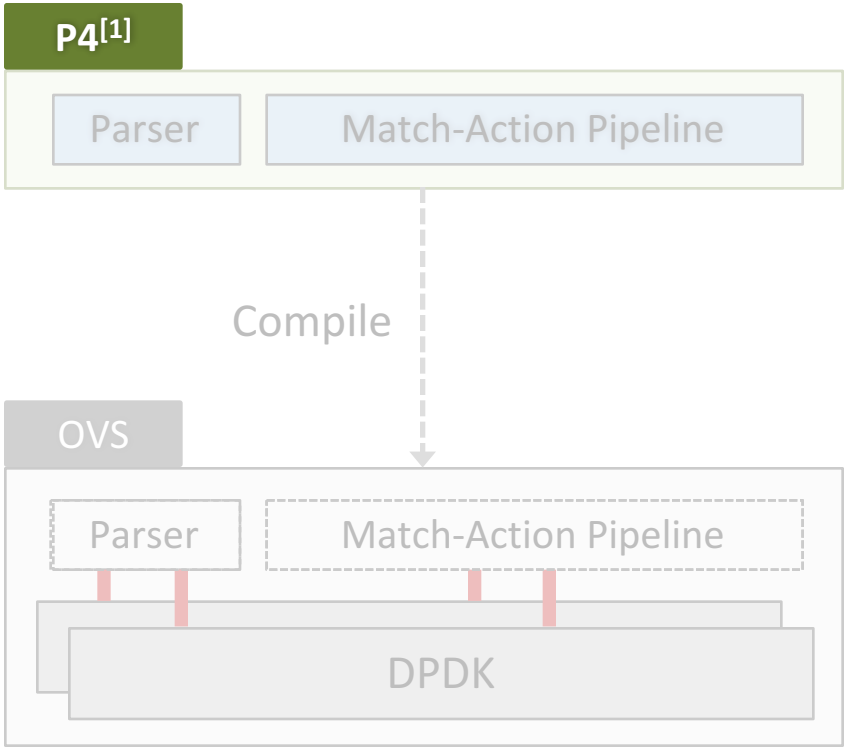


P4 is an **open-source language**.[1]

Easily defines

- **Packet headers and fields**
- **Parser**
- **Actions**
- **Match-Action Tables**

[1] http://p4.org

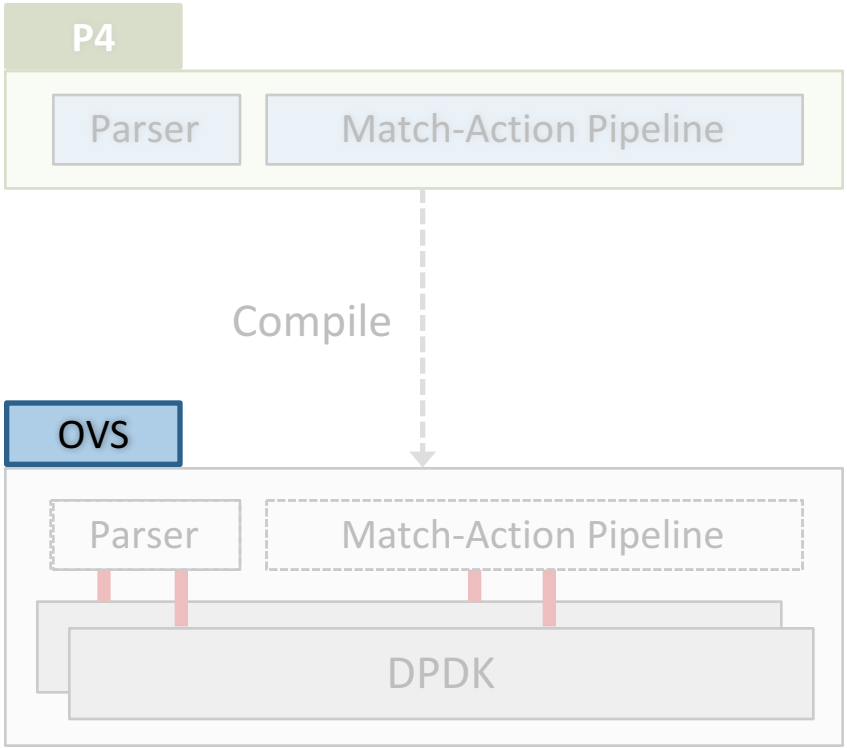


**341** lines of P4 code

Native OVS
Packet Processing Logic

**14,535** lines of C code

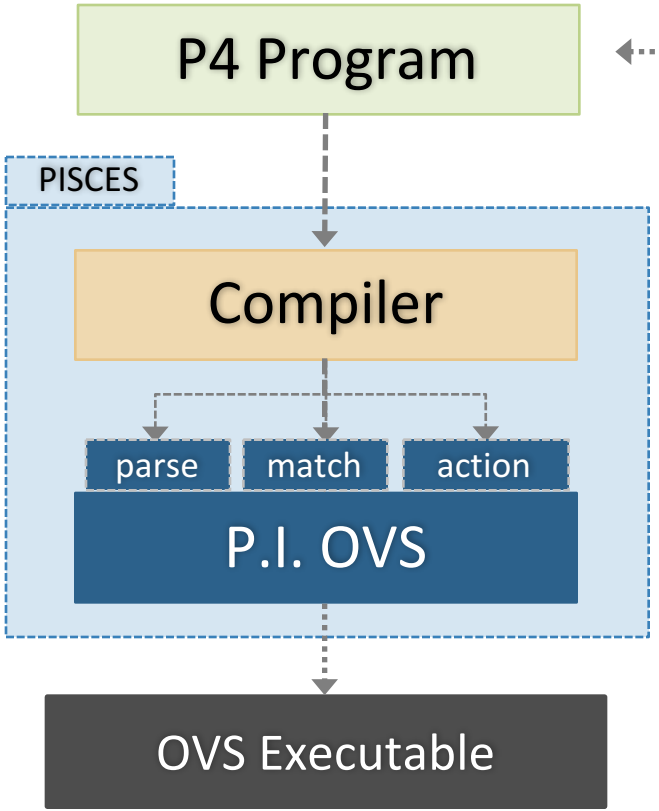


```
header_type tcpv2_t {
  fields {
    srcPort : 16;
    dstPort : 16;
    seqNo : 32;
    ackNo : 32;
    dataOffset : 4;
    res : 4;
    tcp_flags : 8;
    window : 16;
    checksum : 16;
    urgentPtr : 16;
  }
}
parser tcpv2 {
  extract(tcpv2);
  set_metadata(flow.tcp_flags,
               tcpv2.tcp_flags);
  return ingress;
}
...
```

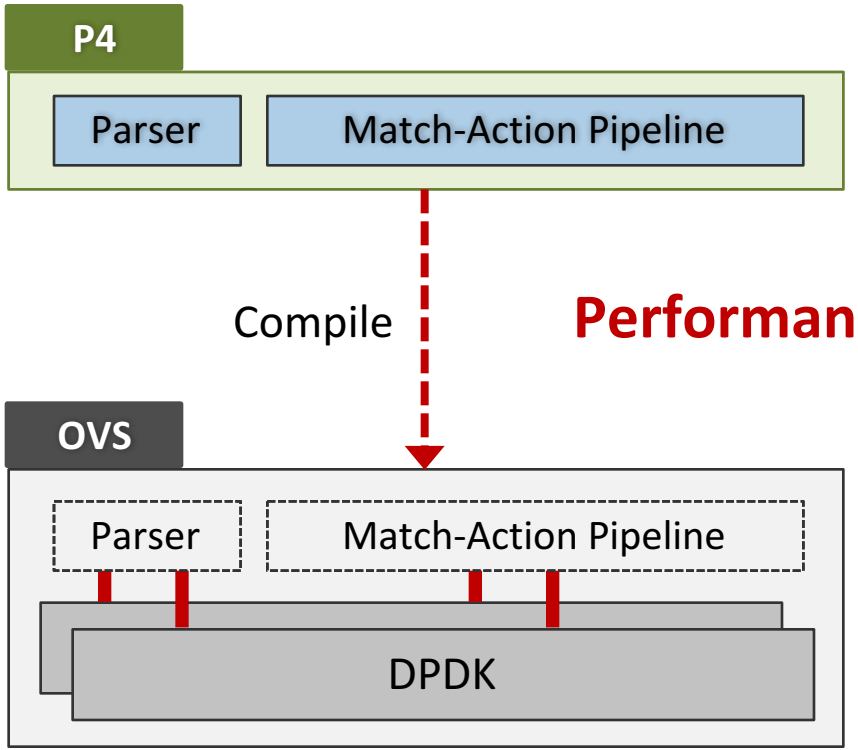


**Performance Overhead?**

Throughput on **Eth + IPv4 + ACL** benchmark application

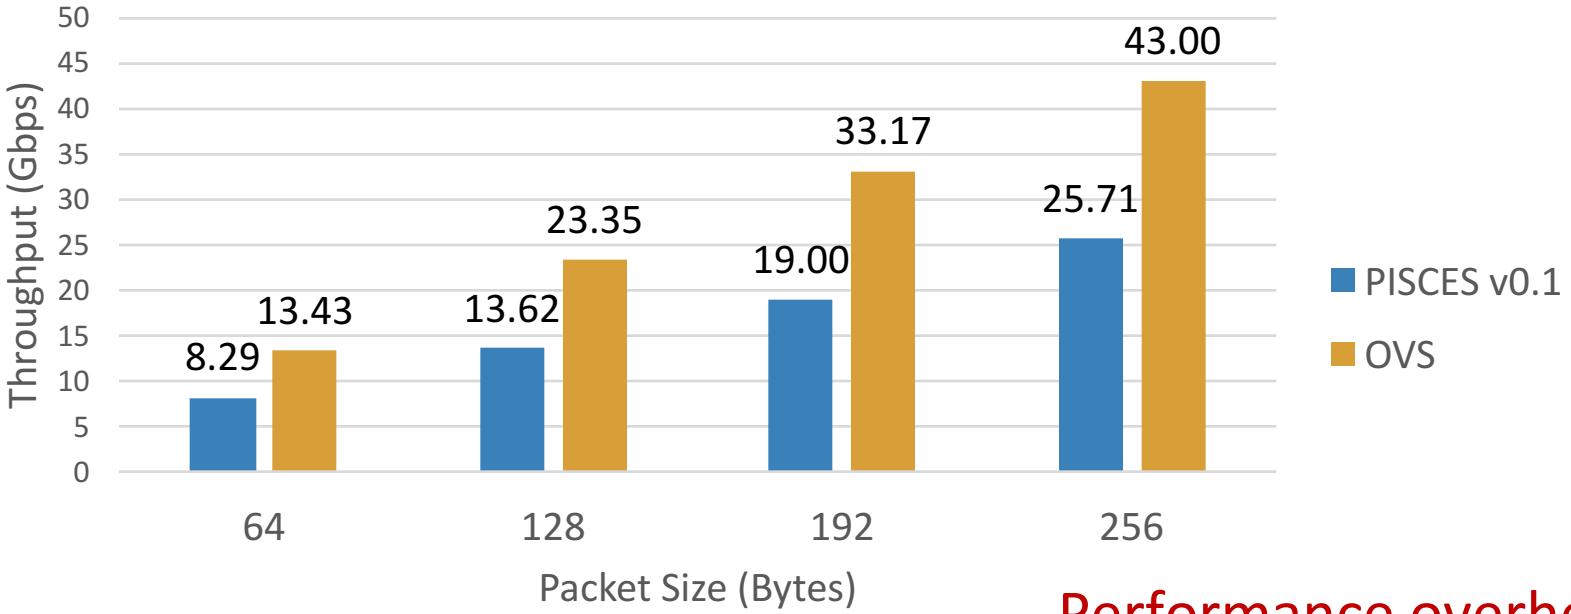


Performance overhead of

**~40%**

# Cause for the Overhead

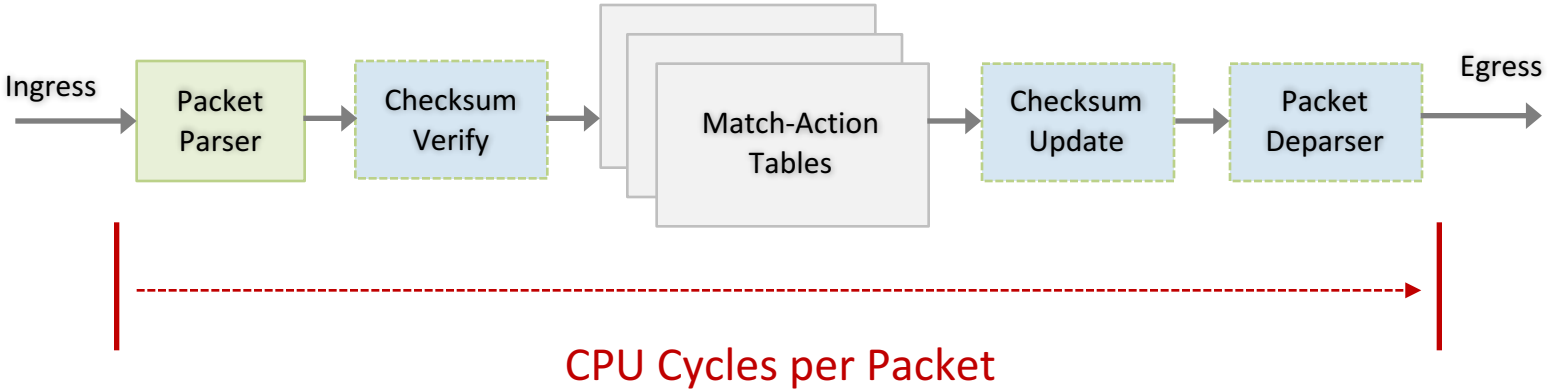

Extra CPU cycles are consumed by

- Checksum computation
- Packet header editing mode

and more…

# Checksum Inefficiencies

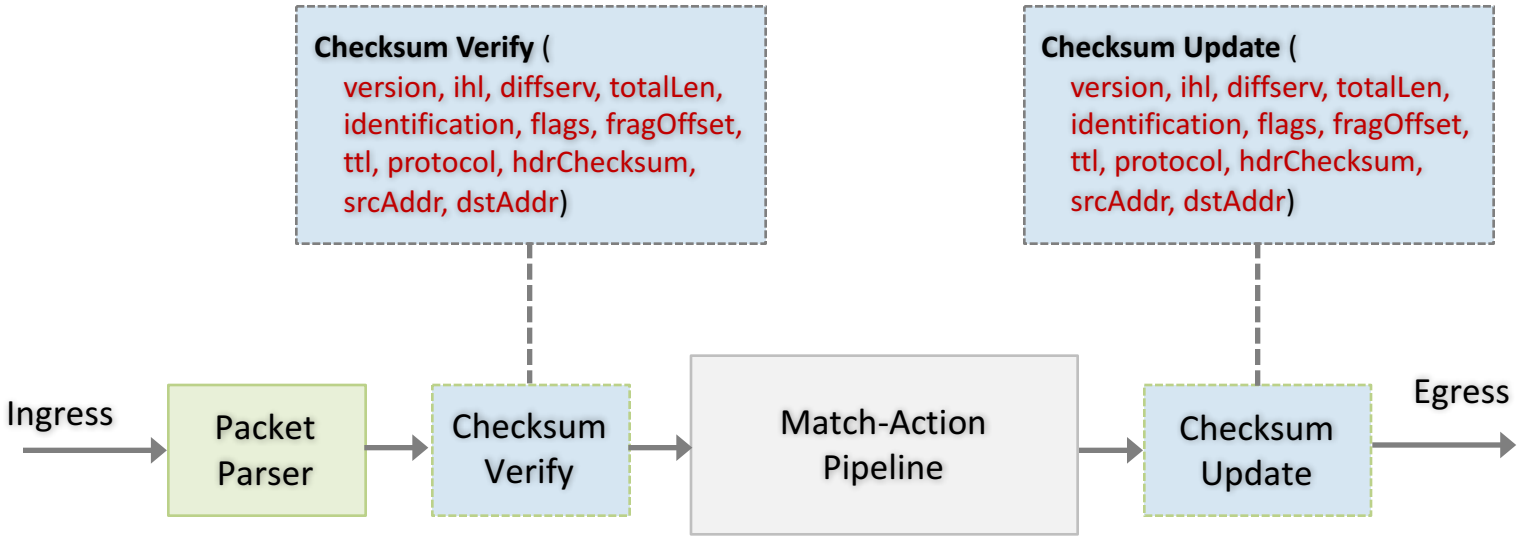

# Checksum Inefficiencies

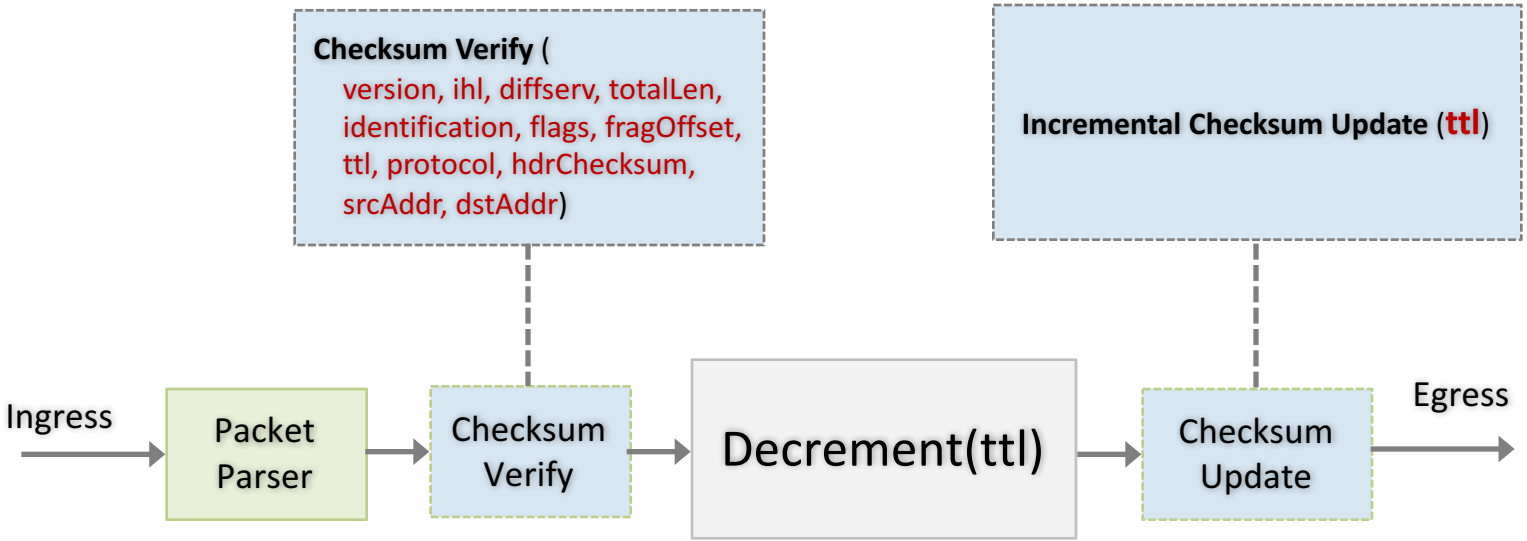

# Post-Pipeline Editing

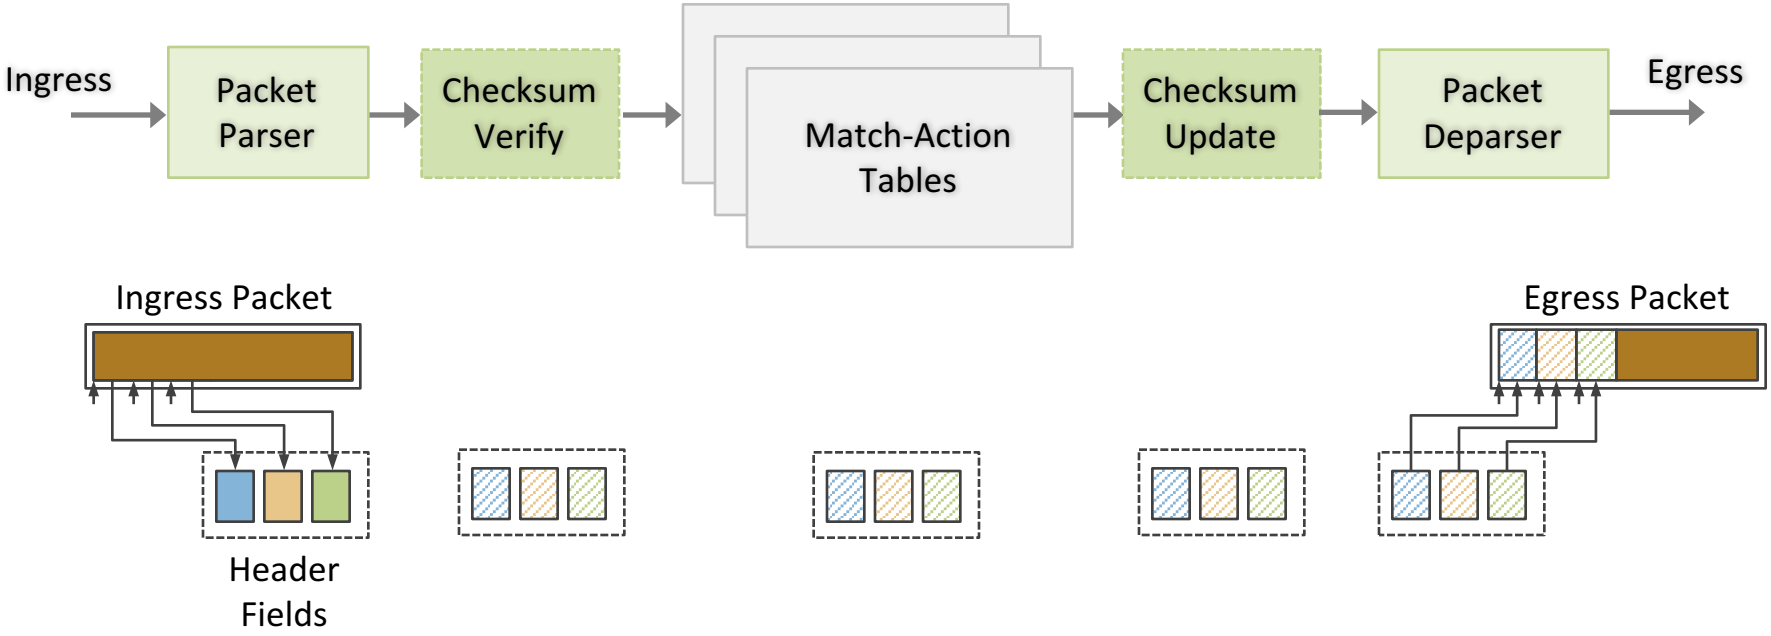

# Inline Editing

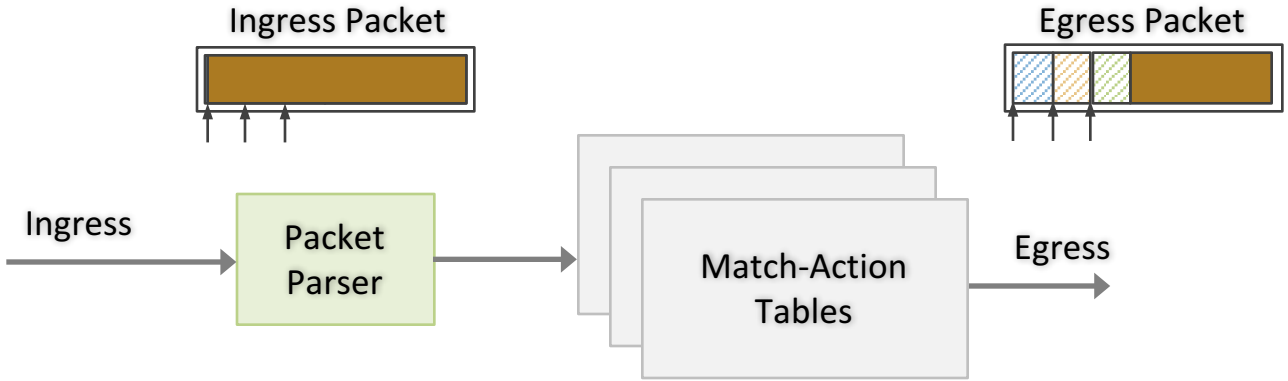

| Editing Mode | Advantage | Disadvantage |
| --- | --- | --- |
| Post-Pipeline | | Extra copy of headers |
| Inline | No extra copy of headers | |

PISCES automatically chooses between

- **Inline Editing**
- **Post-pipeline Editing**

Throughput on **Eth + IPv4 + ACL** benchmark application

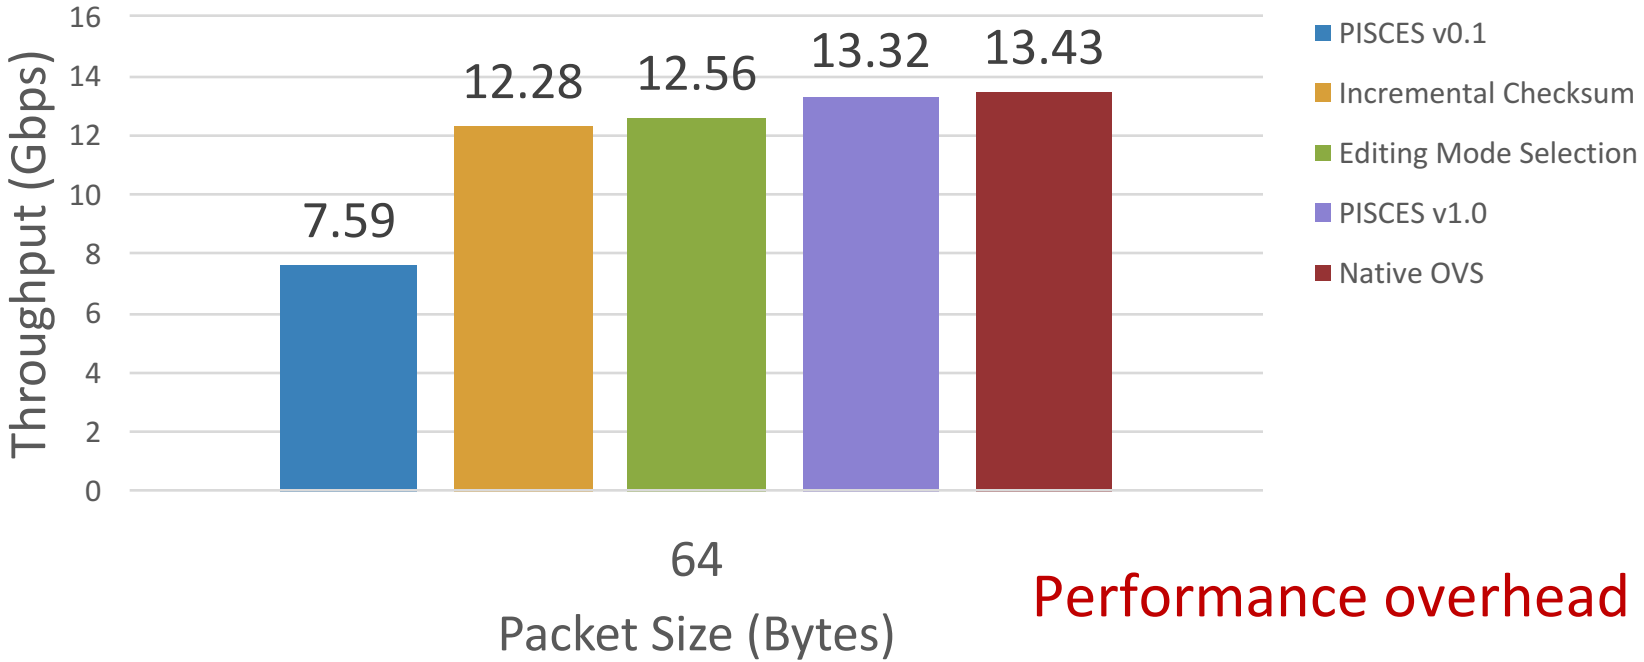


Performance overhead of

**< 2%**

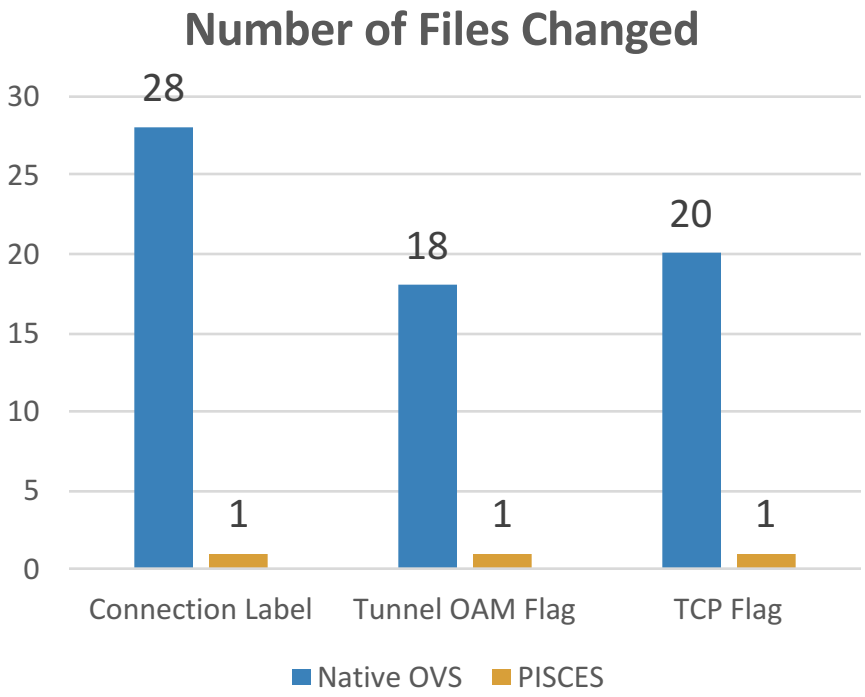
Number of Files Changed
30
25
20
15
10
5
0
28
1
18
1
20
1
Connection Label
Tunnel OAM Flag
TCP Flag
Native OVS
PISCES

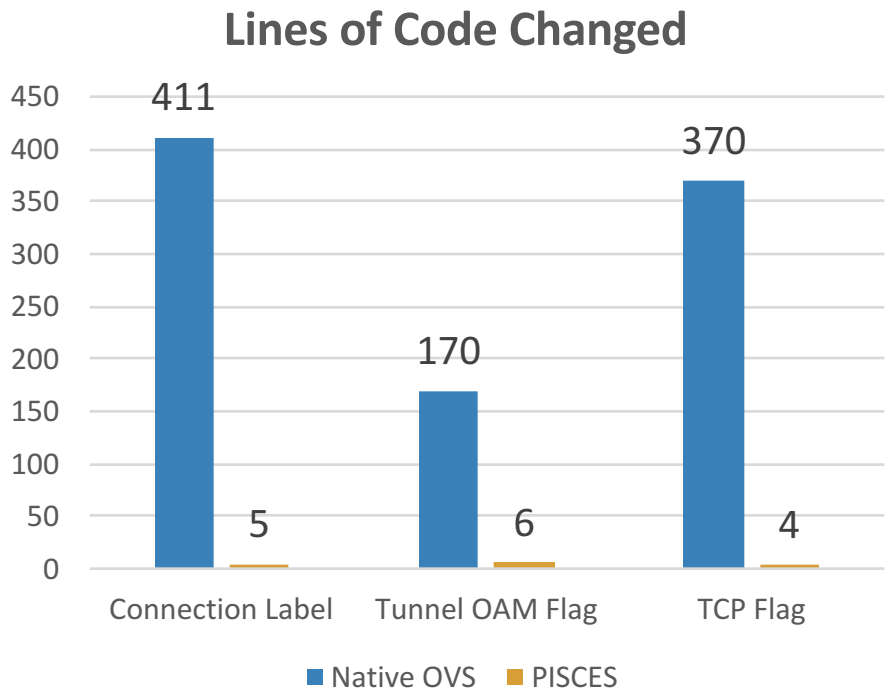
Lines of Code Changed
450
400
350
300
250
200
150
100
50
0
411
5
170
6
370
4
Connection Label
Tunnel OAM Flag
TCP Flag
Native OVS
PISCES

- A method to **quickly develop and deploy** packet processing logic on a software switch
- With hardly **any performance cost!**


**Learn more** and **Try** PISCES here:
http://pisces.cs.princeton.edu